

ECE 6504

Embedded and Controls Systems Security: Stack Smashing III

Reading:

T: *The Geometry of Innocent Flesh on the Bone*

N: *Code Injection Attacks on Harvard-Architecture Devices and Can DREs Provide Long-Lasting Security?*

“If at first you don't succeed, destroy all evidence that you tried.”

–Steven Wright

preventing code injection attacks:
lessons from history

older exploit code:

```
execve( "/bin/sh", NULL, NULL );
```



kernel executes

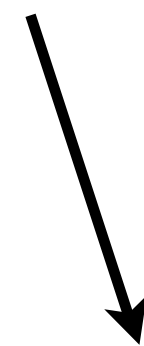
program → a 'system' call

```
#include <stdlib.h>

int main() {
    execve("/bin/sh", NULL, NULL);
}
```

```
#include <stdlib.h>
```

```
int main() {  
    execve("/bin/sh", NULL, NULL);  
}
```

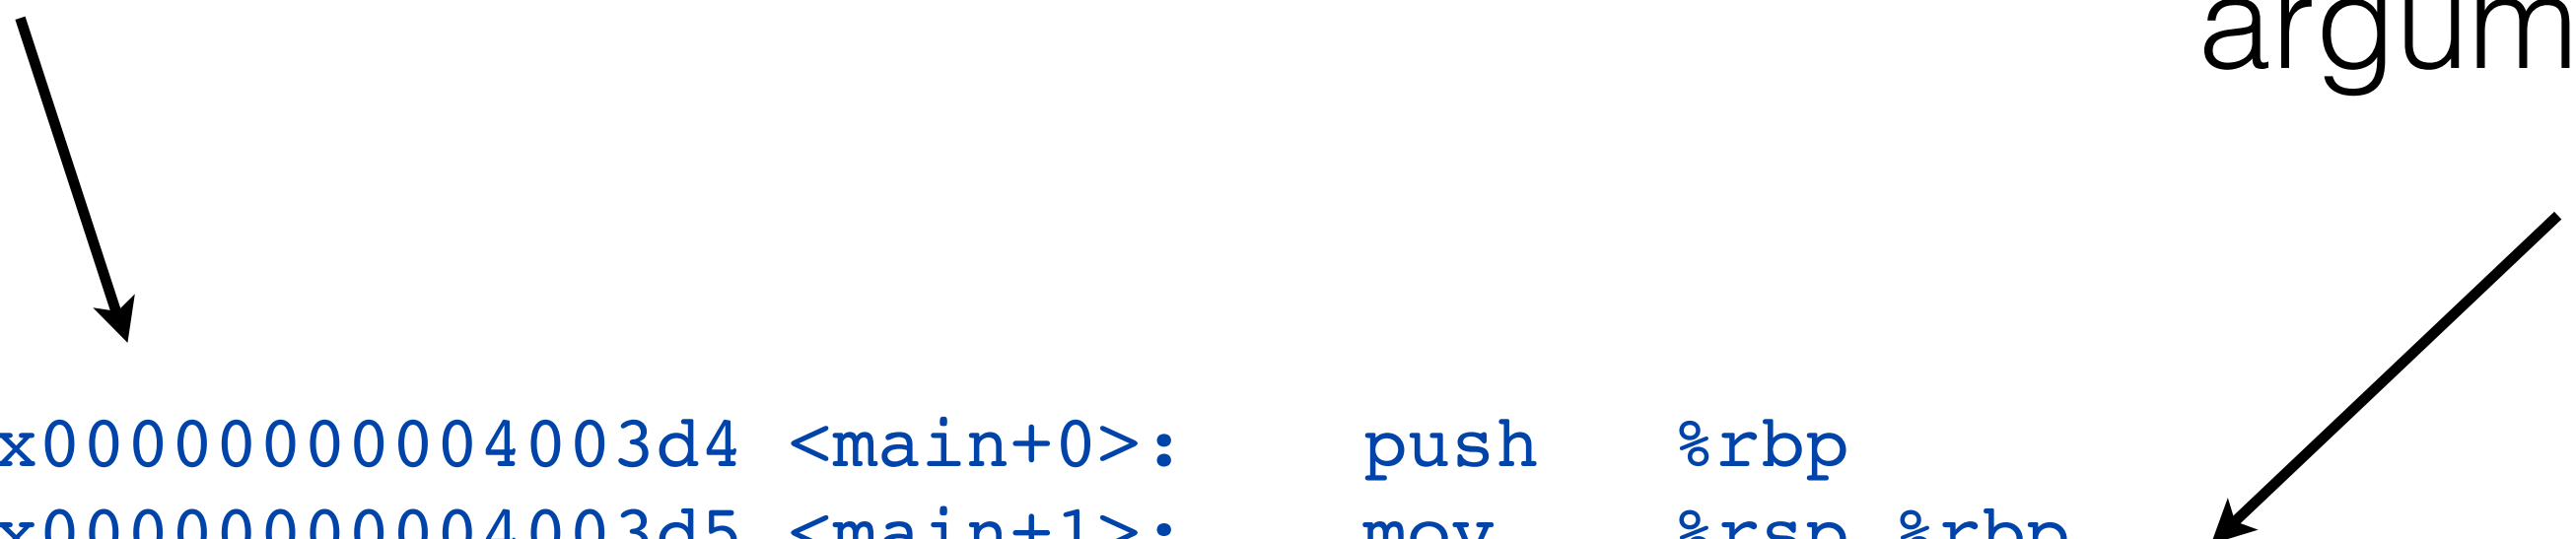


0x00000000004003d4	<main+0>:	push	%rbp
0x00000000004003d5	<main+1>:	mov	%rsp,%rbp
0x00000000004003d8	<main+4>:	mov	\$0x0,%rdx
0x00000000004003dd	<main+9>:	mov	\$0x0,%rsi
0x00000000004003e2	<main+14>:	mov	\$0x46c610,%rdi
0x00000000004003e7	<main+19>:	callq	0x40ad30 <execve>
0x00000000004003ec	<main+24>:	leaveq	
0x00000000004003ed	<main+25>:	retq	

```
#include <stdlib.h>
```

```
int main() {  
    execve("/bin/sh", NULL, NULL);  
}
```

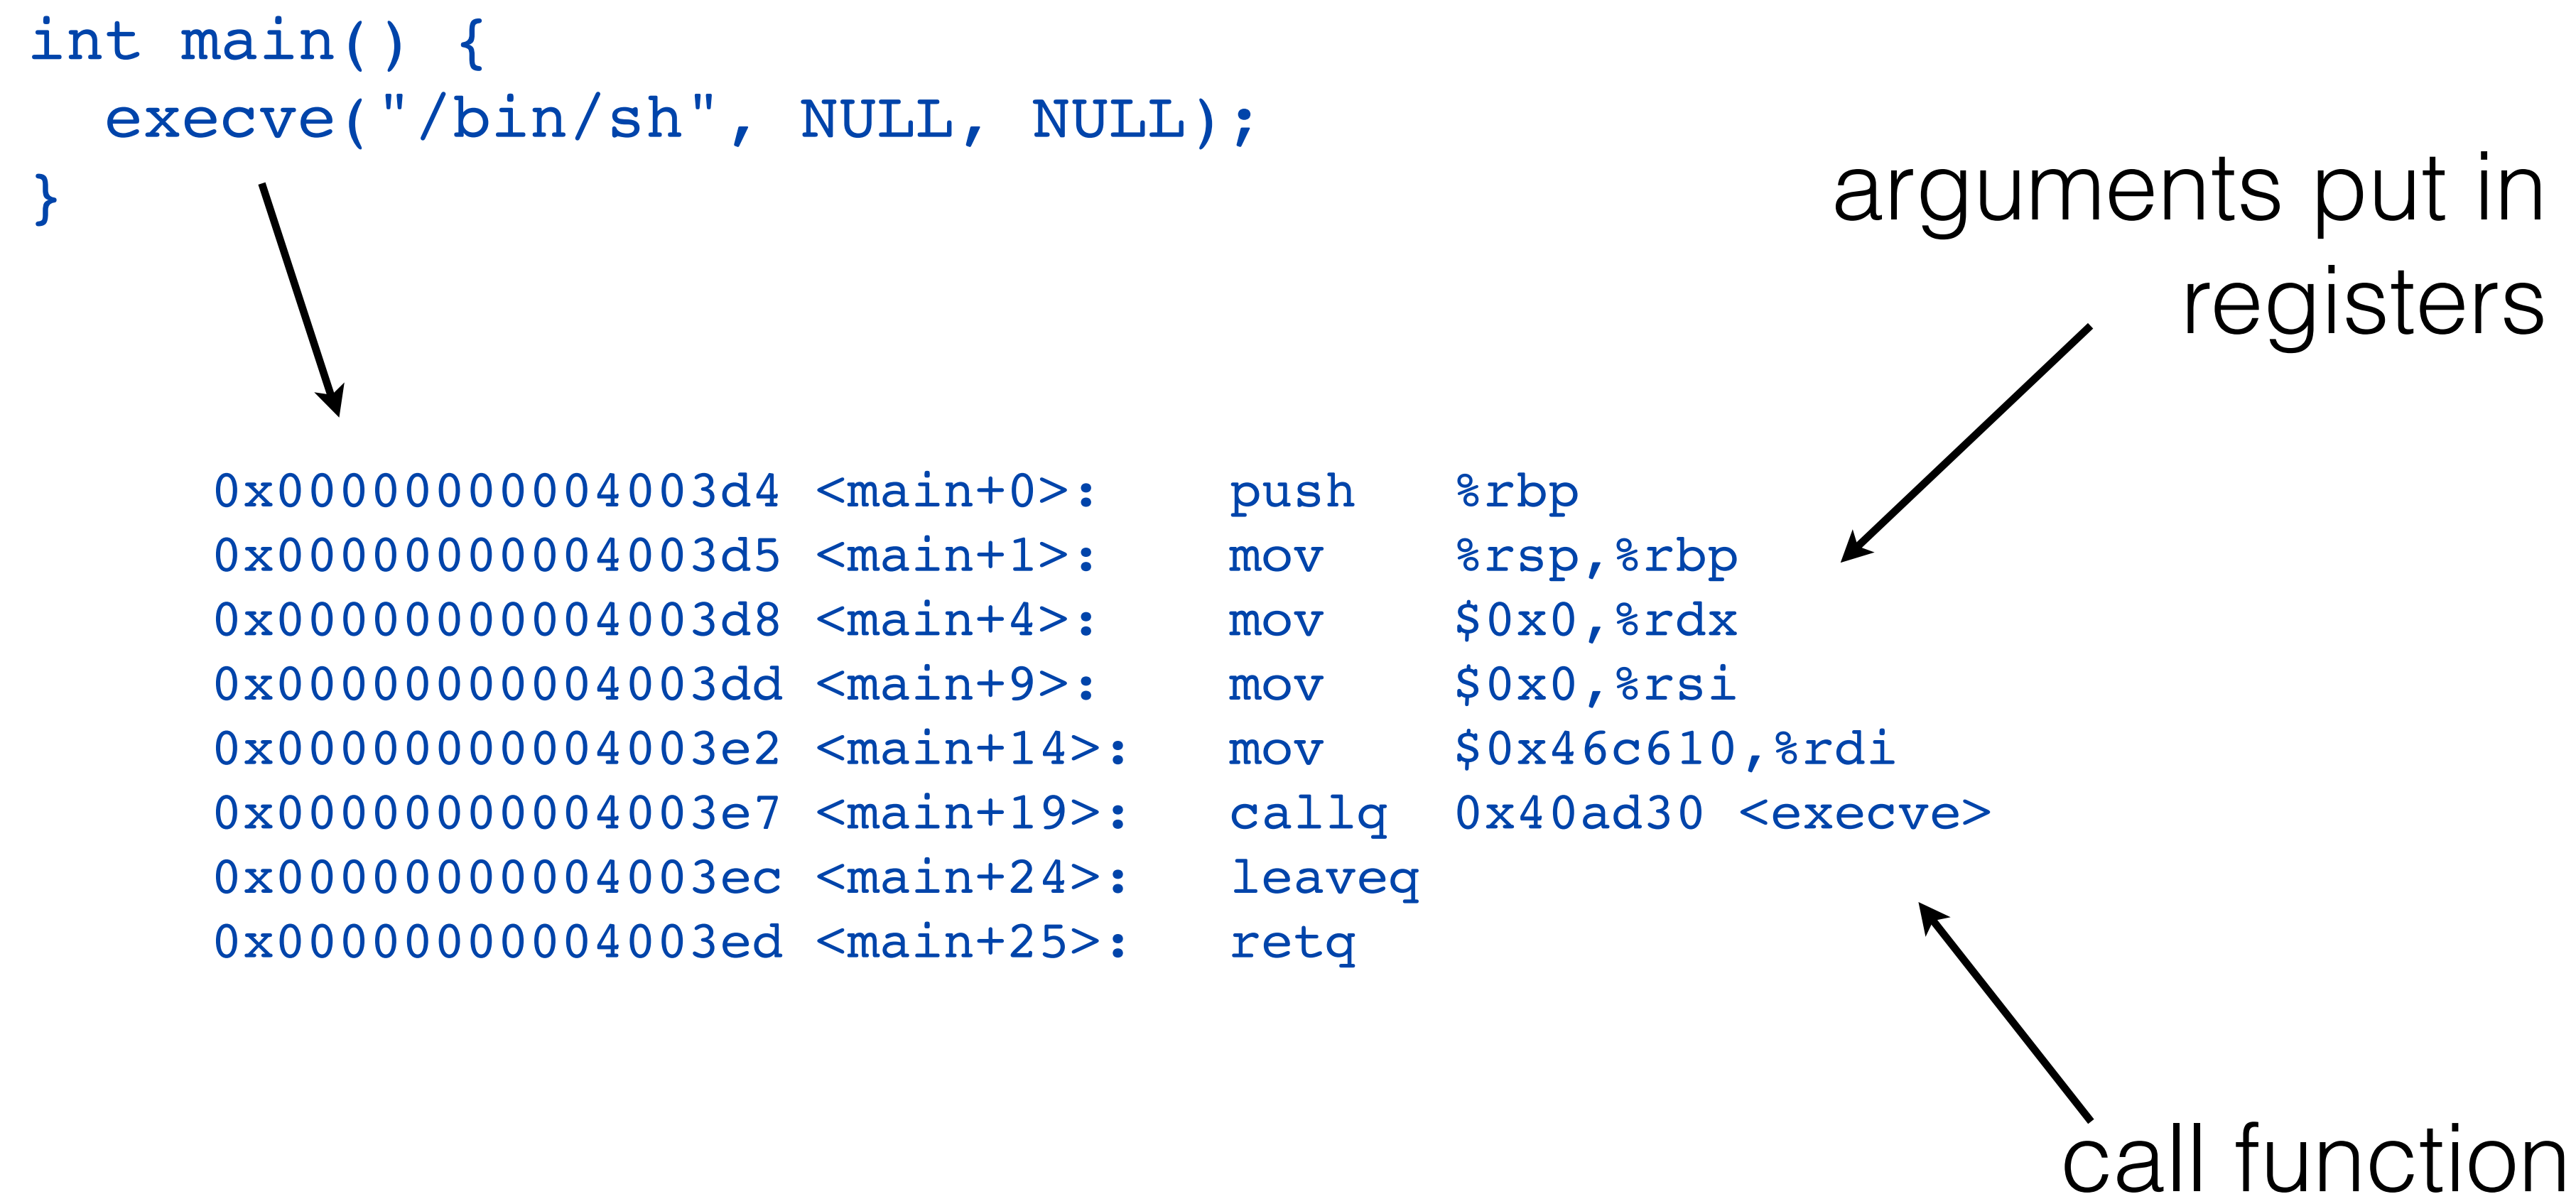
arguments put in
registers



```
0x00000000004003d4 <main+0>:  push    %rbp  
0x00000000004003d5 <main+1>:  mov     %rsp,%rbp  
0x00000000004003d8 <main+4>:  mov     $0x0,%rdx  
0x00000000004003dd <main+9>:  mov     $0x0,%rsi  
0x00000000004003e2 <main+14>: mov     $0x46c610,%rdi  
0x00000000004003e7 <main+19>: callq   0x40ad30 <execve>  
0x00000000004003ec <main+24>: leaveq    
0x00000000004003ed <main+25>:  retq    
```

```
#include <stdlib.h>
```

```
int main() {  
    execve("/bin/sh", NULL, NULL);  
}
```

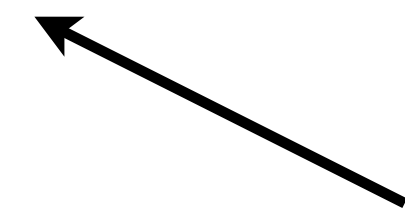


```
0x00000000004003d4 <main+0>:  push    %rbp  
0x00000000004003d5 <main+1>:  mov     %rsp,%rbp  
0x00000000004003d8 <main+4>:  mov     $0x0,%rdx  
0x00000000004003dd <main+9>:  mov     $0x0,%rsi  
0x00000000004003e2 <main+14>: mov     $0x46c610,%rdi  
0x00000000004003e7 <main+19>: callq   0x40ad30 <execve>  
0x00000000004003ec <main+24>: leaveq    
0x00000000004003ed <main+25>:  retq    
```

arguments put in
registers

call function

system calls (**execve**):
pass arguments via registers



(most OSs; e.g. Linux and Windows)

system calls (**execve**):
pass arguments via registers

(most OSs; e.g. Linux and Windows)

function calls (**system**):
pass arguments via stack

system calls (**execve**):
pass arguments via registers

(most OSs; e.g. Linux and Windows)

function calls (**system**):
pass arguments via stack

Q: assuming $W \wedge X$, why must shellcode use
system instead of **execve**

Q: assuming $W \wedge X$, why must shellcode use
`system` instead of `execve`

A: because $W \wedge X$ won't allow execution of stack

Q: assuming $W \wedge X$, why must shellcode use
`system` instead of `execve`

A: because $W \wedge X$ won't allow execution of stack



no way to get arguments into registers

Q: assuming $W \wedge X$, why must shellcode use
`system` instead of `execve`

A: because $W \wedge X$ won't allow execution of stack



no way to get arguments into registers



which `execve` requires

attacker: fine just use `system`

it's in libc



attacker: fine just use `system`

it's in libc

attacker: fine just use `system`

`system("/bin/sh")`

```
graph TD; A[it's in libc] --> B[attacker: fine just use system]; B --> C[system( "/bin/sh" )];
```

attacker:



replace 'diploma' with **execve**

attacker:



replace 'diploma' with **execve**

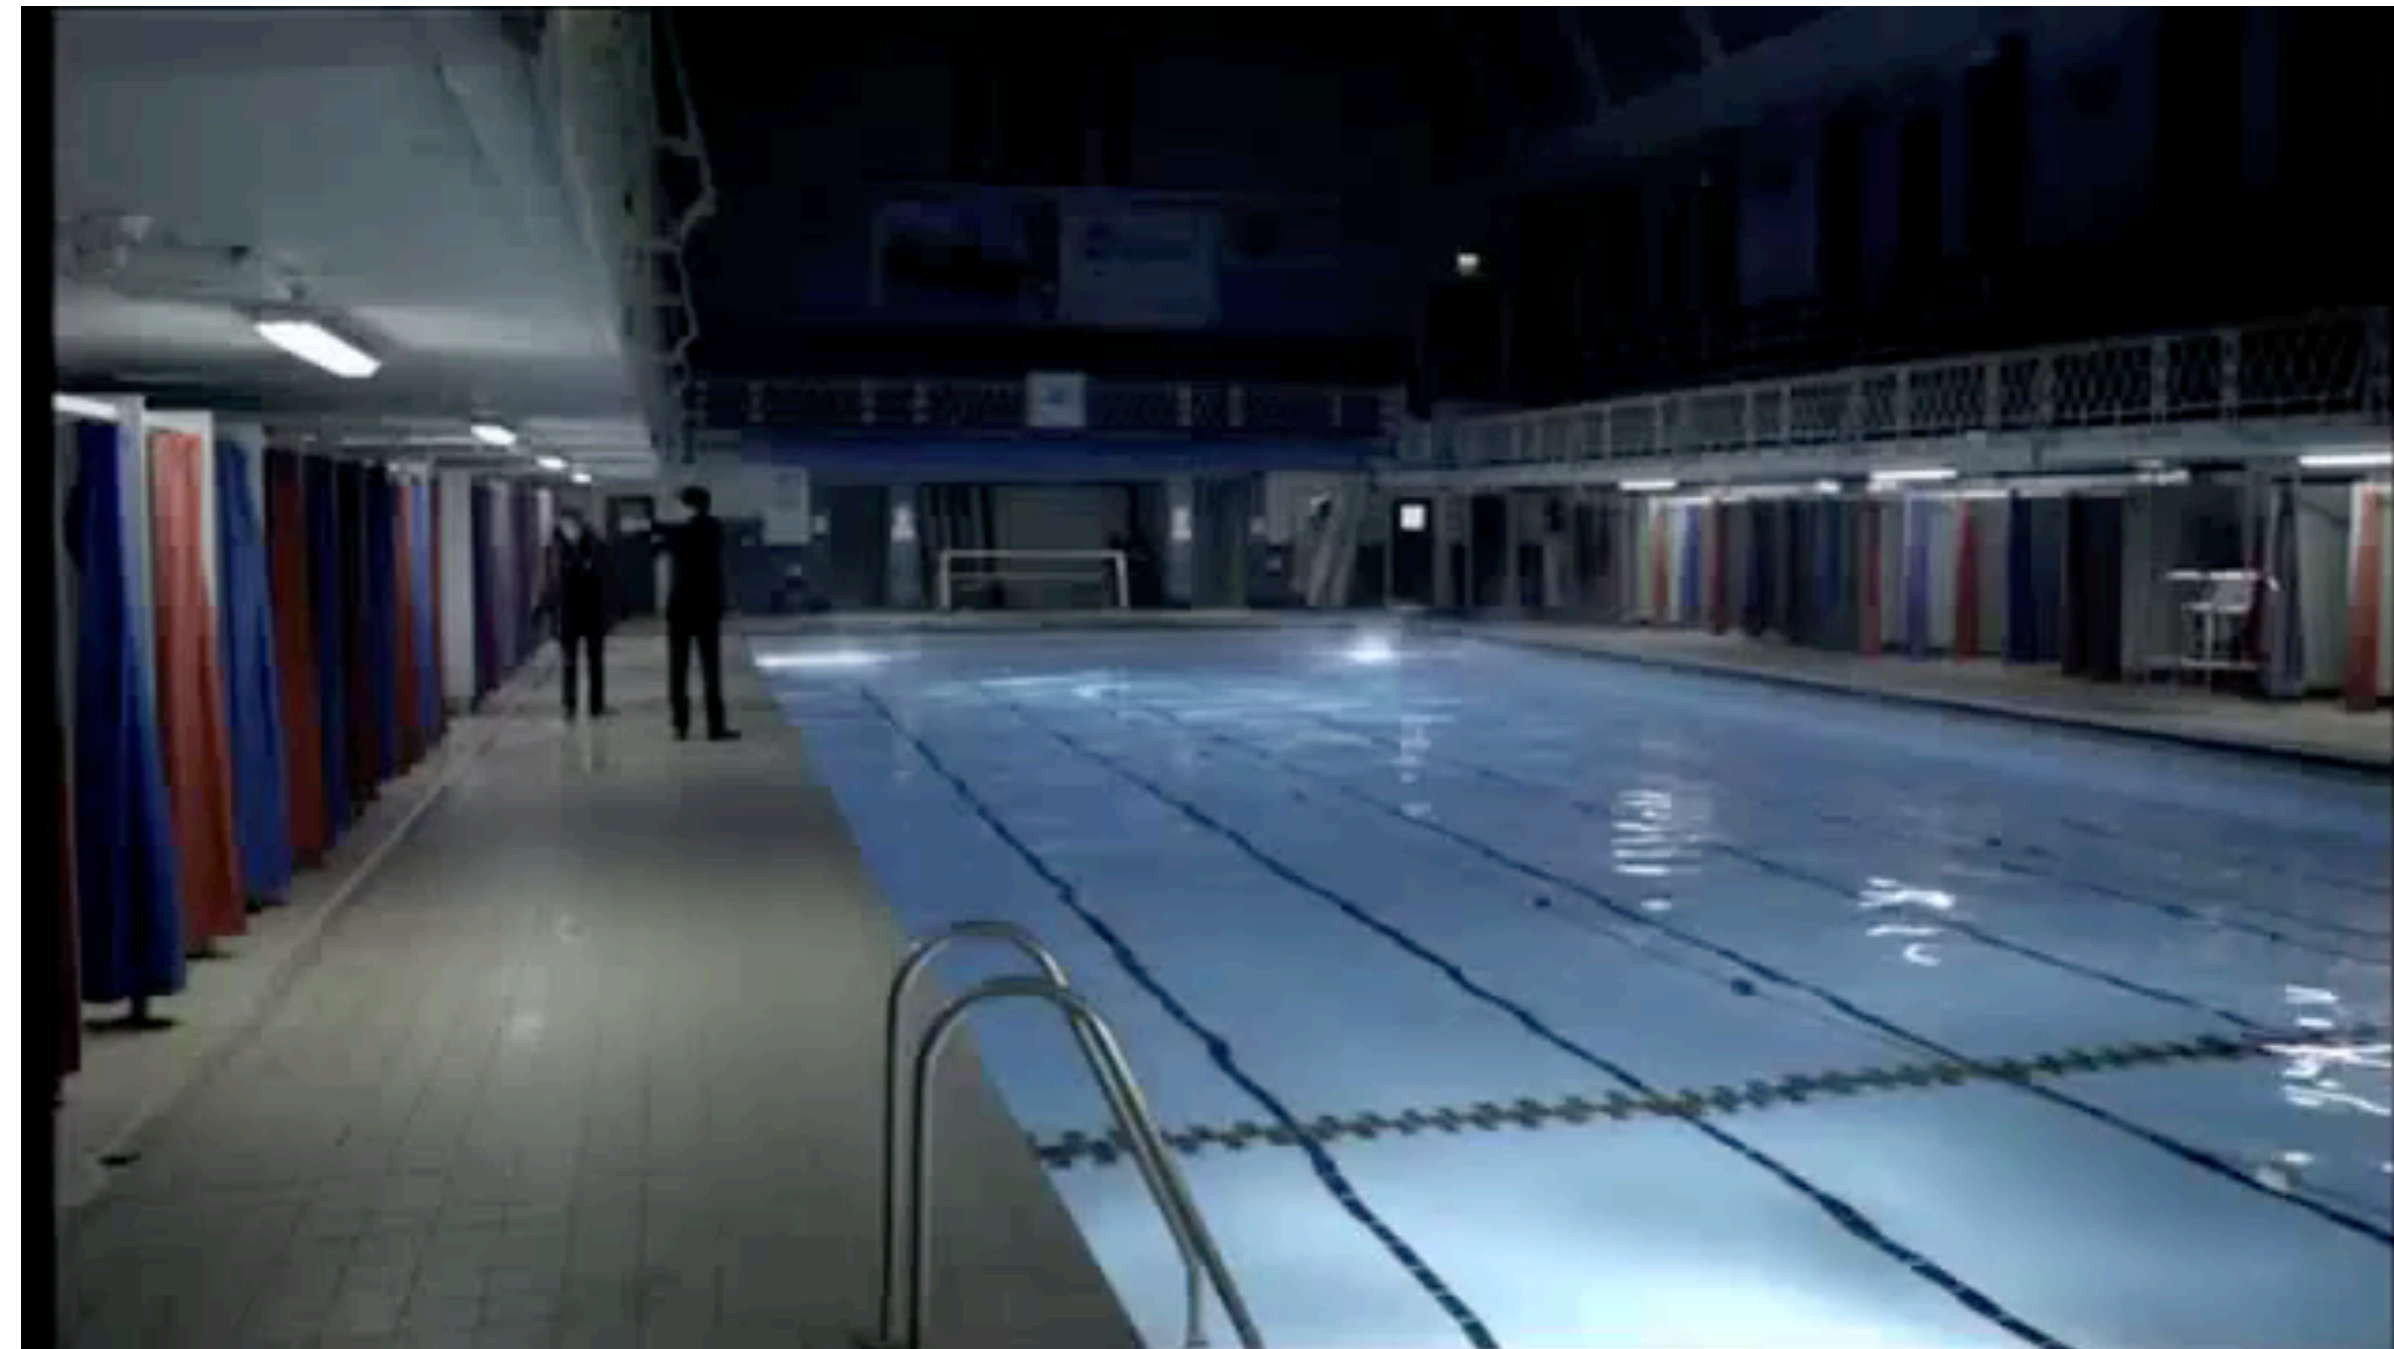
attacker: fine just use `system`



defender: let's get rid of that

excise `system` when not needed

attacker: fine just use `system`



defender: let's get rid of that

excise `system` when not needed

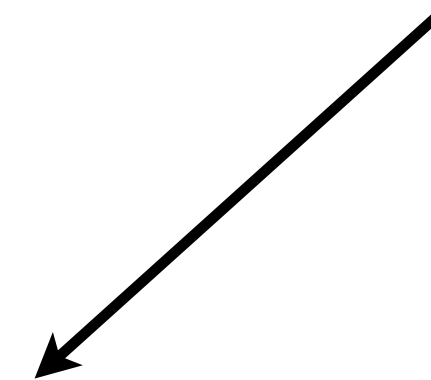
defences:

1. excise `system` when not needed
2. require all function calls to use registers for passing

defences:

1. excise **system** when not needed
2. require all function calls to use registers for passing

i.e. operate like system calls



defences:

1. excise **system** when not needed
2. require all function calls to use registers for passing

i.e. operate like system calls

need extra registers: x86-64

attacker response?:



alas, no

attacker response?:



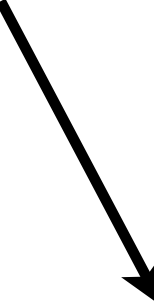
alas, no

attacker:
useful programs make system calls

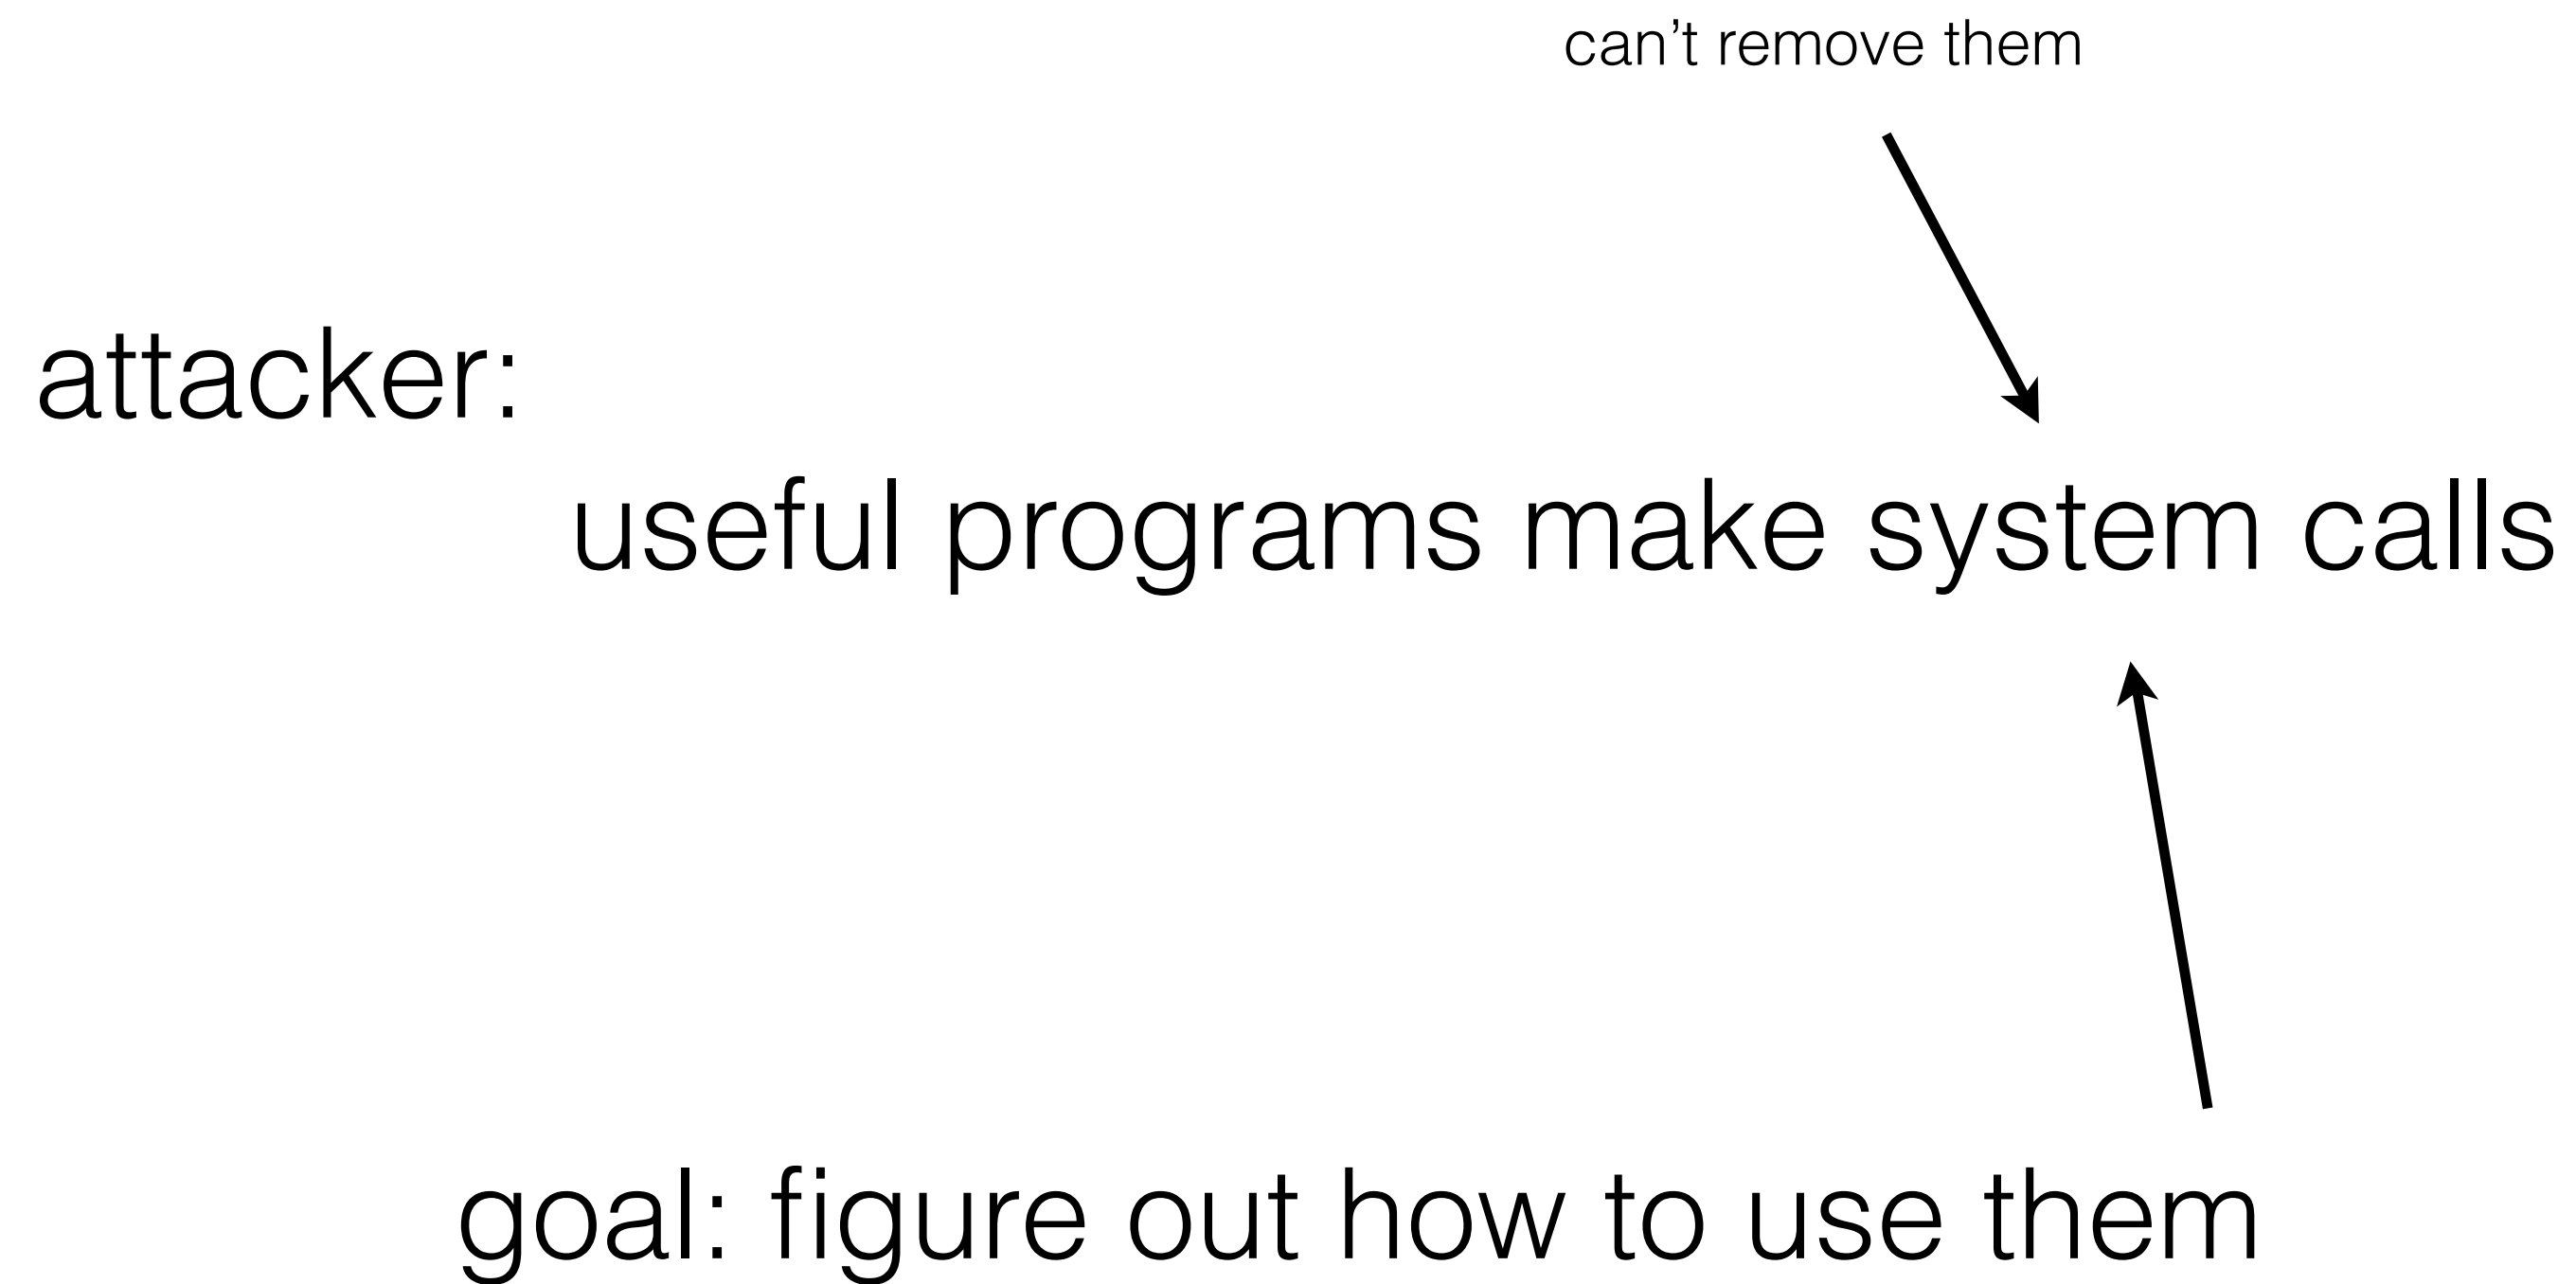
can't remove them

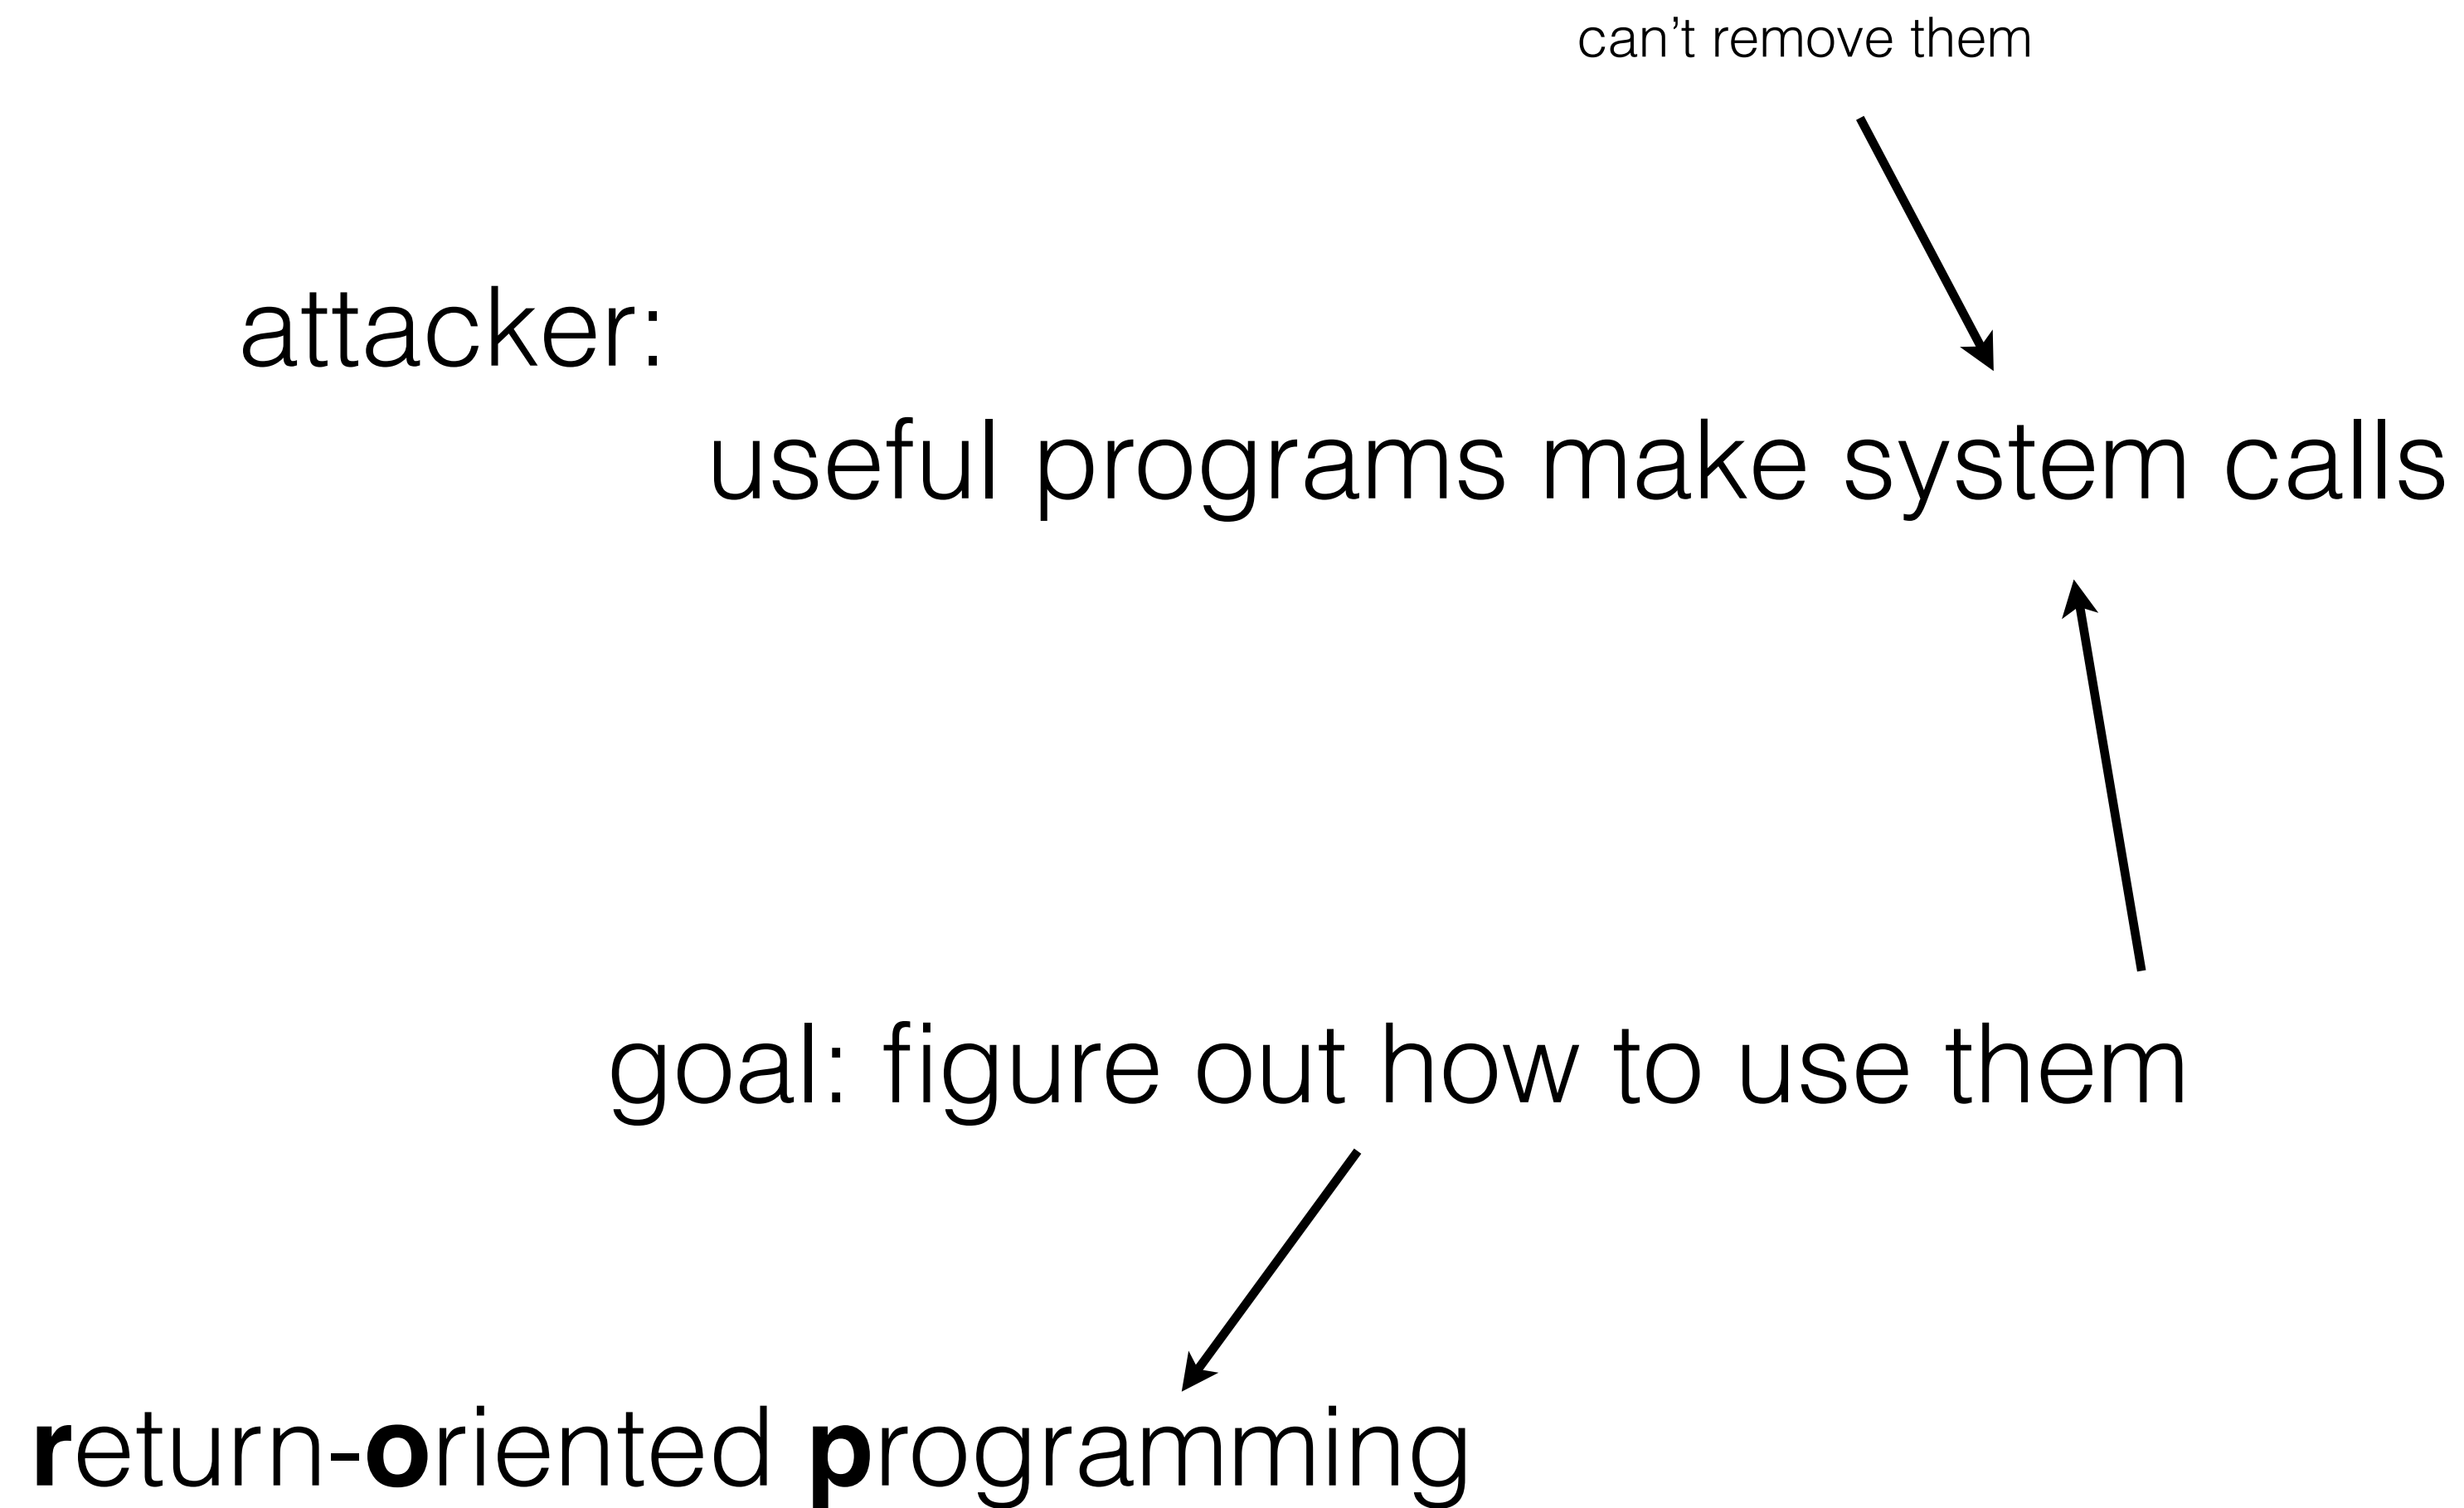
attacker:

useful programs make system calls



```
graph TD; A[can't remove them] --> B[useful programs make system calls];
```





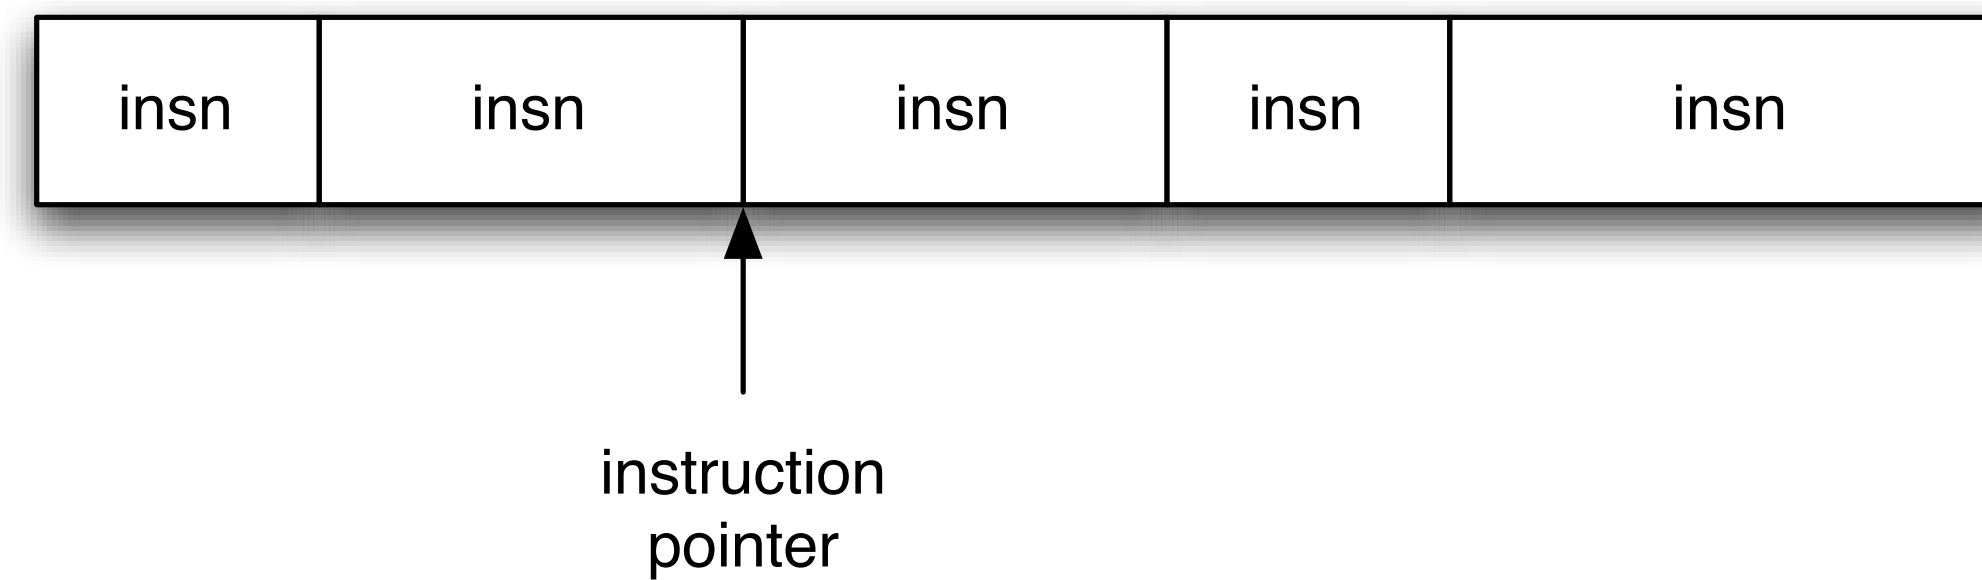
return-oriented programming...



return-oriented programming...



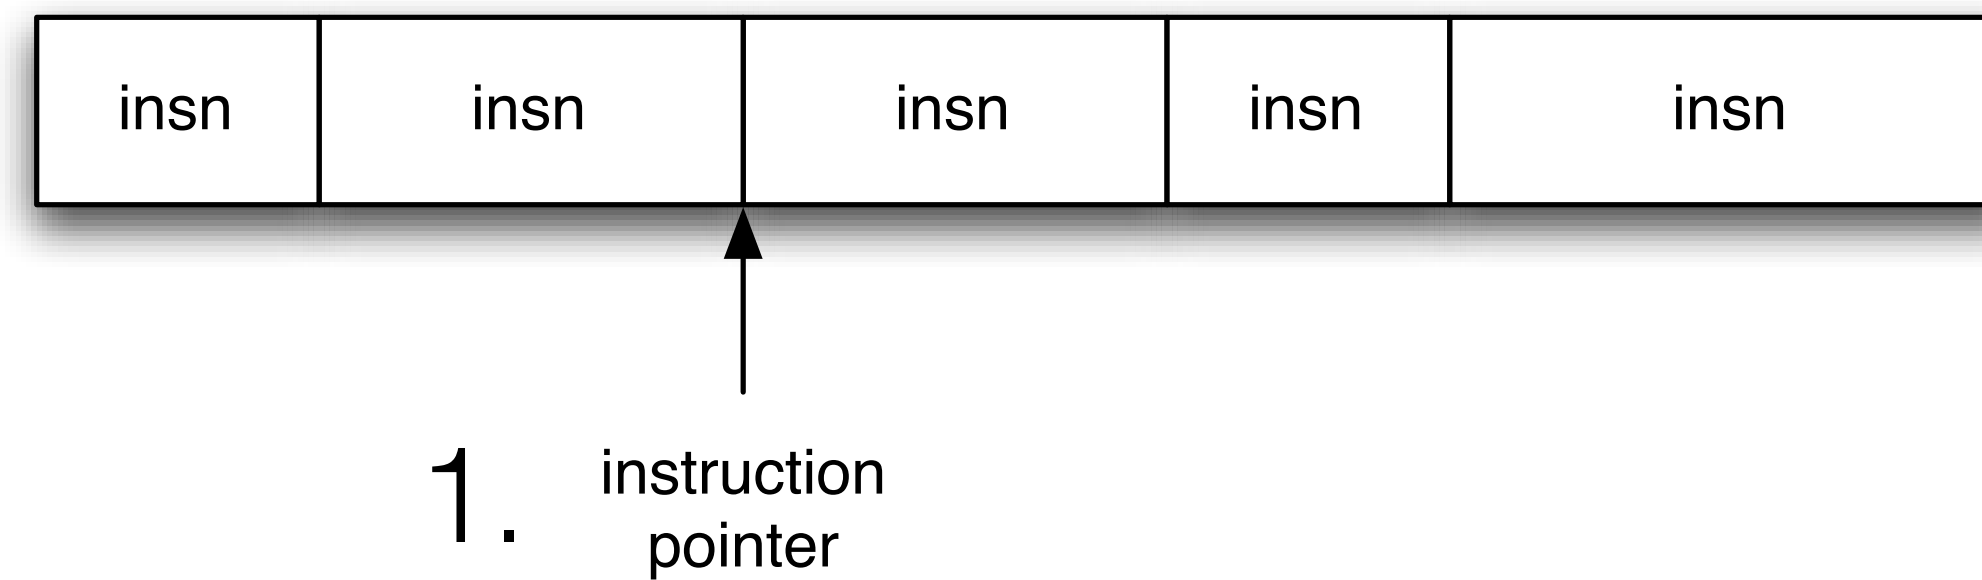
ordinary programming:



procedure:

1. IP points at instruction
2. instruction executed
3. IP incremented by length of instruction

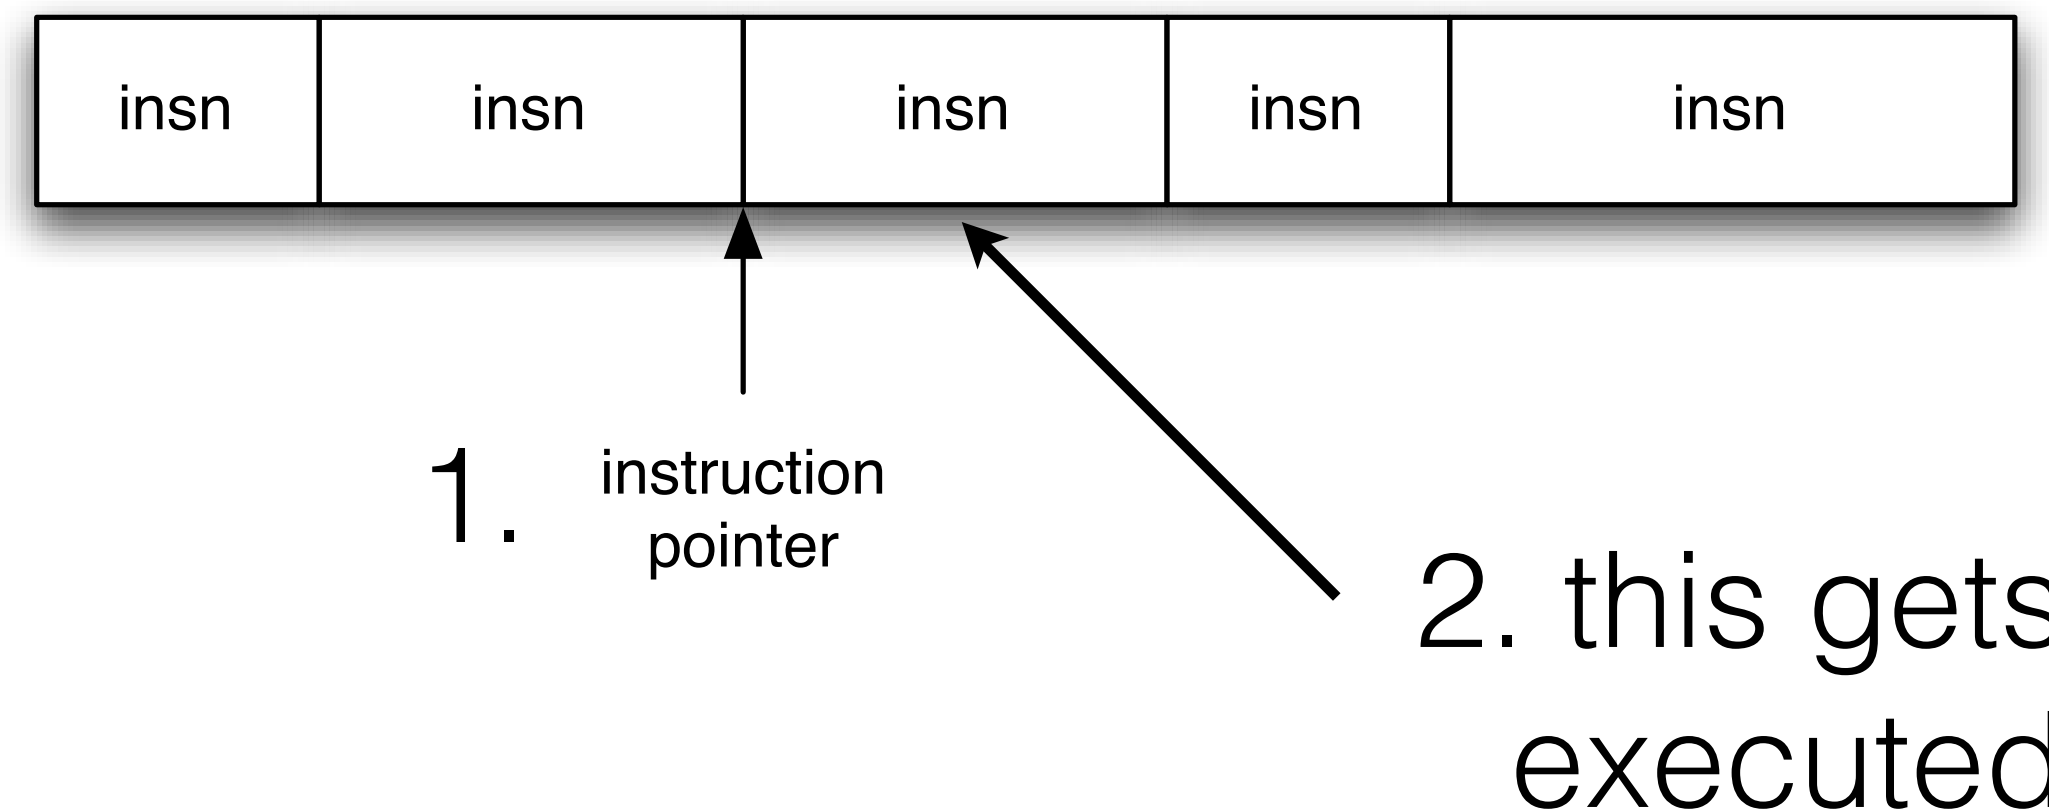
ordinary programming:



procedure:

1. IP points at instruction
2. instruction executed
3. IP incremented by length of instruction

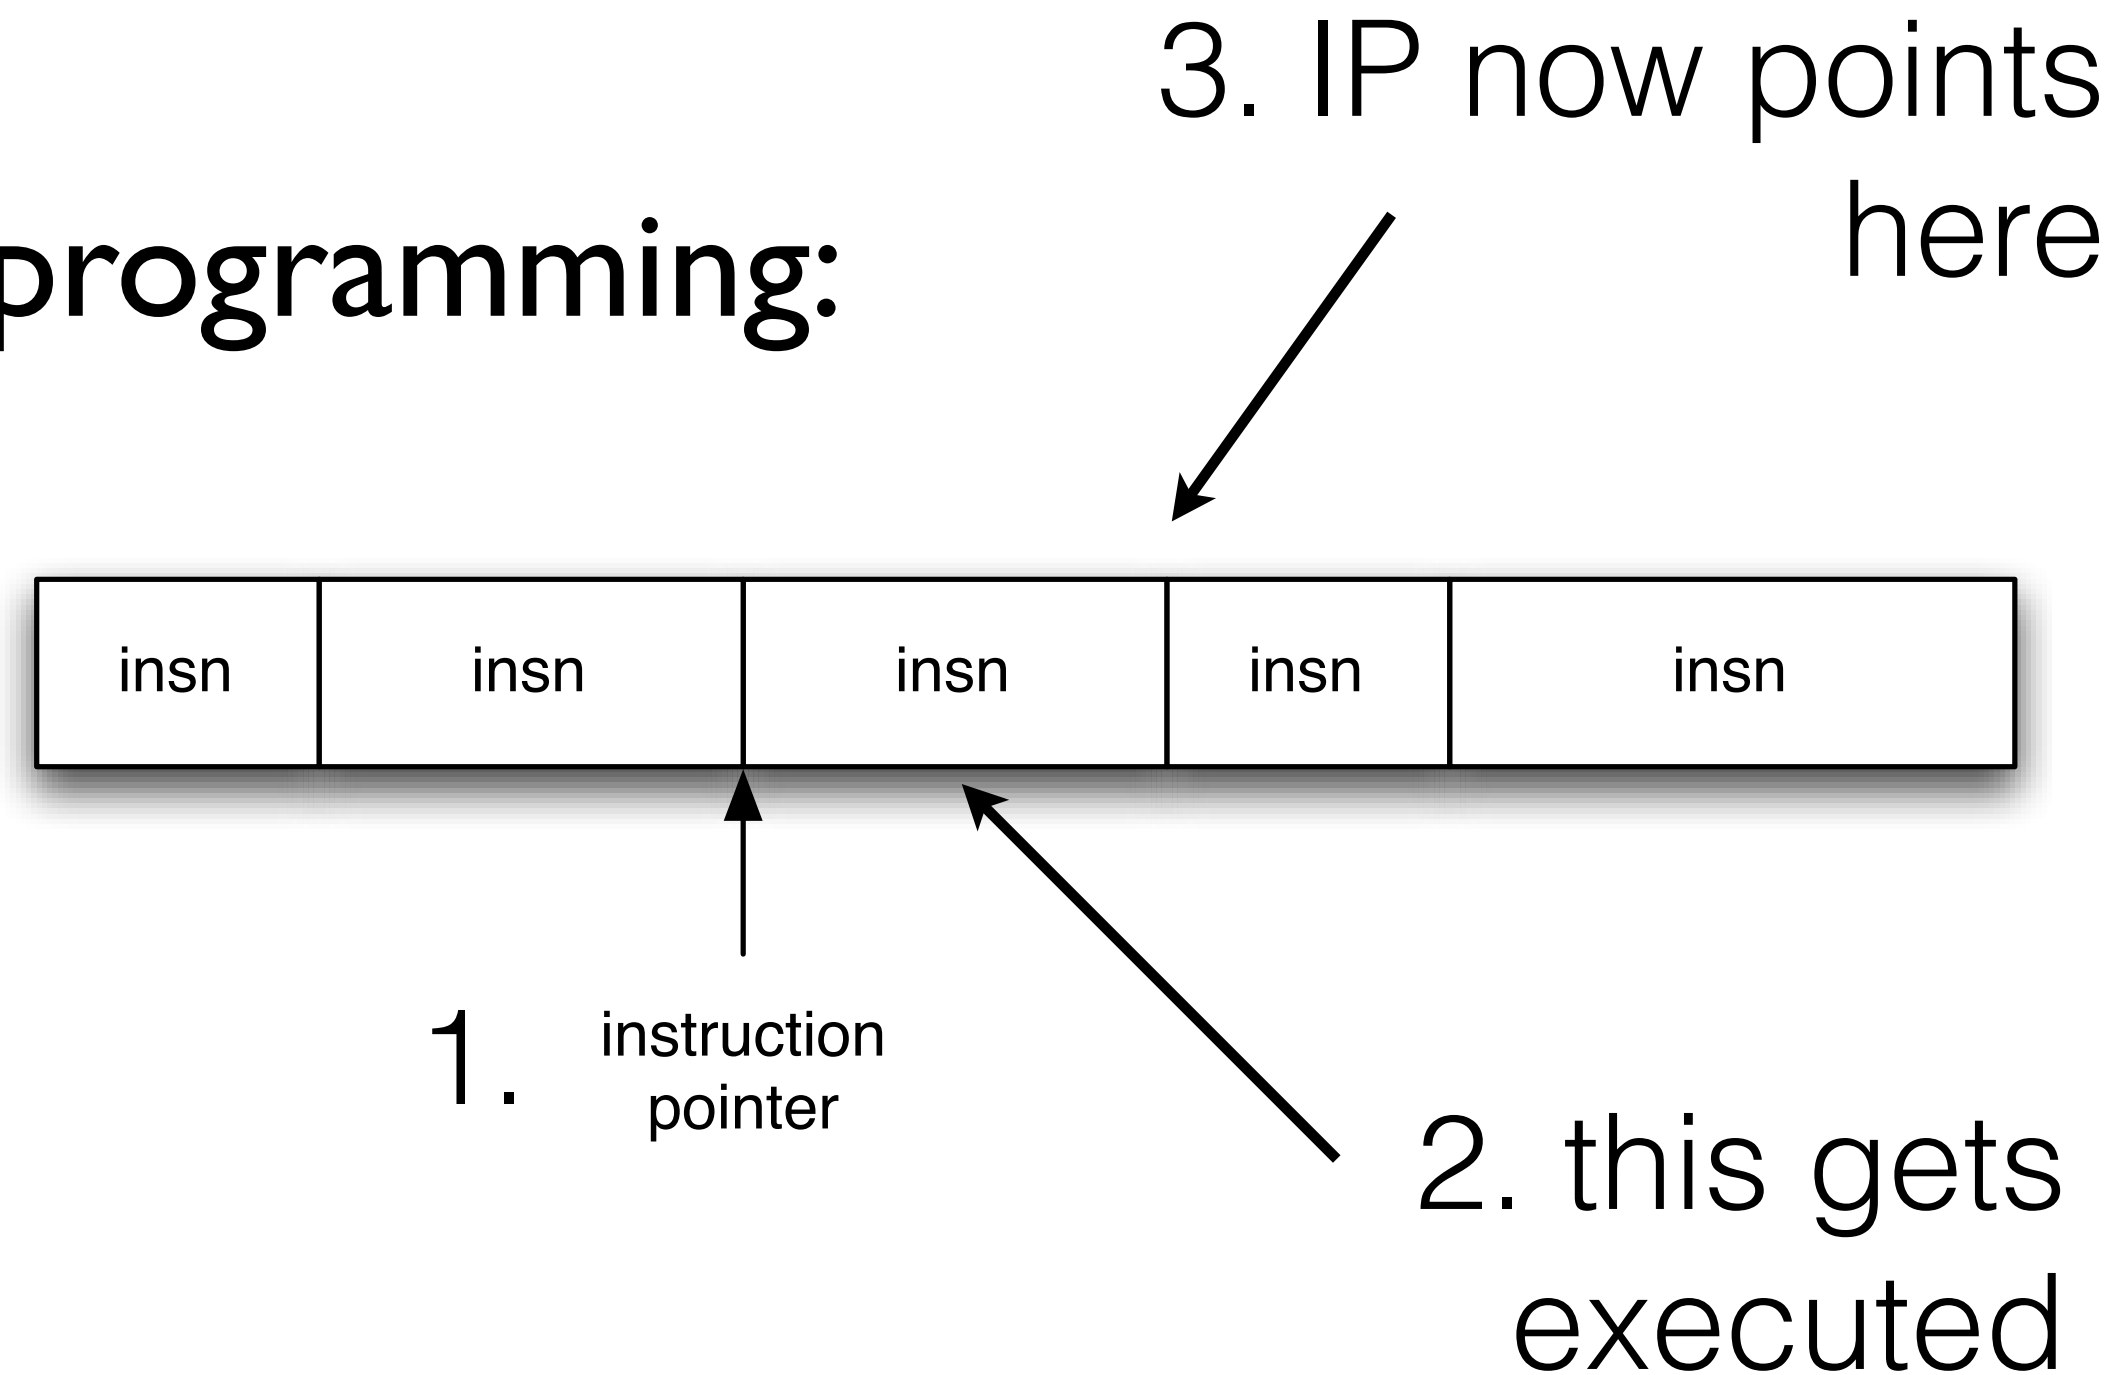
ordinary programming:



procedure:

1. IP points at instruction
2. instruction executed
3. IP incremented by length of instruction

ordinary programming:



procedure:

1. IP points at instruction
2. instruction executed
3. IP incremented by length of instruction

return-oriented programming (ROP)

setup stack to point to instructions that end in return

return-oriented programming (ROP)

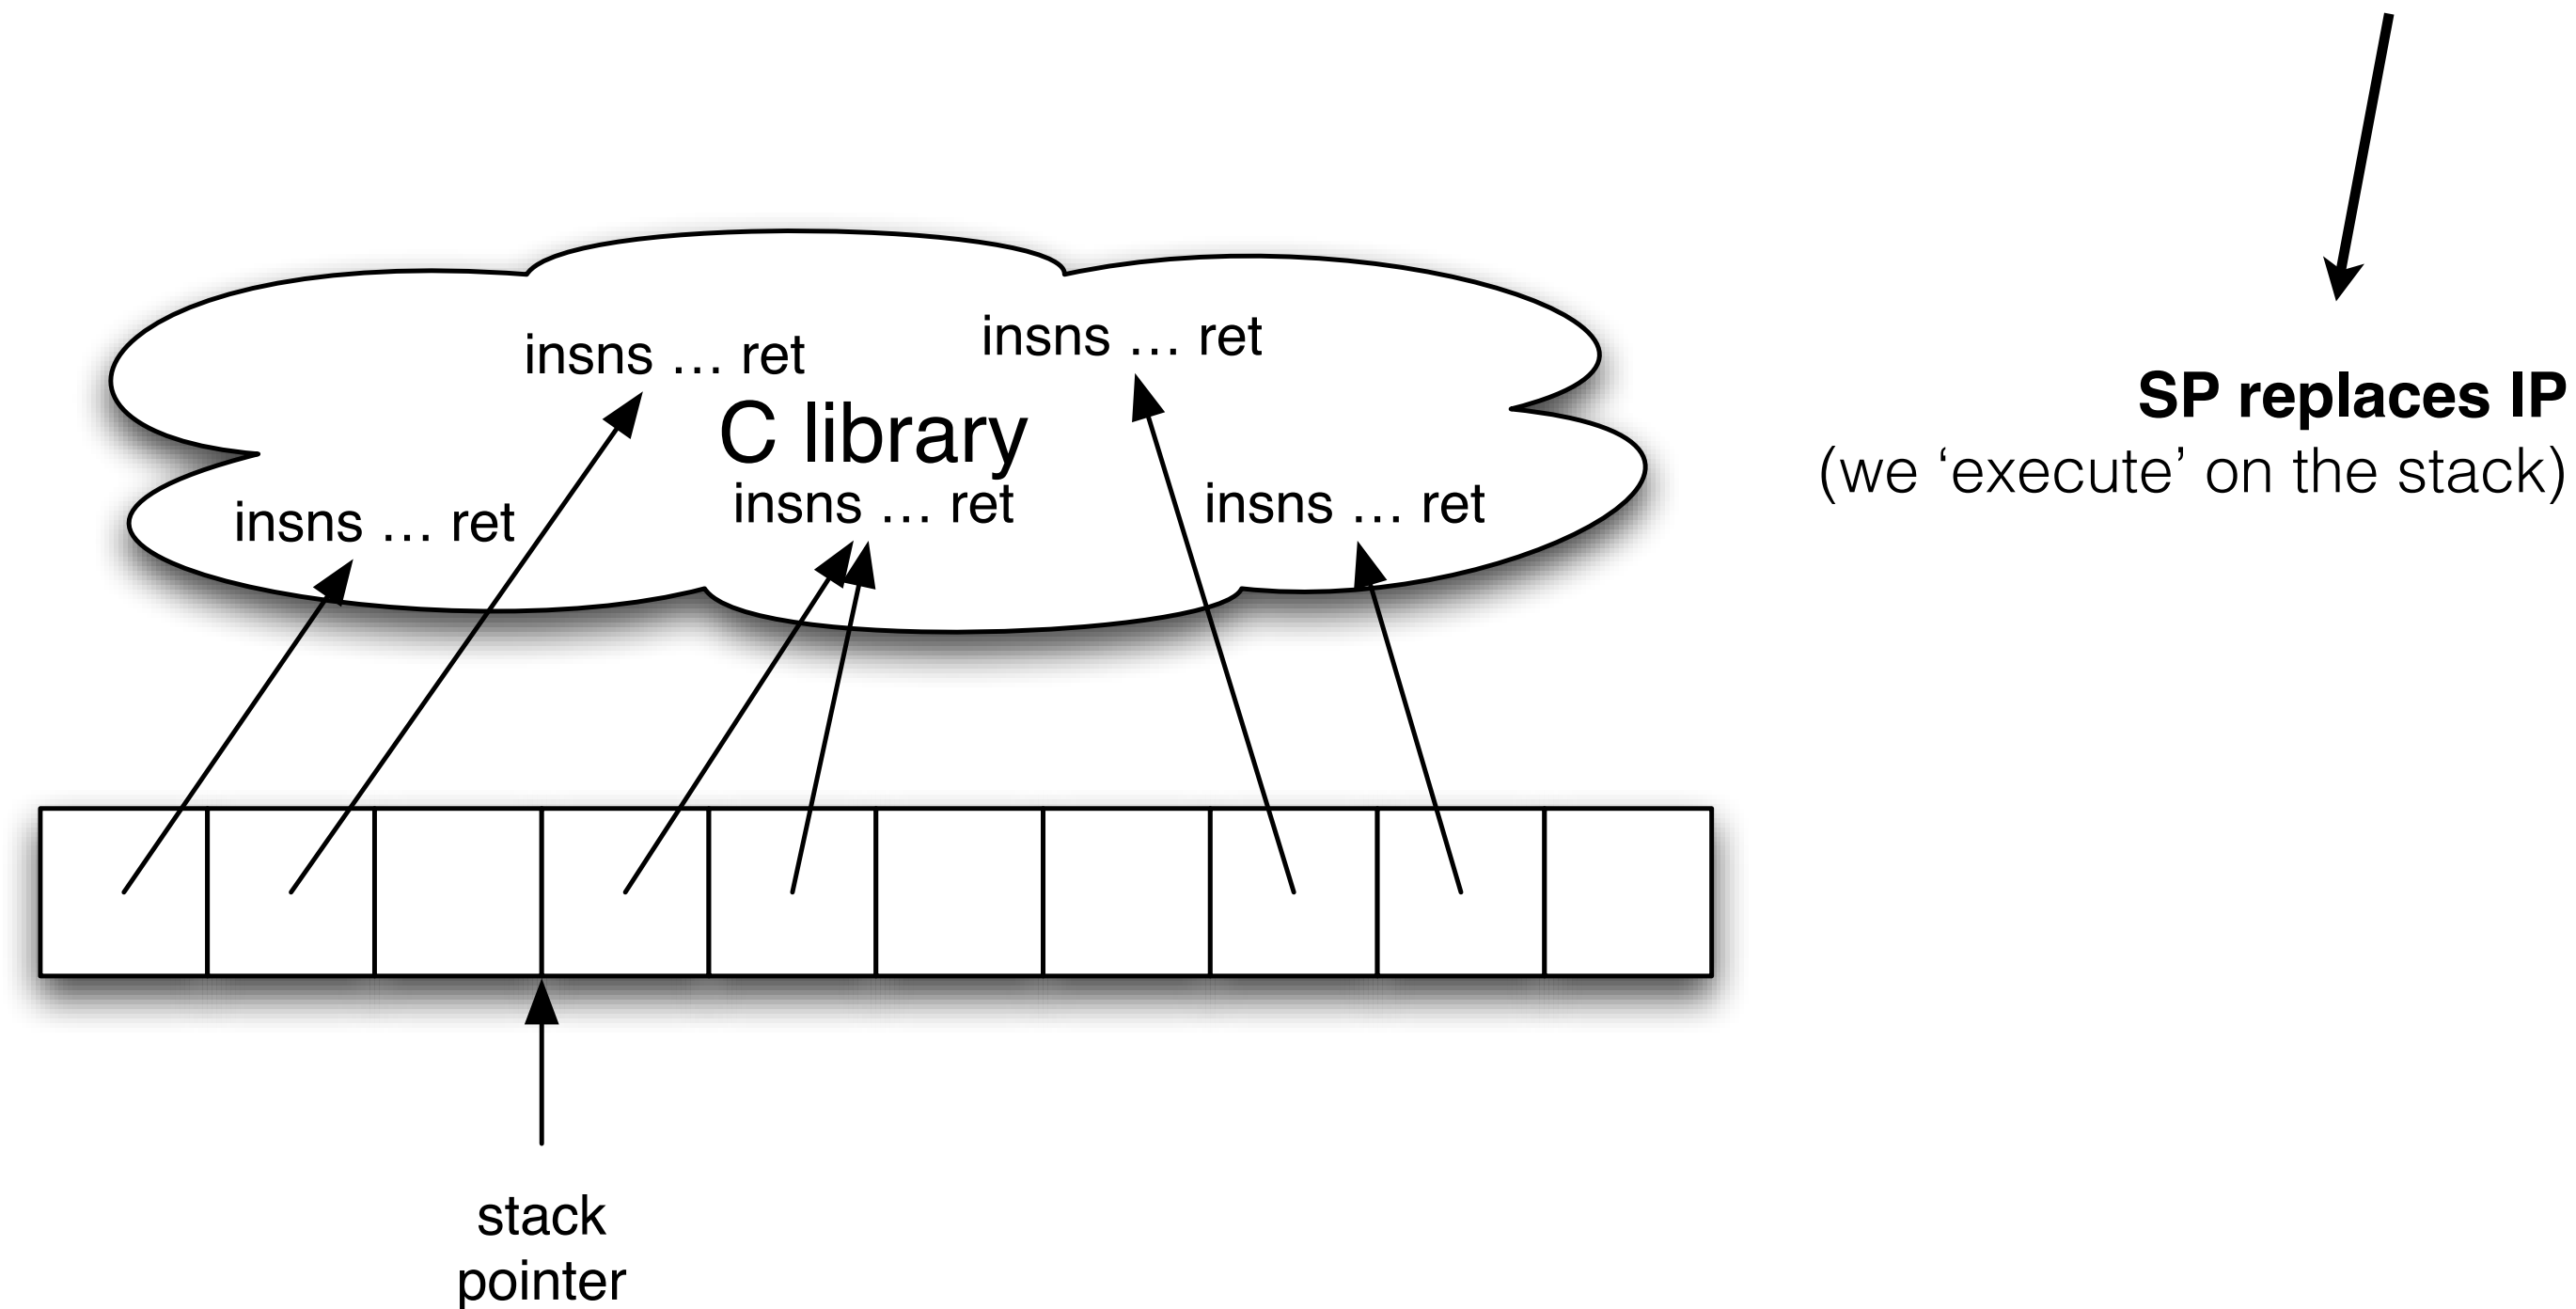
setup stack to point to instructions that end in return



SP replaces IP
(we 'execute' on the stack)

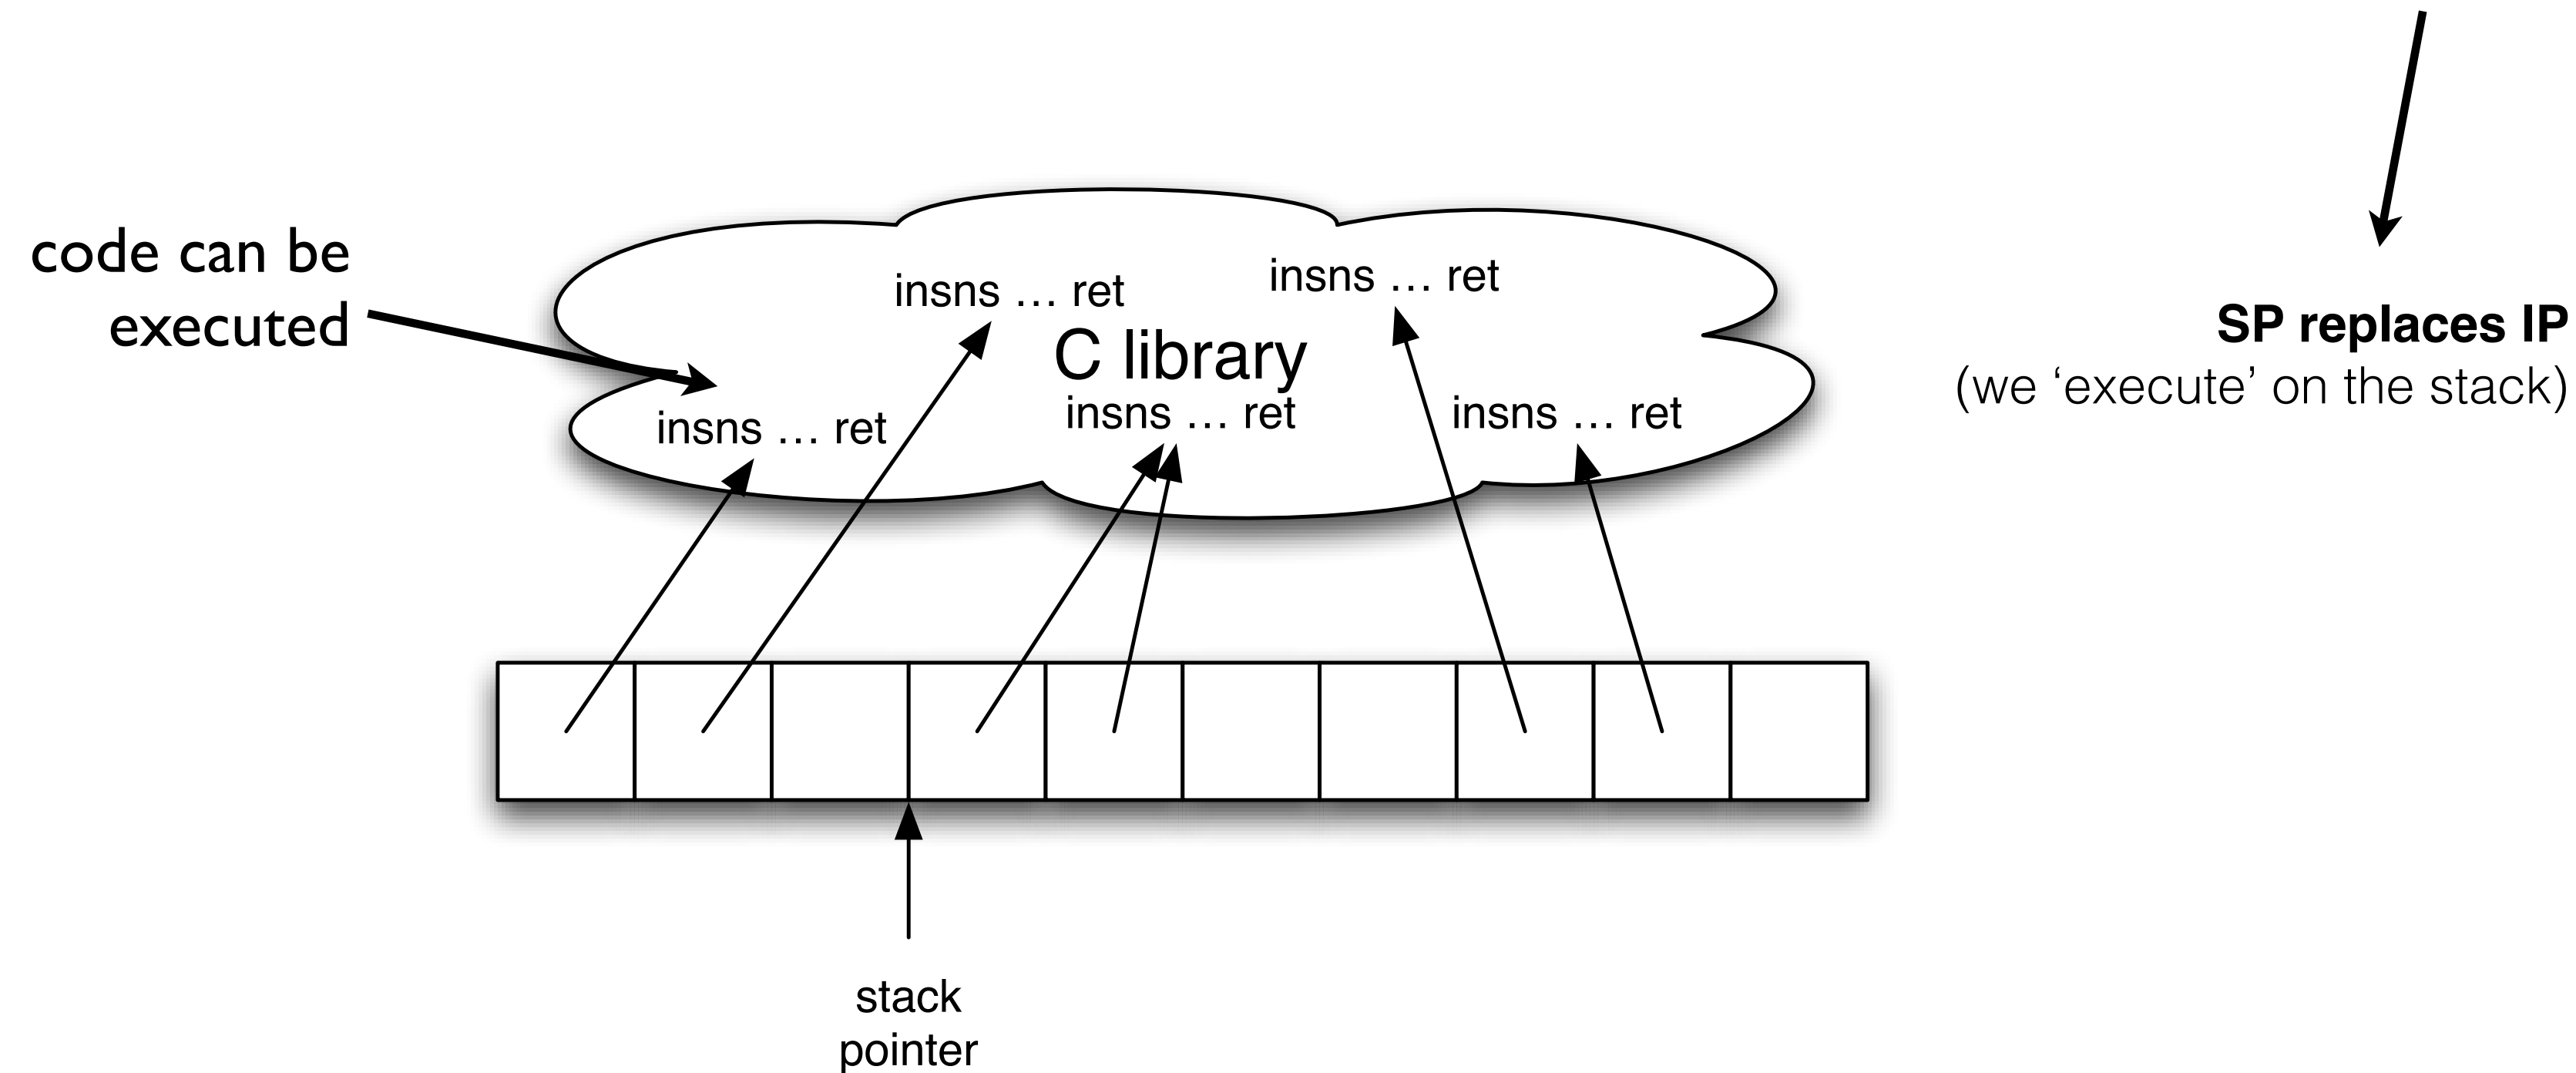
return-oriented programming (ROP)

setup stack to point to instructions that end in return



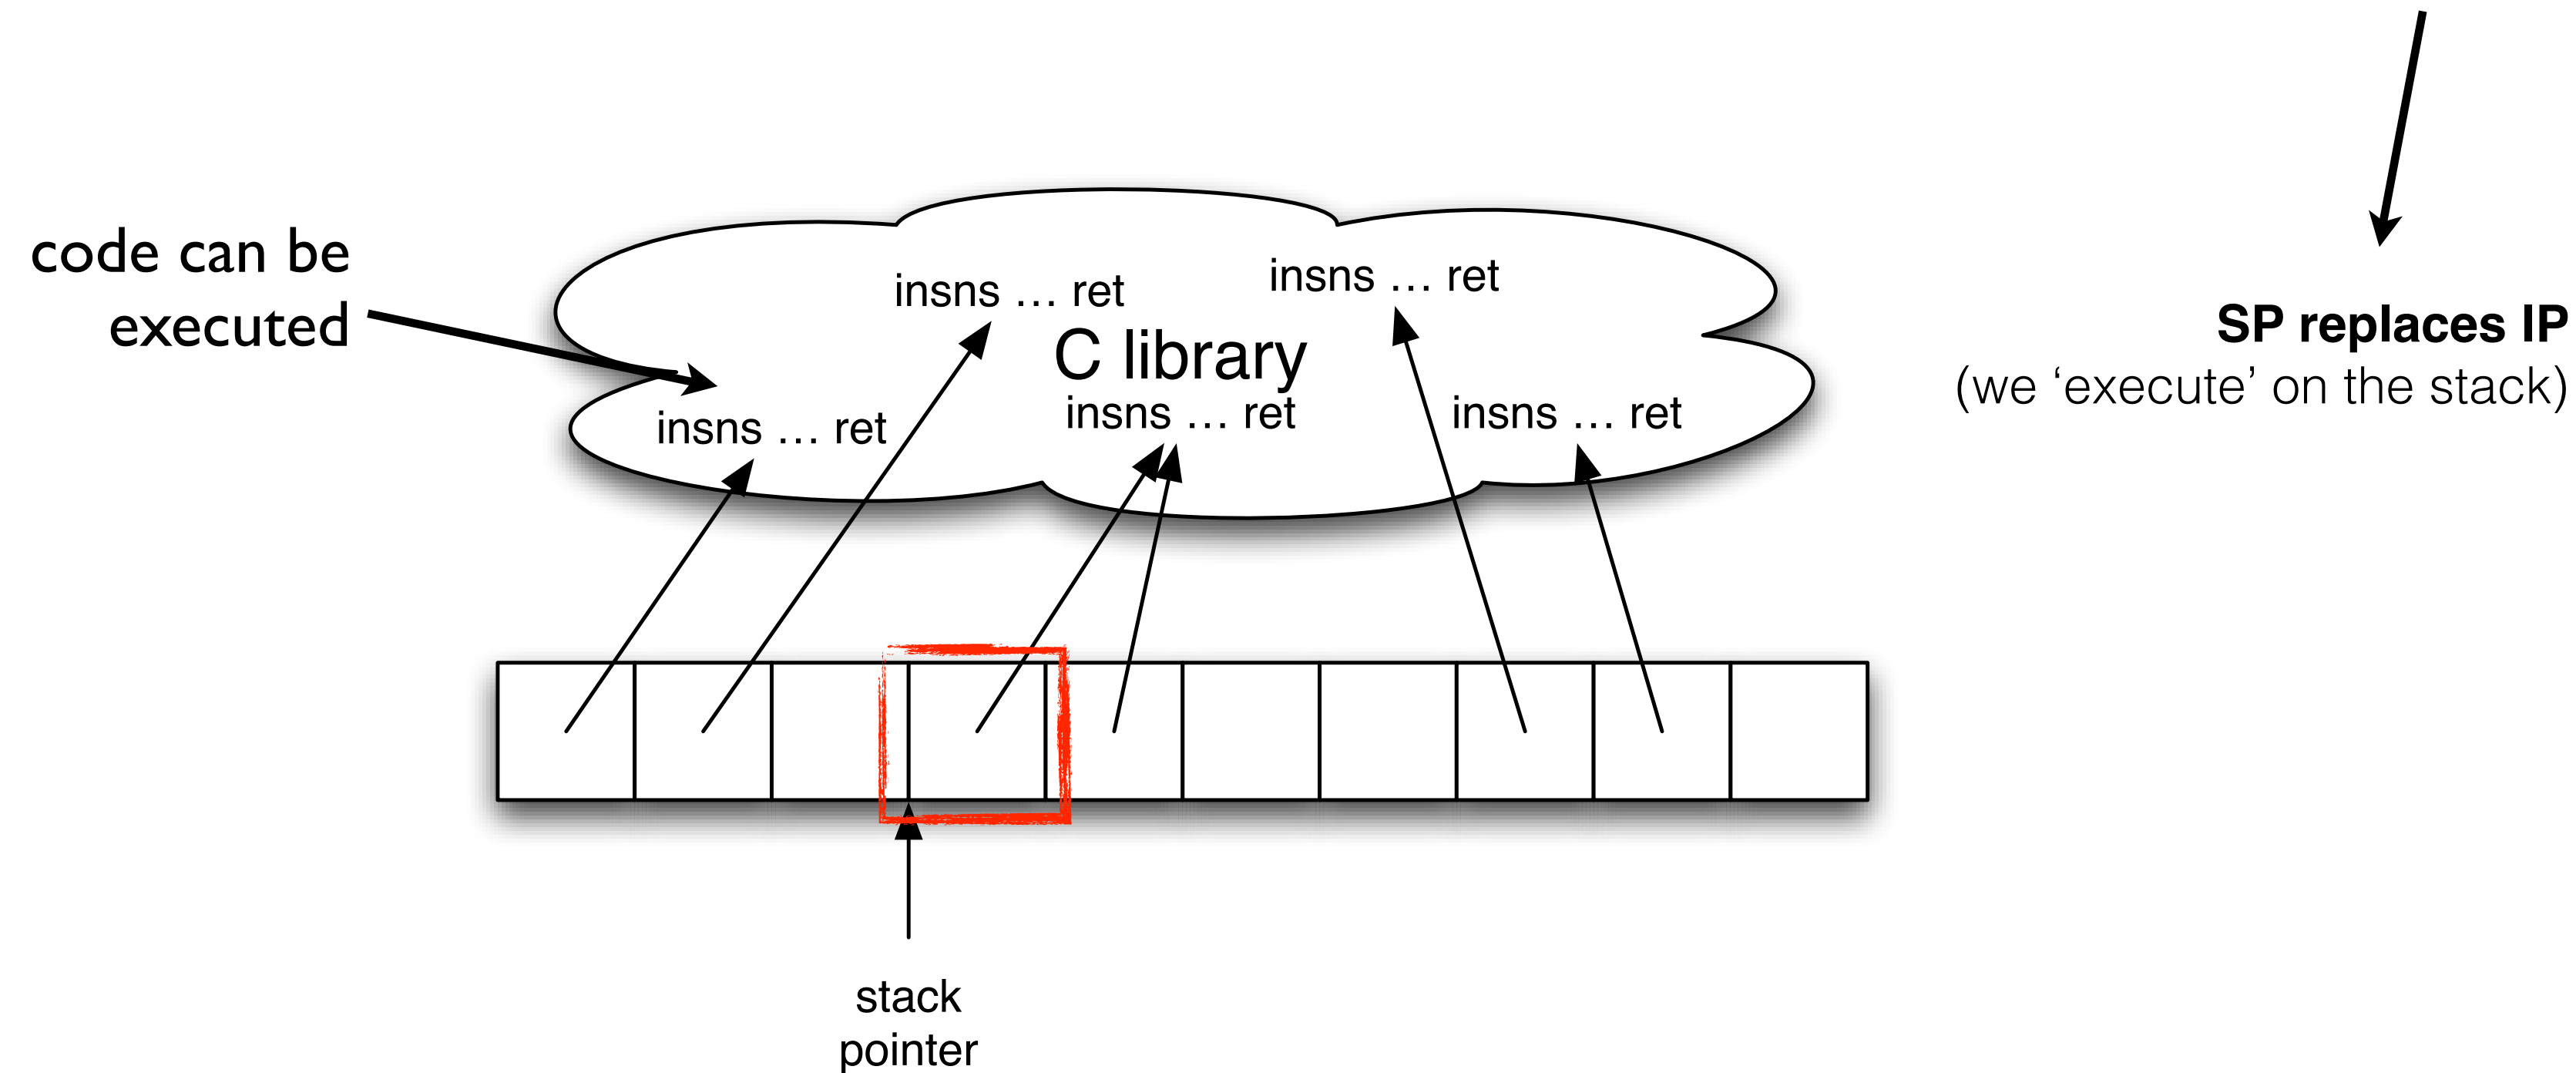
return-oriented programming (ROP)

setup stack to point to instructions that end in return



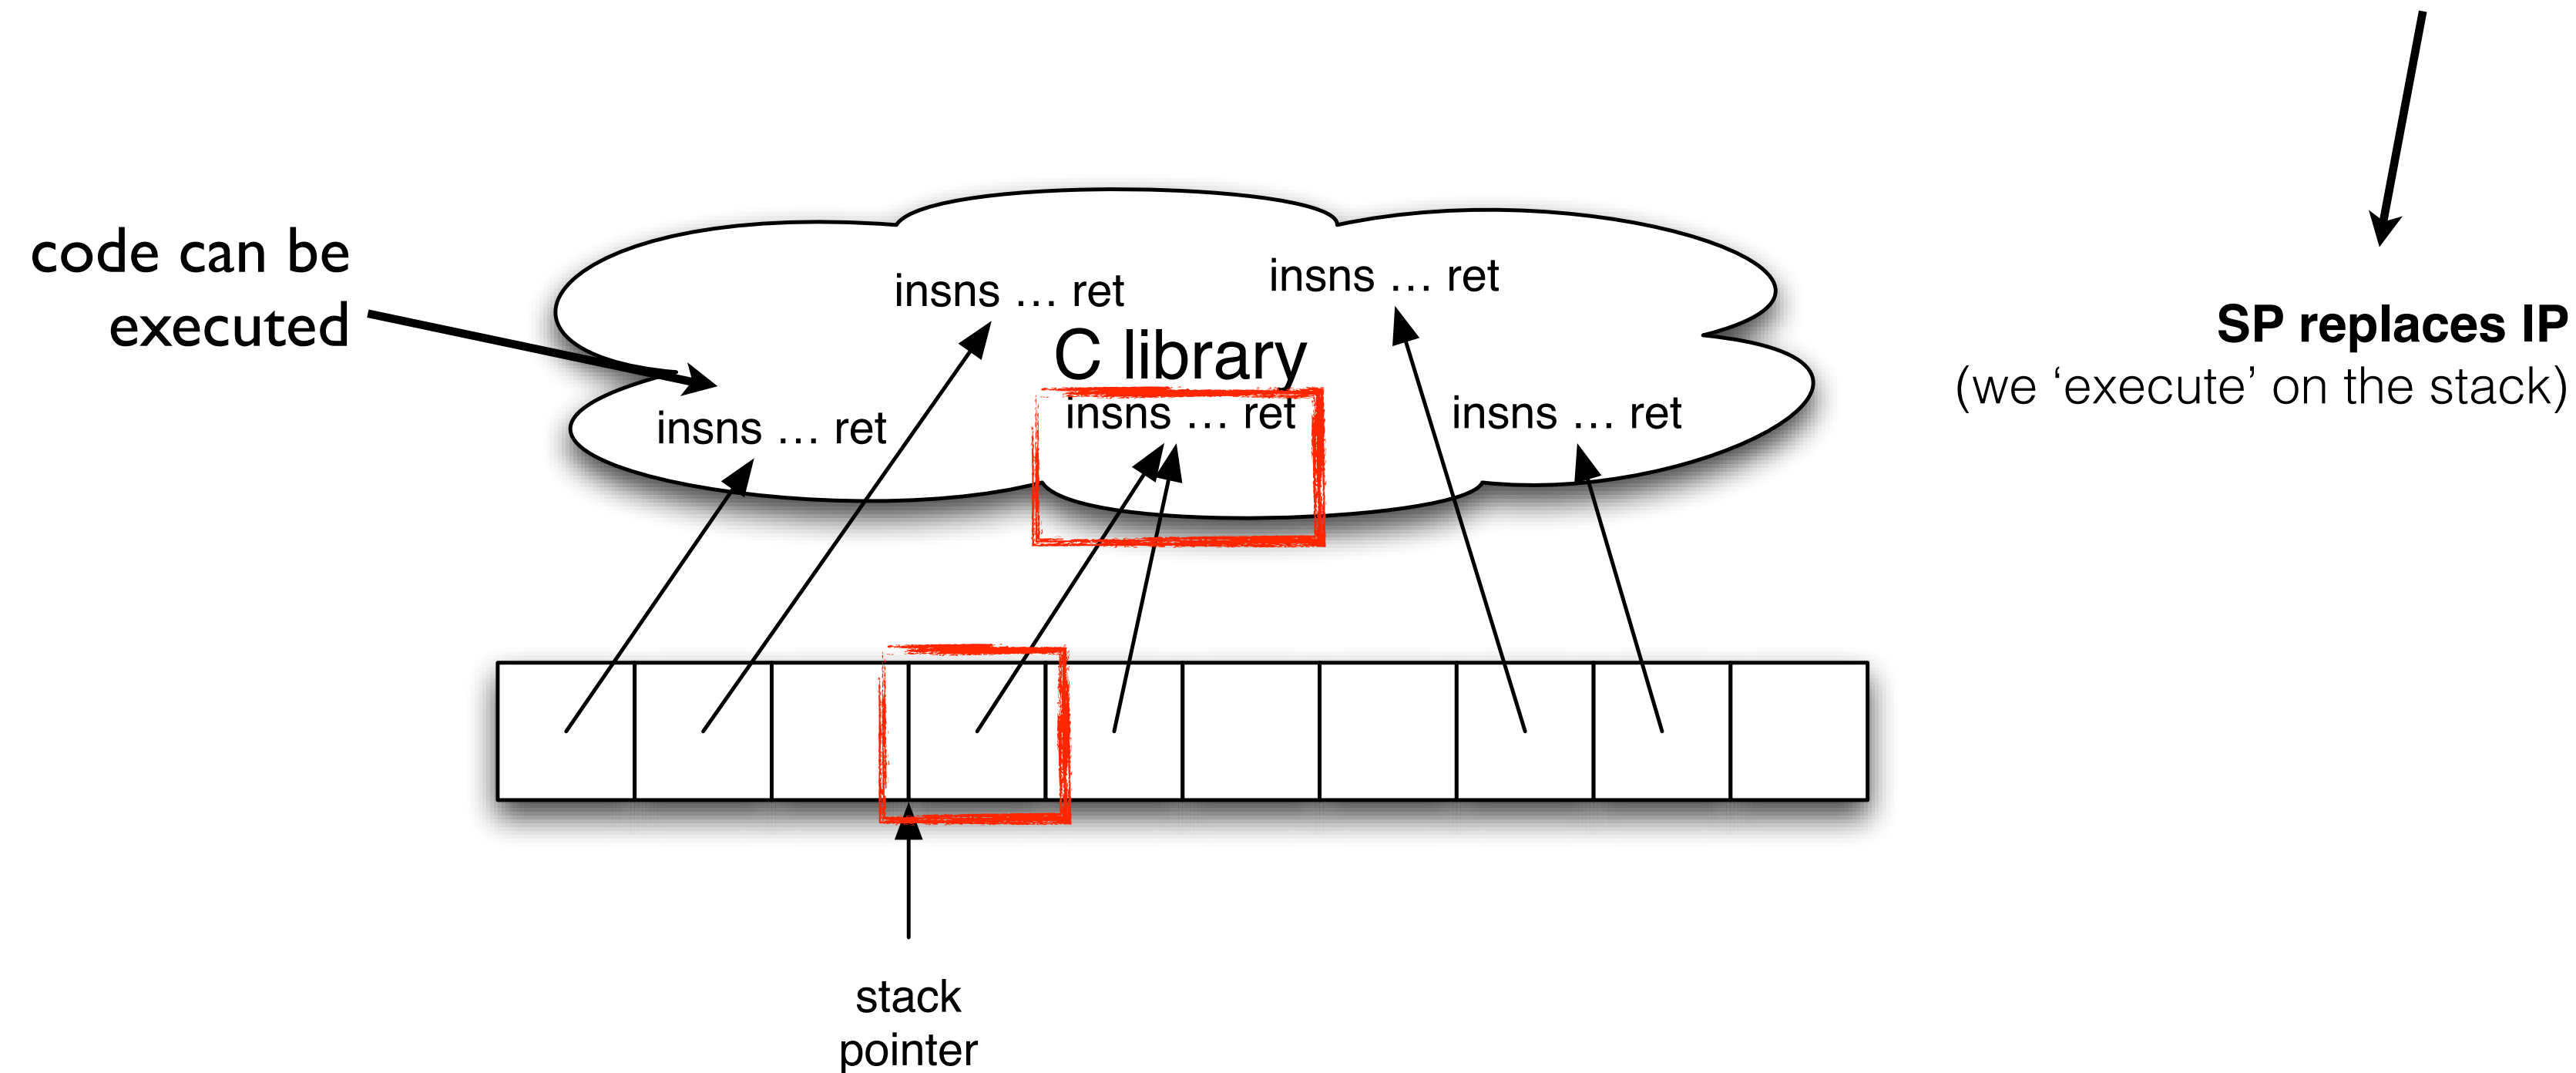
return-oriented programming (ROP)

setup stack to point to instructions that end in return



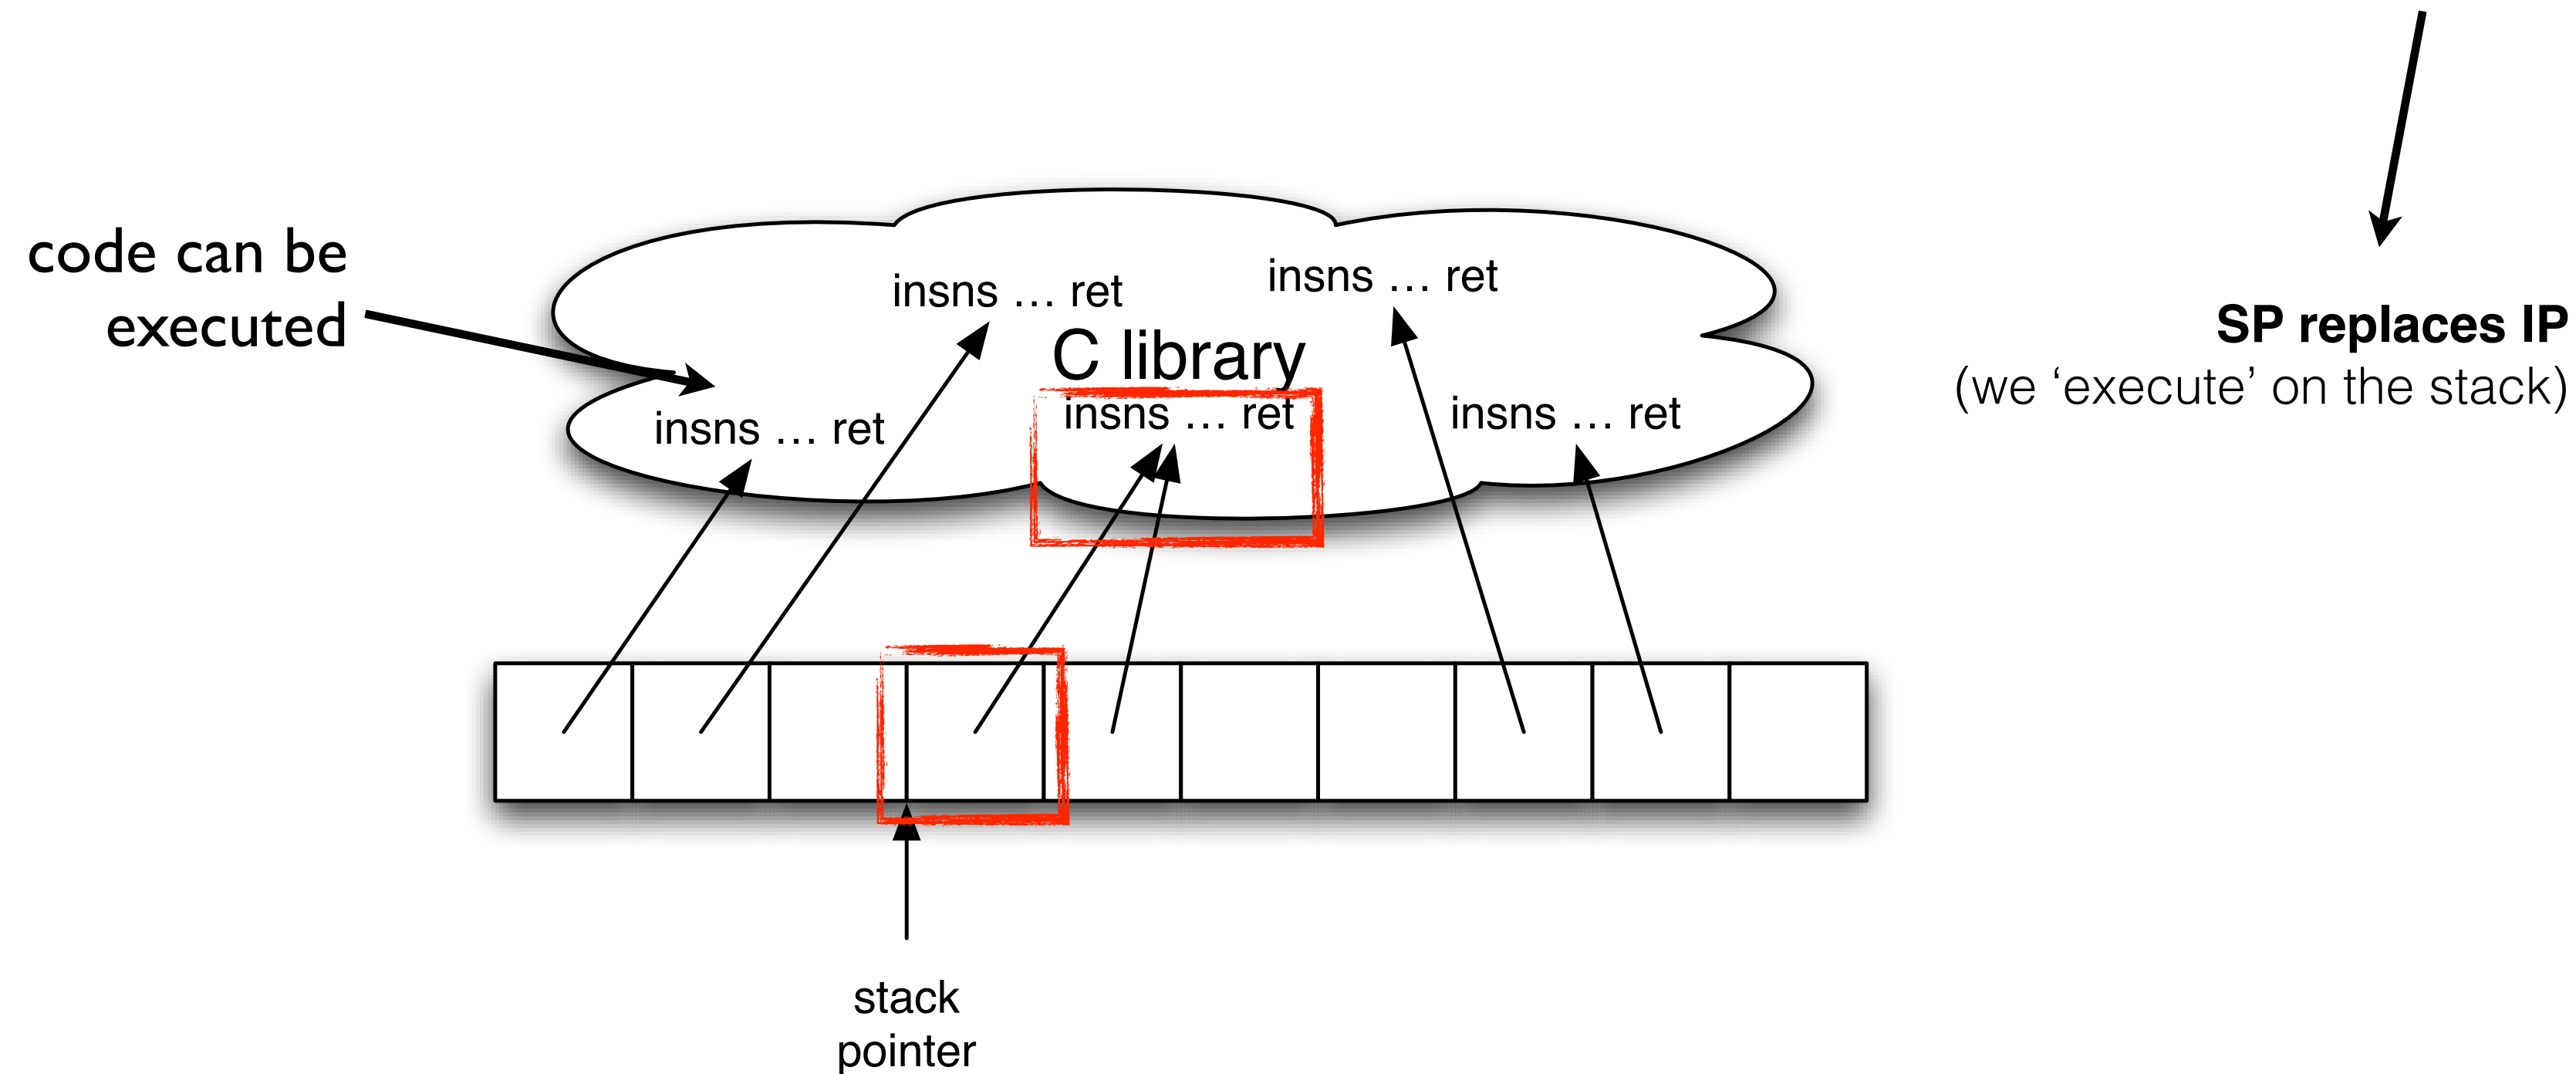
return-oriented programming (ROP)

setup stack to point to instructions that end in return



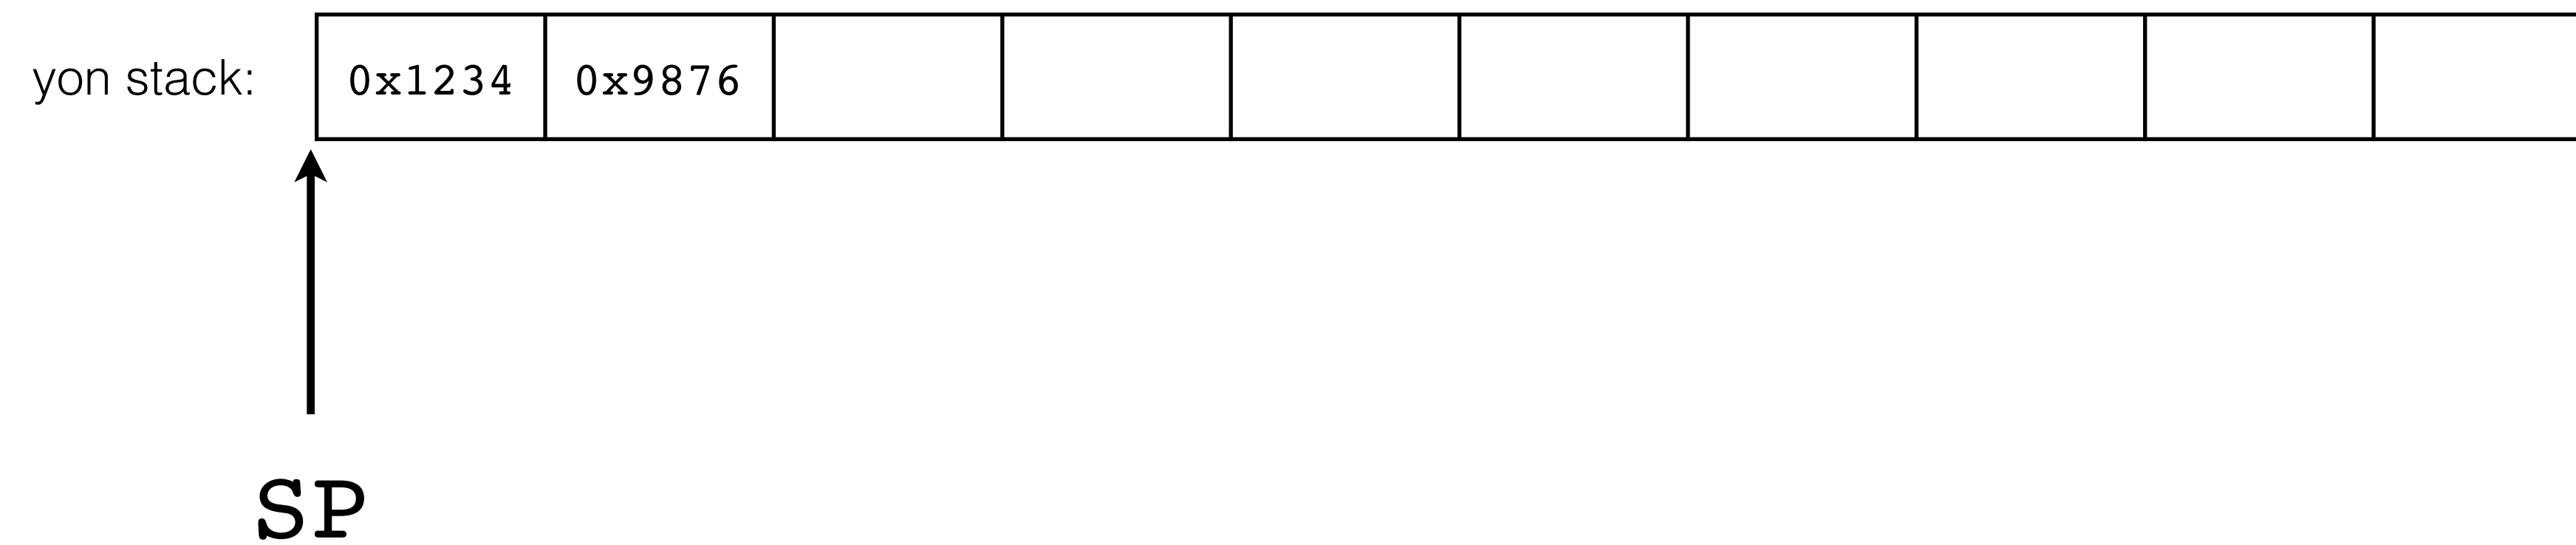
return-oriented programming (ROP)

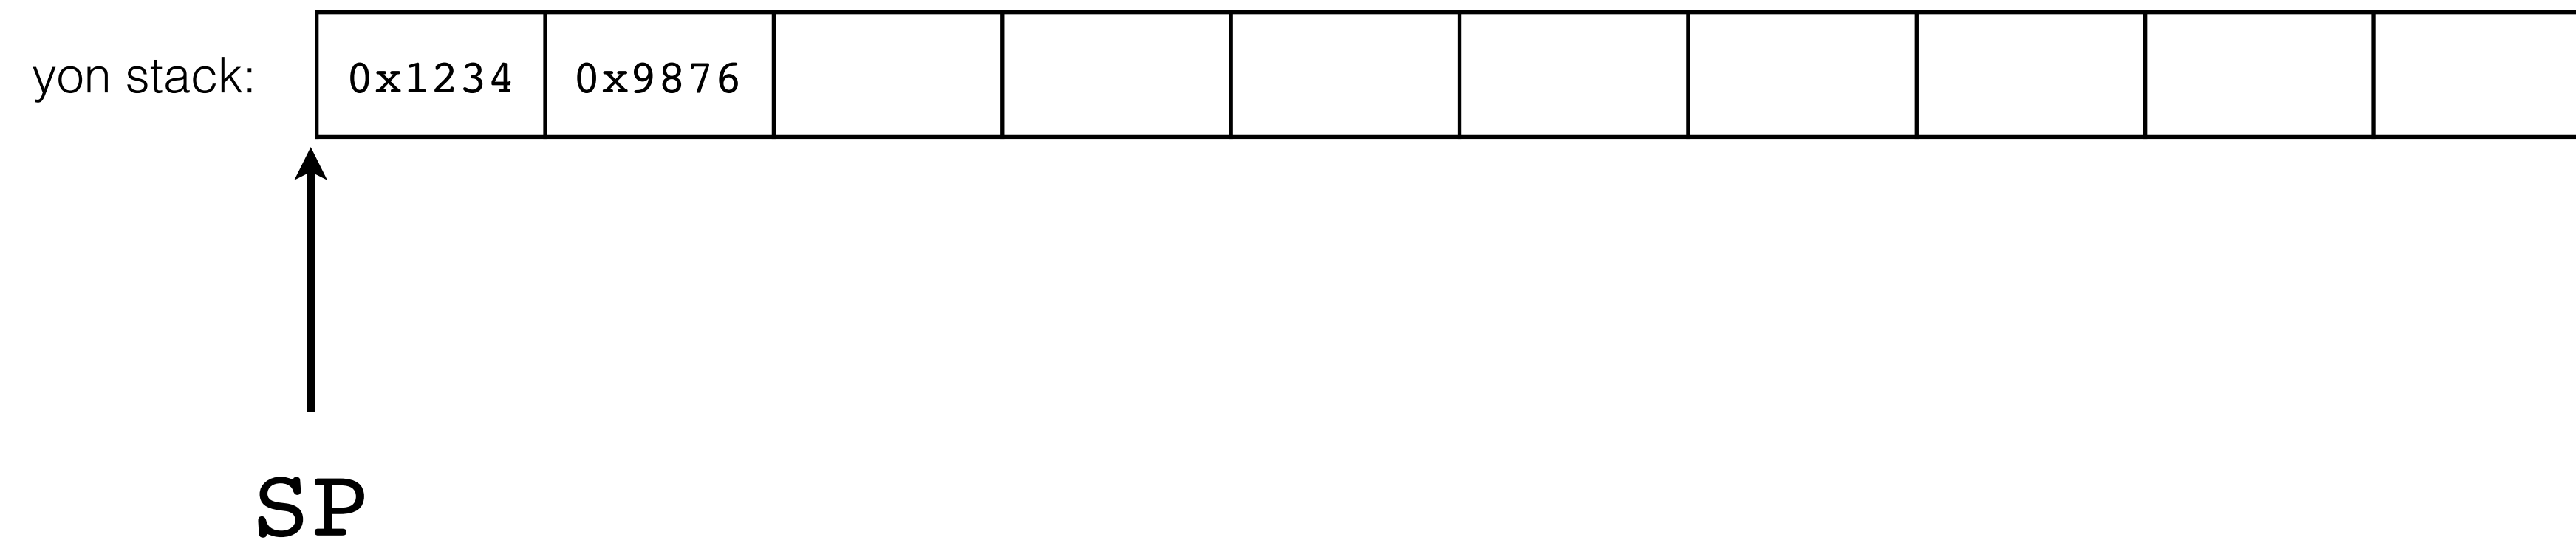
setup stack to point to instructions that end in return



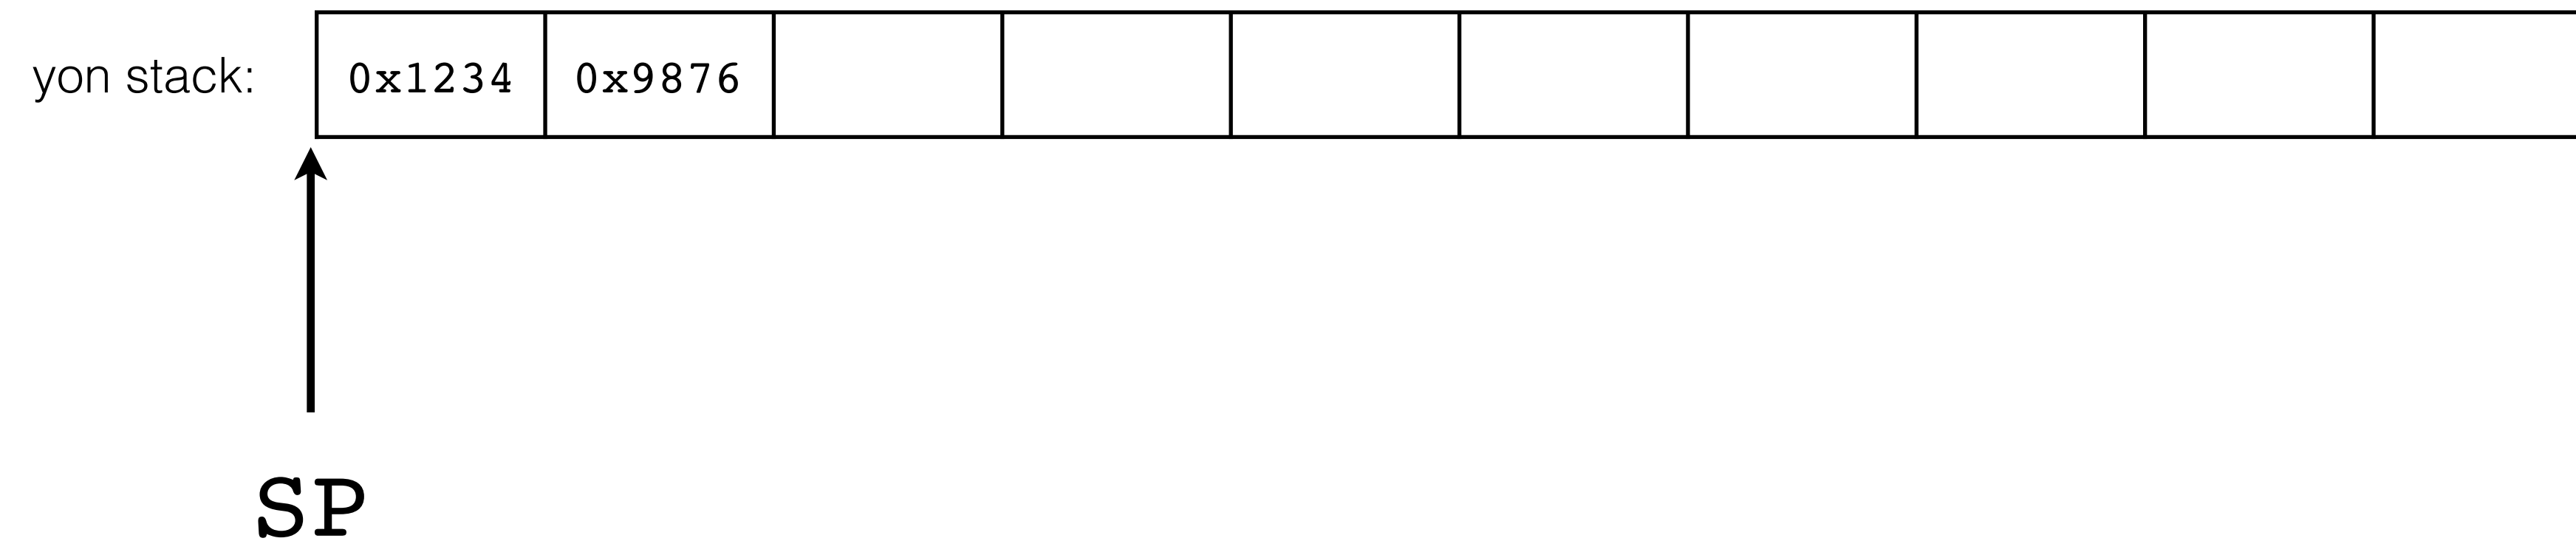
procedure:

1. setup stack via buffer overflow
2. issue return
3. pop addr from stack
4. IP executes instruction at addr





1. exploited function issues `ret`

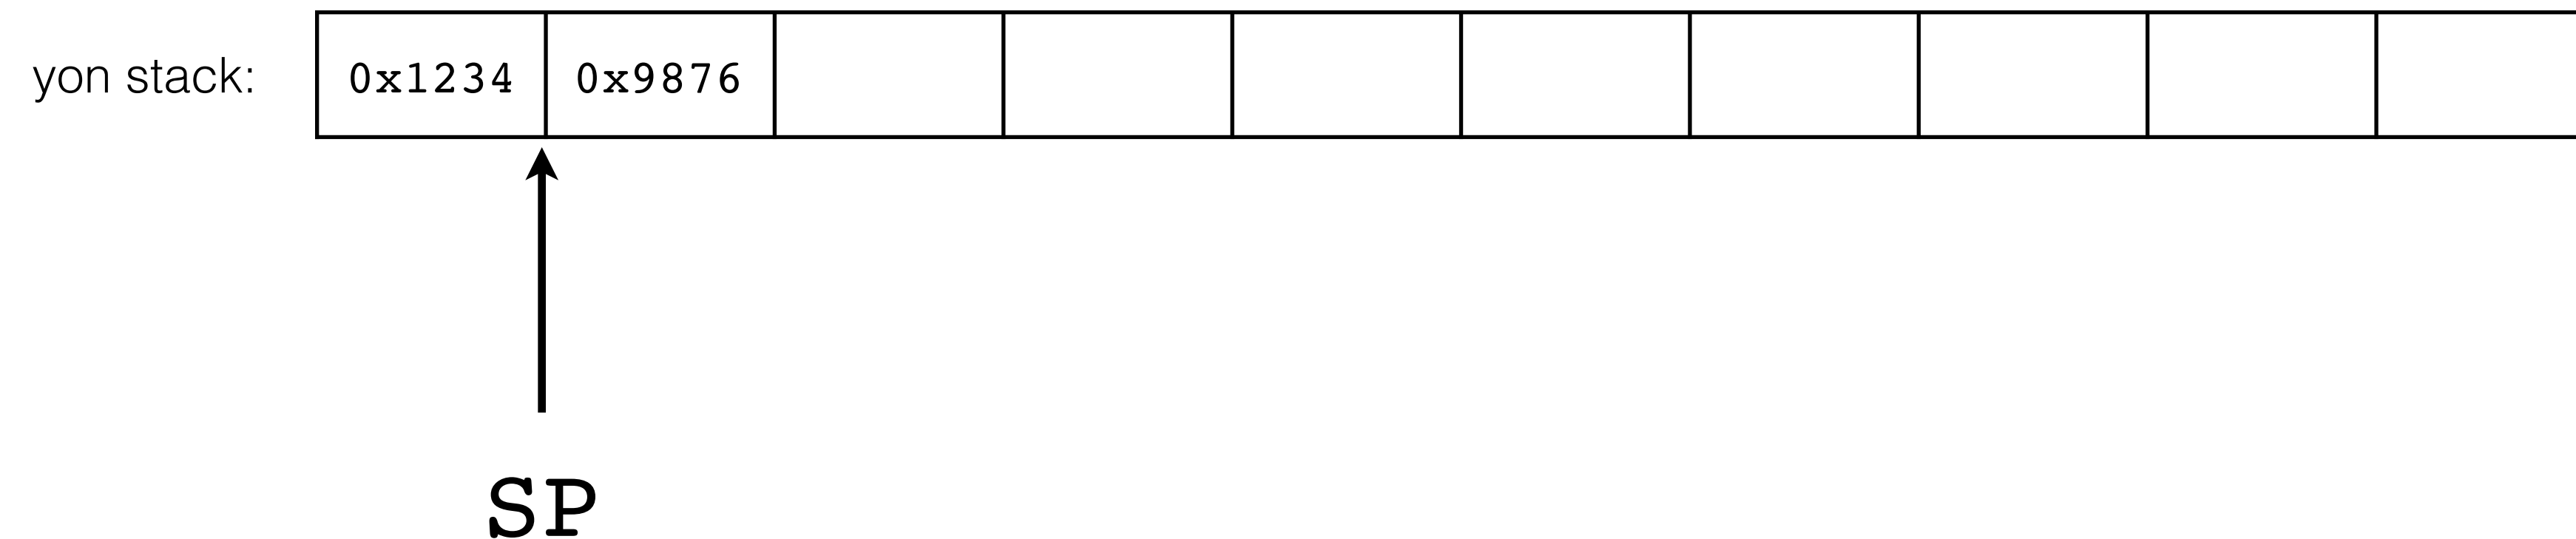


1. exploited function issues **ret**
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

yon stack:

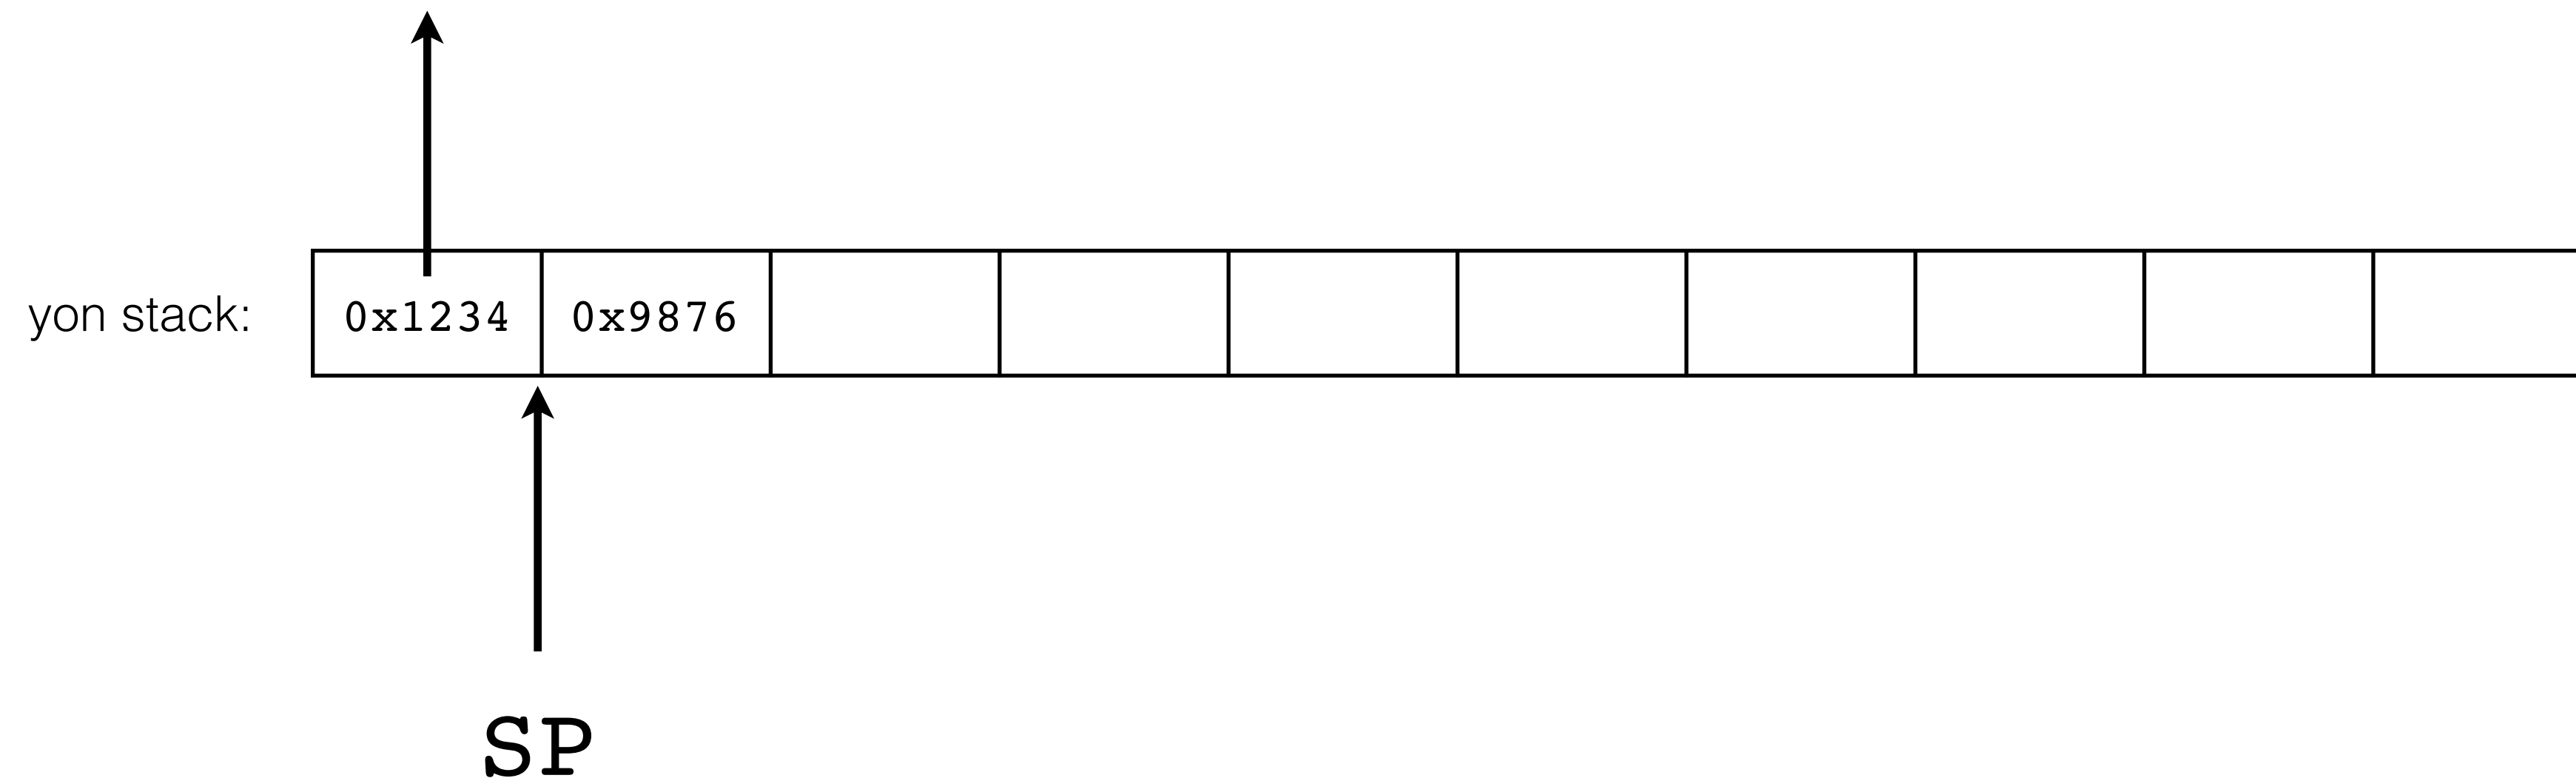
0x1234	0x9876								
--------	--------	--	--	--	--	--	--	--	--

1. exploited function issues **ret**
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes



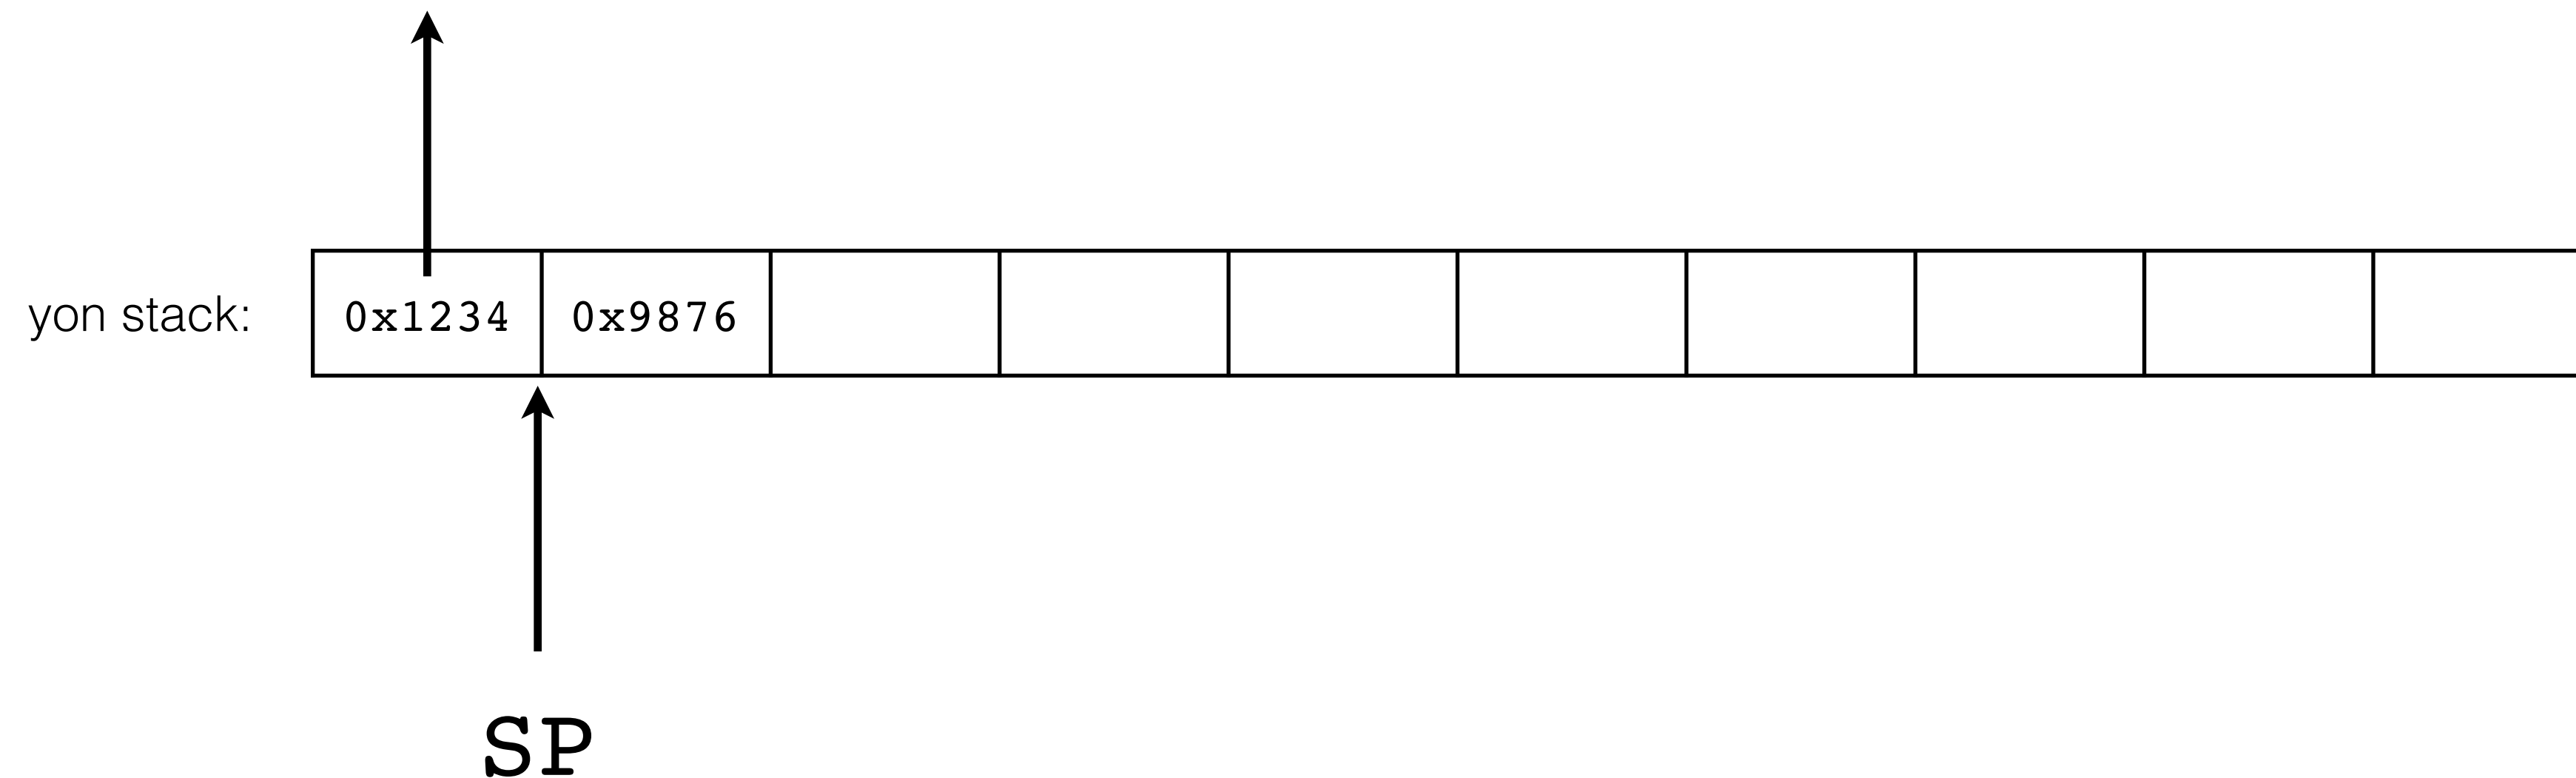
1. exploited function issues **ret**
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

IP → 0x1234 mov ...
0x1236 str ...
0x1238 ret ...



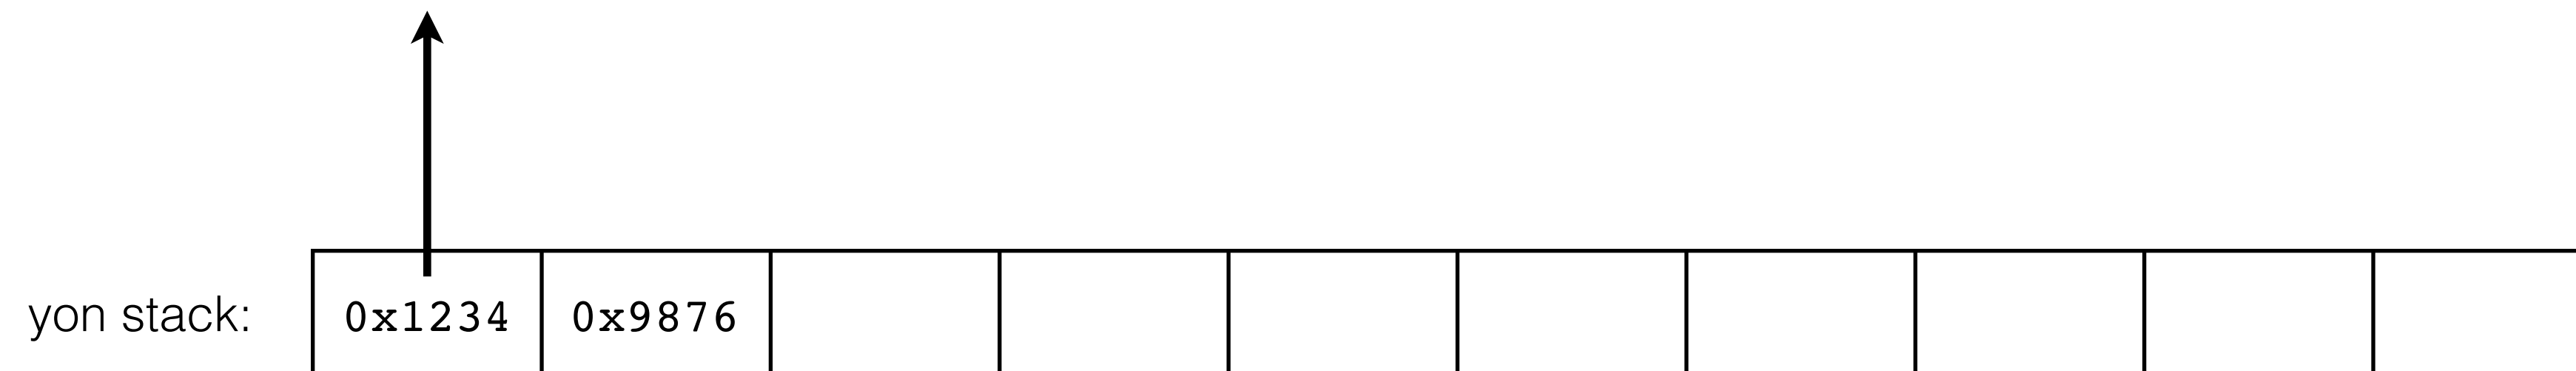
1. exploited function issues **ret**
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

IP → 0x1234 mov ... ← 3. a. execute instructions
0x1236 str ... b. issue `ret`
0x1238 ret ... c. follow step 2



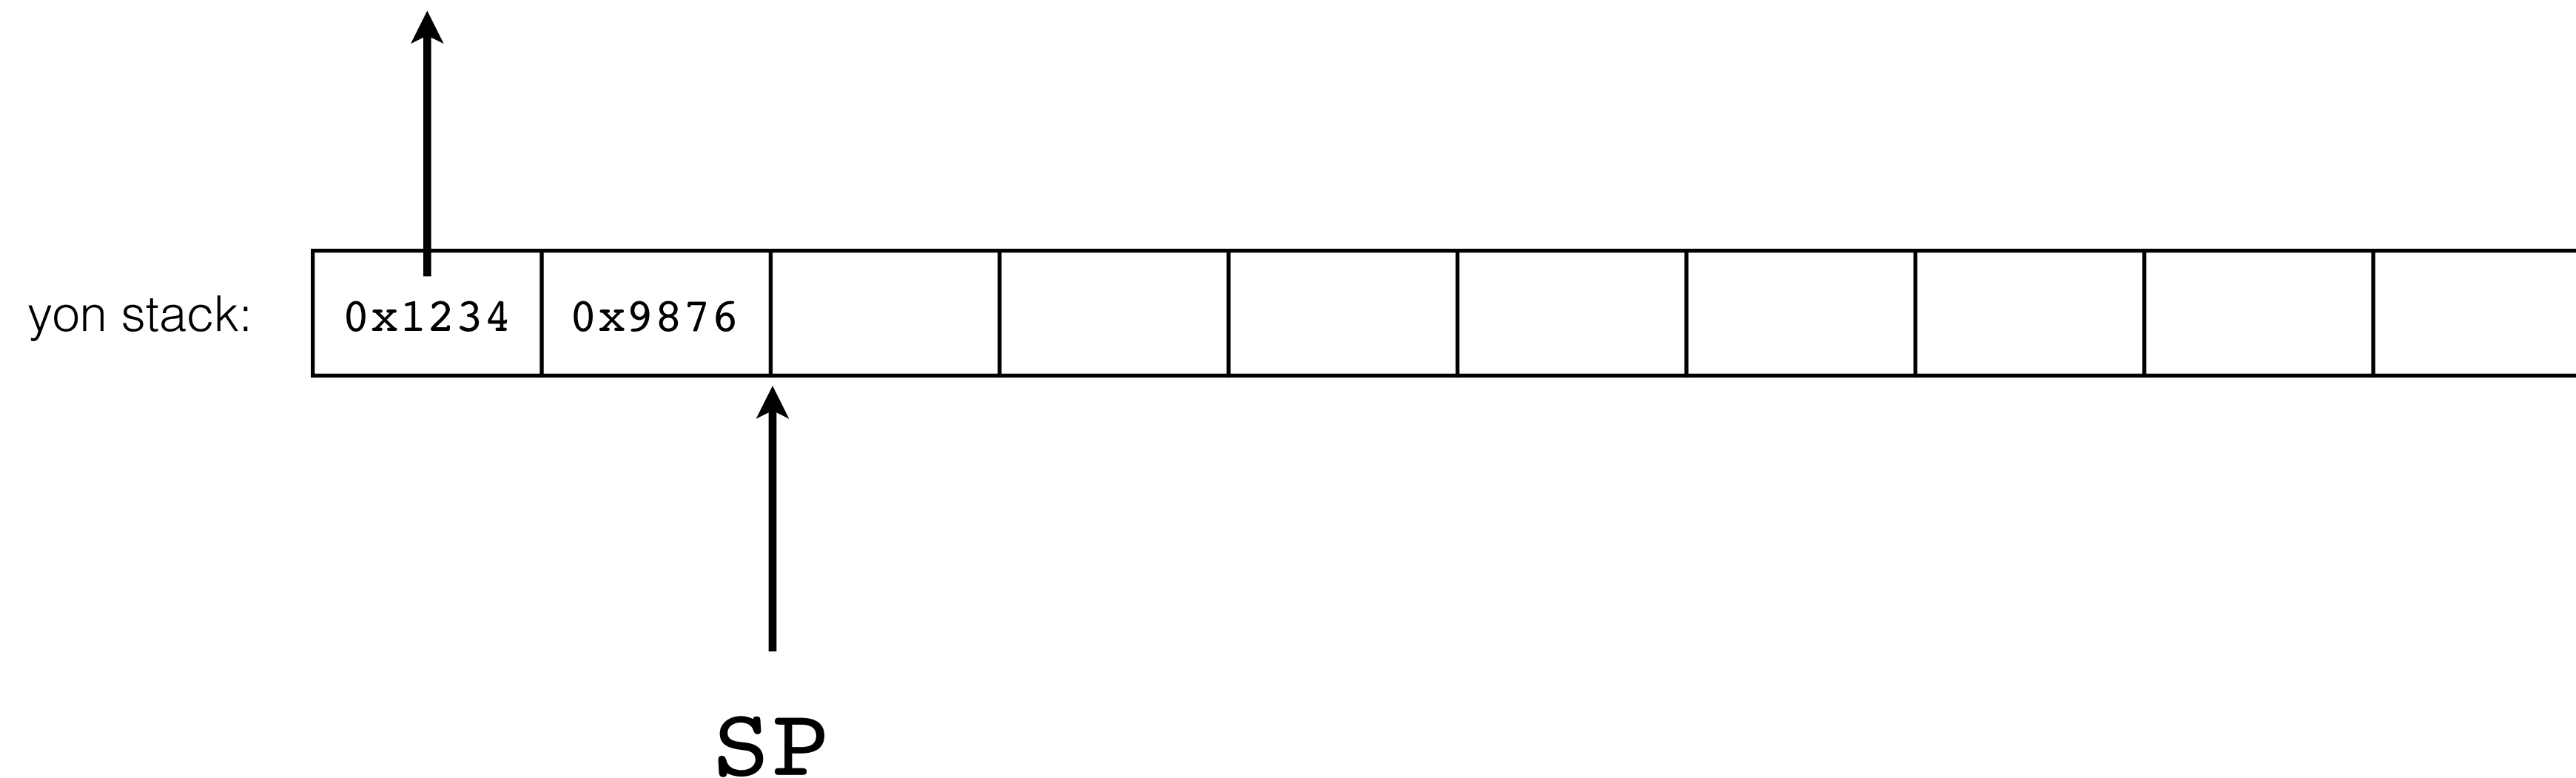
1. exploited function issues `ret`
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

IP → 0x1234 mov ... ← 3. a. execute instructions
0x1236 str ... b. issue `ret`
0x1238 ret ... c. follow step 2



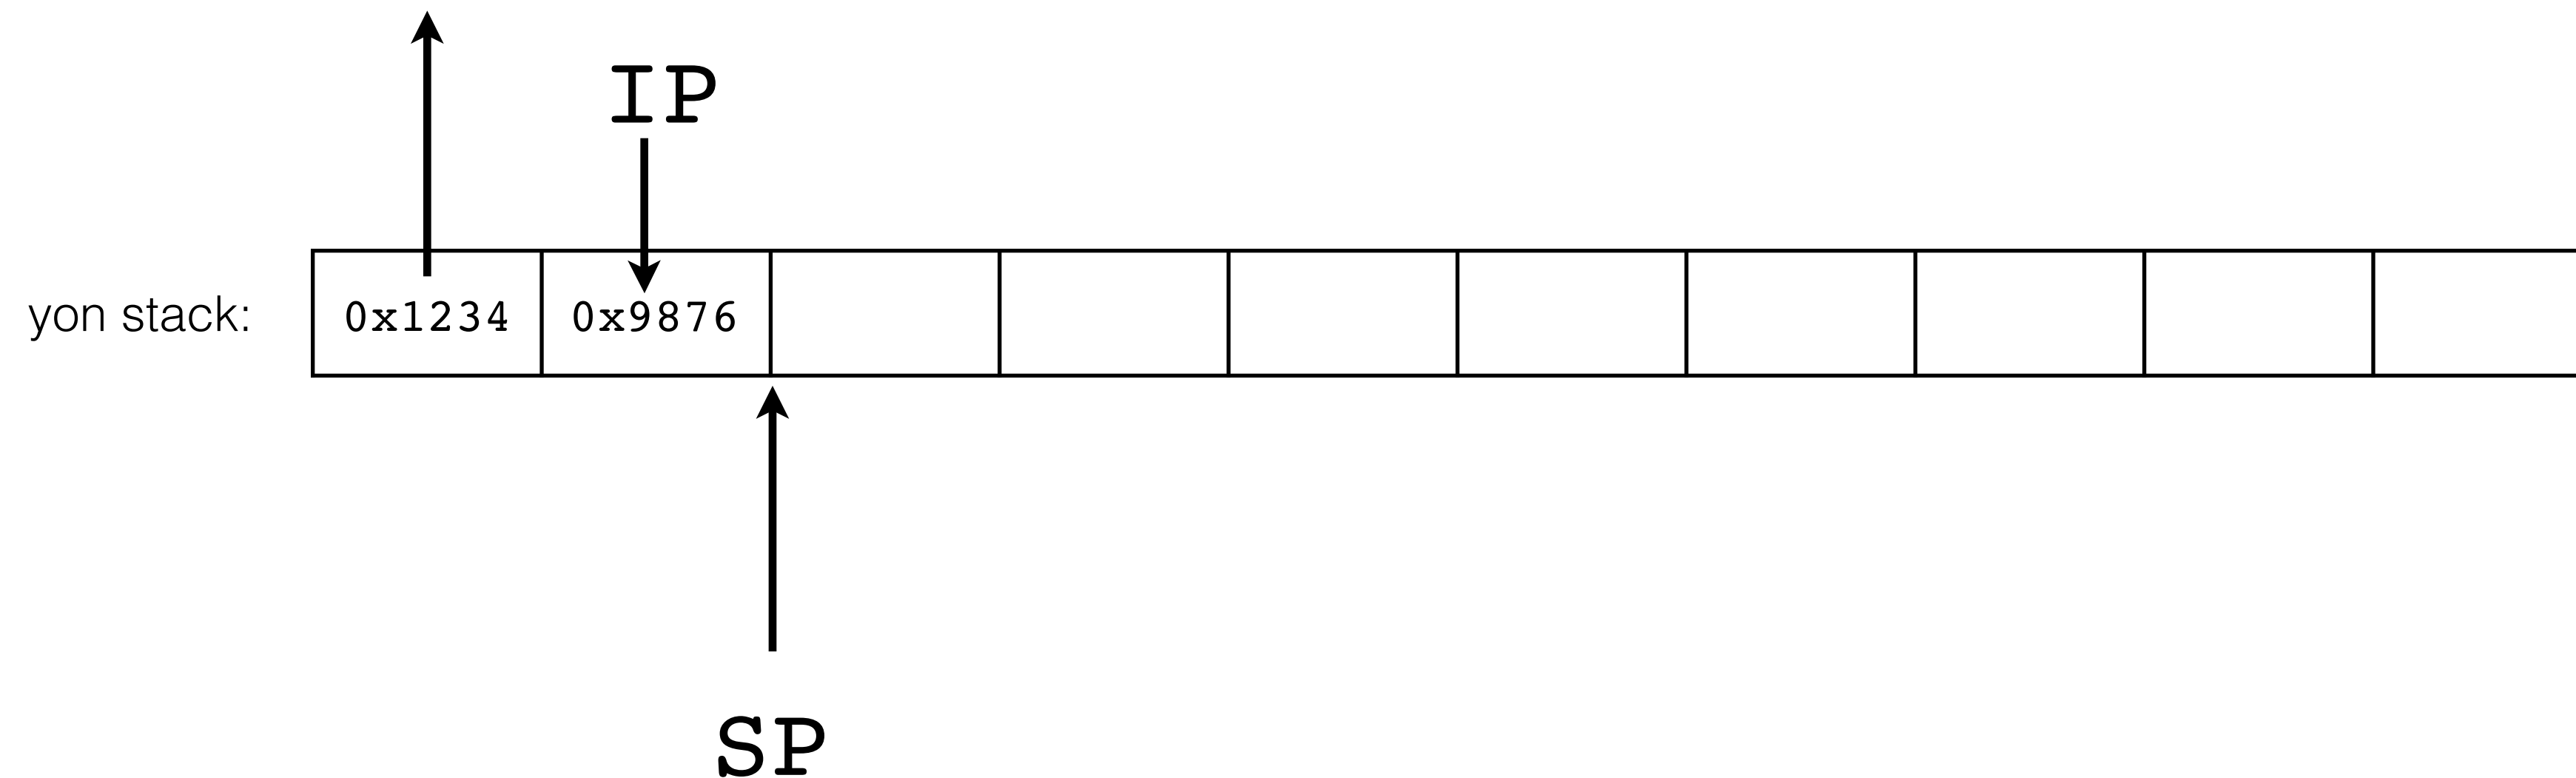
1. exploited function issues `ret`
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

IP → 0x1234 mov ... ← 3. a. execute instructions
0x1236 str ... b. issue `ret`
0x1238 ret ... c. follow step 2



1. exploited function issues `ret`
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

IP → 0x1234 mov ... ← 3. a. execute instructions
0x1236 str ... b. issue `ret`
0x1238 ret ... c. follow step 2



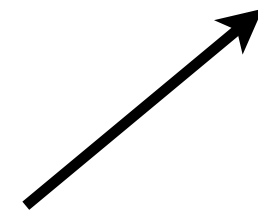
1. exploited function issues `ret`
2. a. two bytes popped off stack
b. SP incremented by two
c. jump to address of bytes

why ROP works:

1. instruction set
2. lots of code in shared libraries

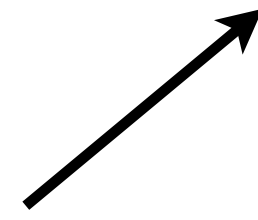
x86 instructions

x86 instructions

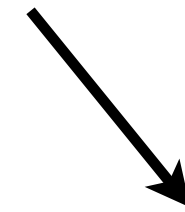


unaligned and variable length

x86 instructions



unaligned and variable length



can jump to point within instruction

x86 instructions

unaligned and variable length

can jump to point within instruction

IP here

f7 c7 07 00 00 00
0f 95 45 c3

test \$0x00000007, %edi
setnzb -61(%ebp)

x86 instructions

unaligned and variable length

can jump to point within instruction

IP here

f7 c7 07 00 00 00
0f 95 45 c3

test \$0x00000007, %edi
setnzb -61(%ebp)

get these instructions

IP here
(one byte over)

c7 07 00 00 00 0f
95
45
c3

movl \$0x0f000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret

x86 instructions

unaligned and variable length

can jump to point within instruction

IP here

f7 c7 07 00 00 00
0f 95 45 c3

test \$0x00000007, %edi
setnzb -61(%ebp)

get these instructions

IP here
(one byte over)

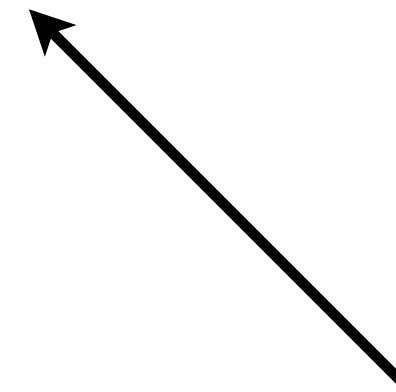
c7 07 00 00 00 0f
95
45
c3

movl \$0x0f000000, (%edi)
xchg %ebp, %eax
inc %ebp
ret

the cool thing: can still work on sensible instruction sets (RISC)

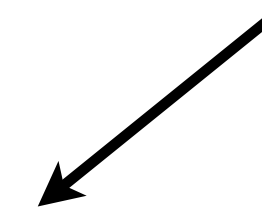
`libc`: lots of instructions

`libc`: lots of instructions



search for sequences that end in `0xc3`

`ret`



`libc`: lots of instructions

search for sequences that end in `0xc3`

`ret`

string together on stack to do useful things

defender is thinking...



...this could be bad

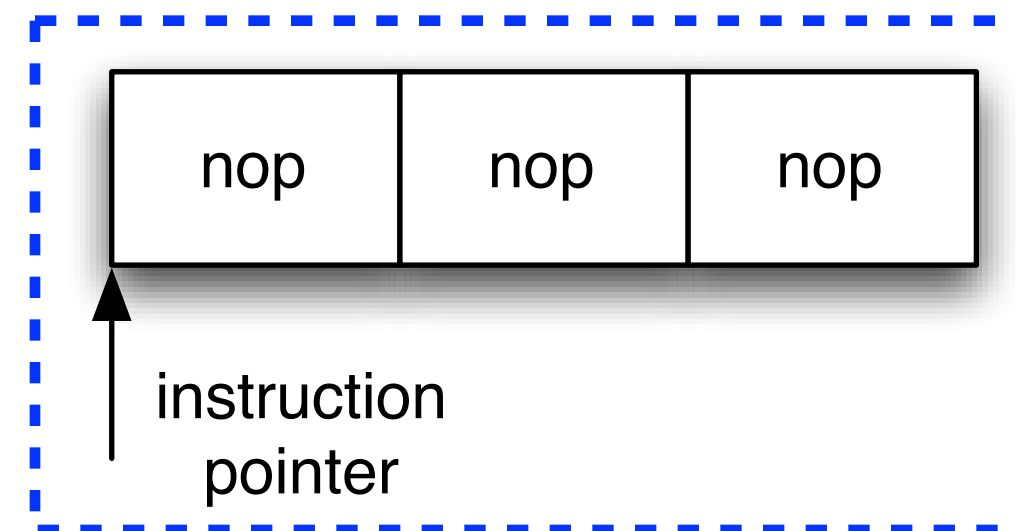
defender is thinking...



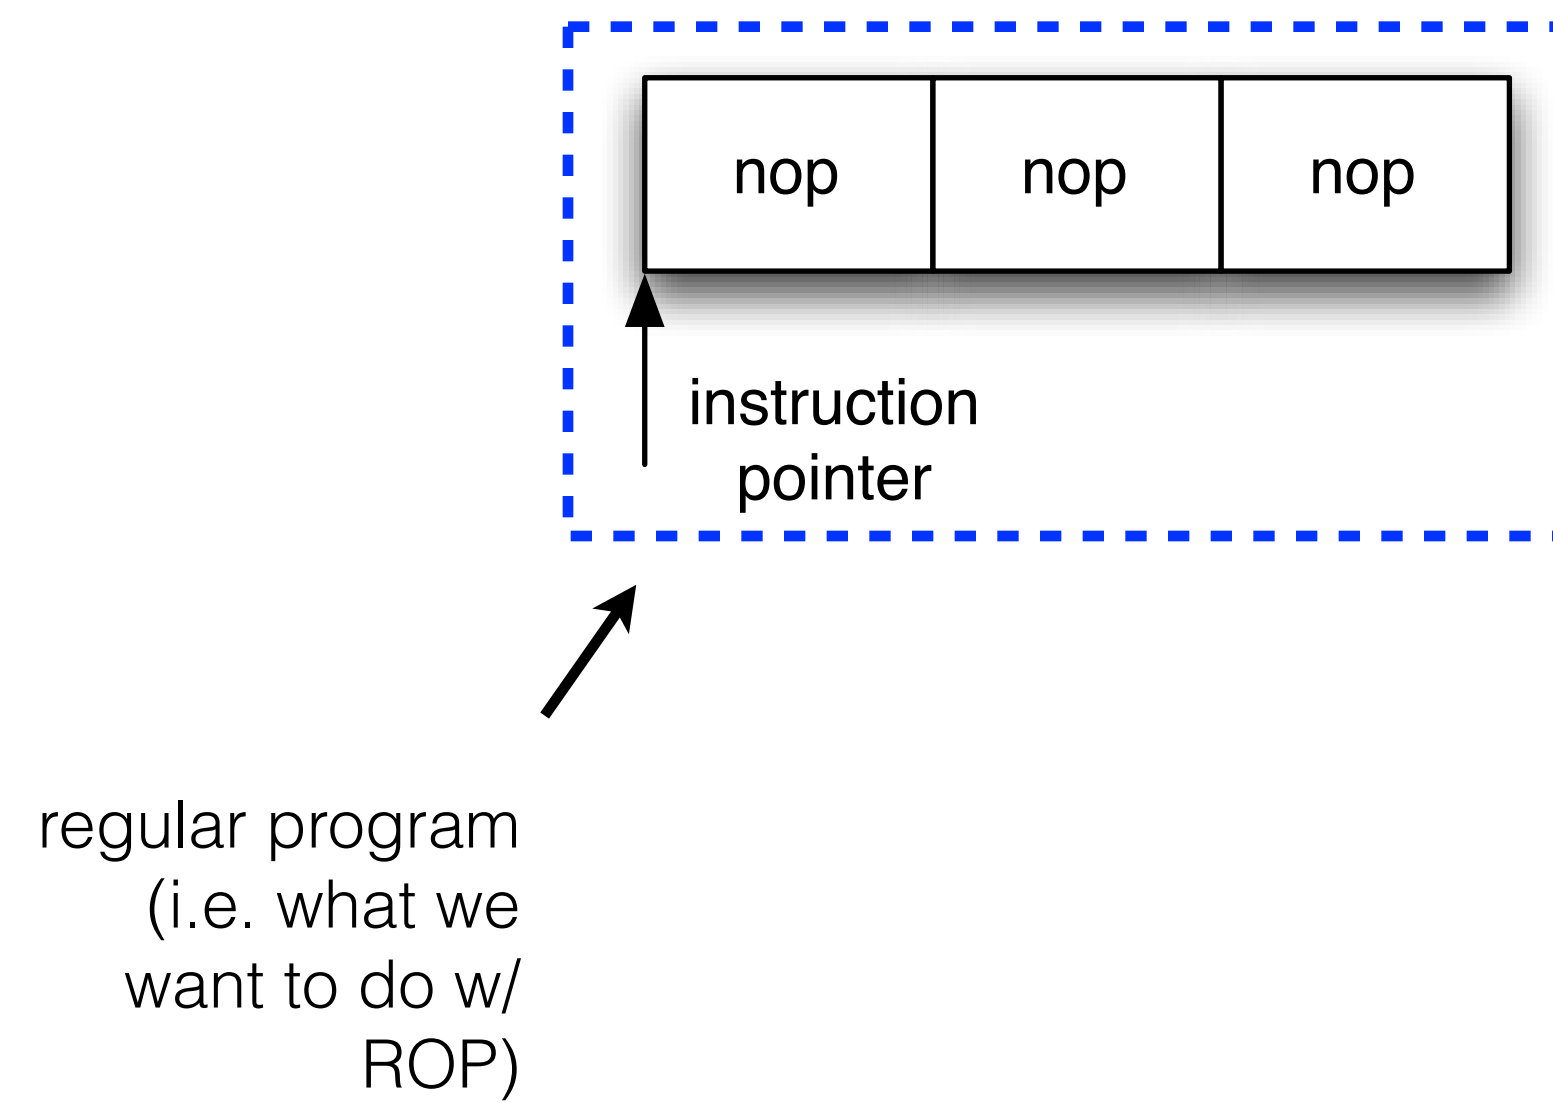
...this could be bad

ROP NOP:

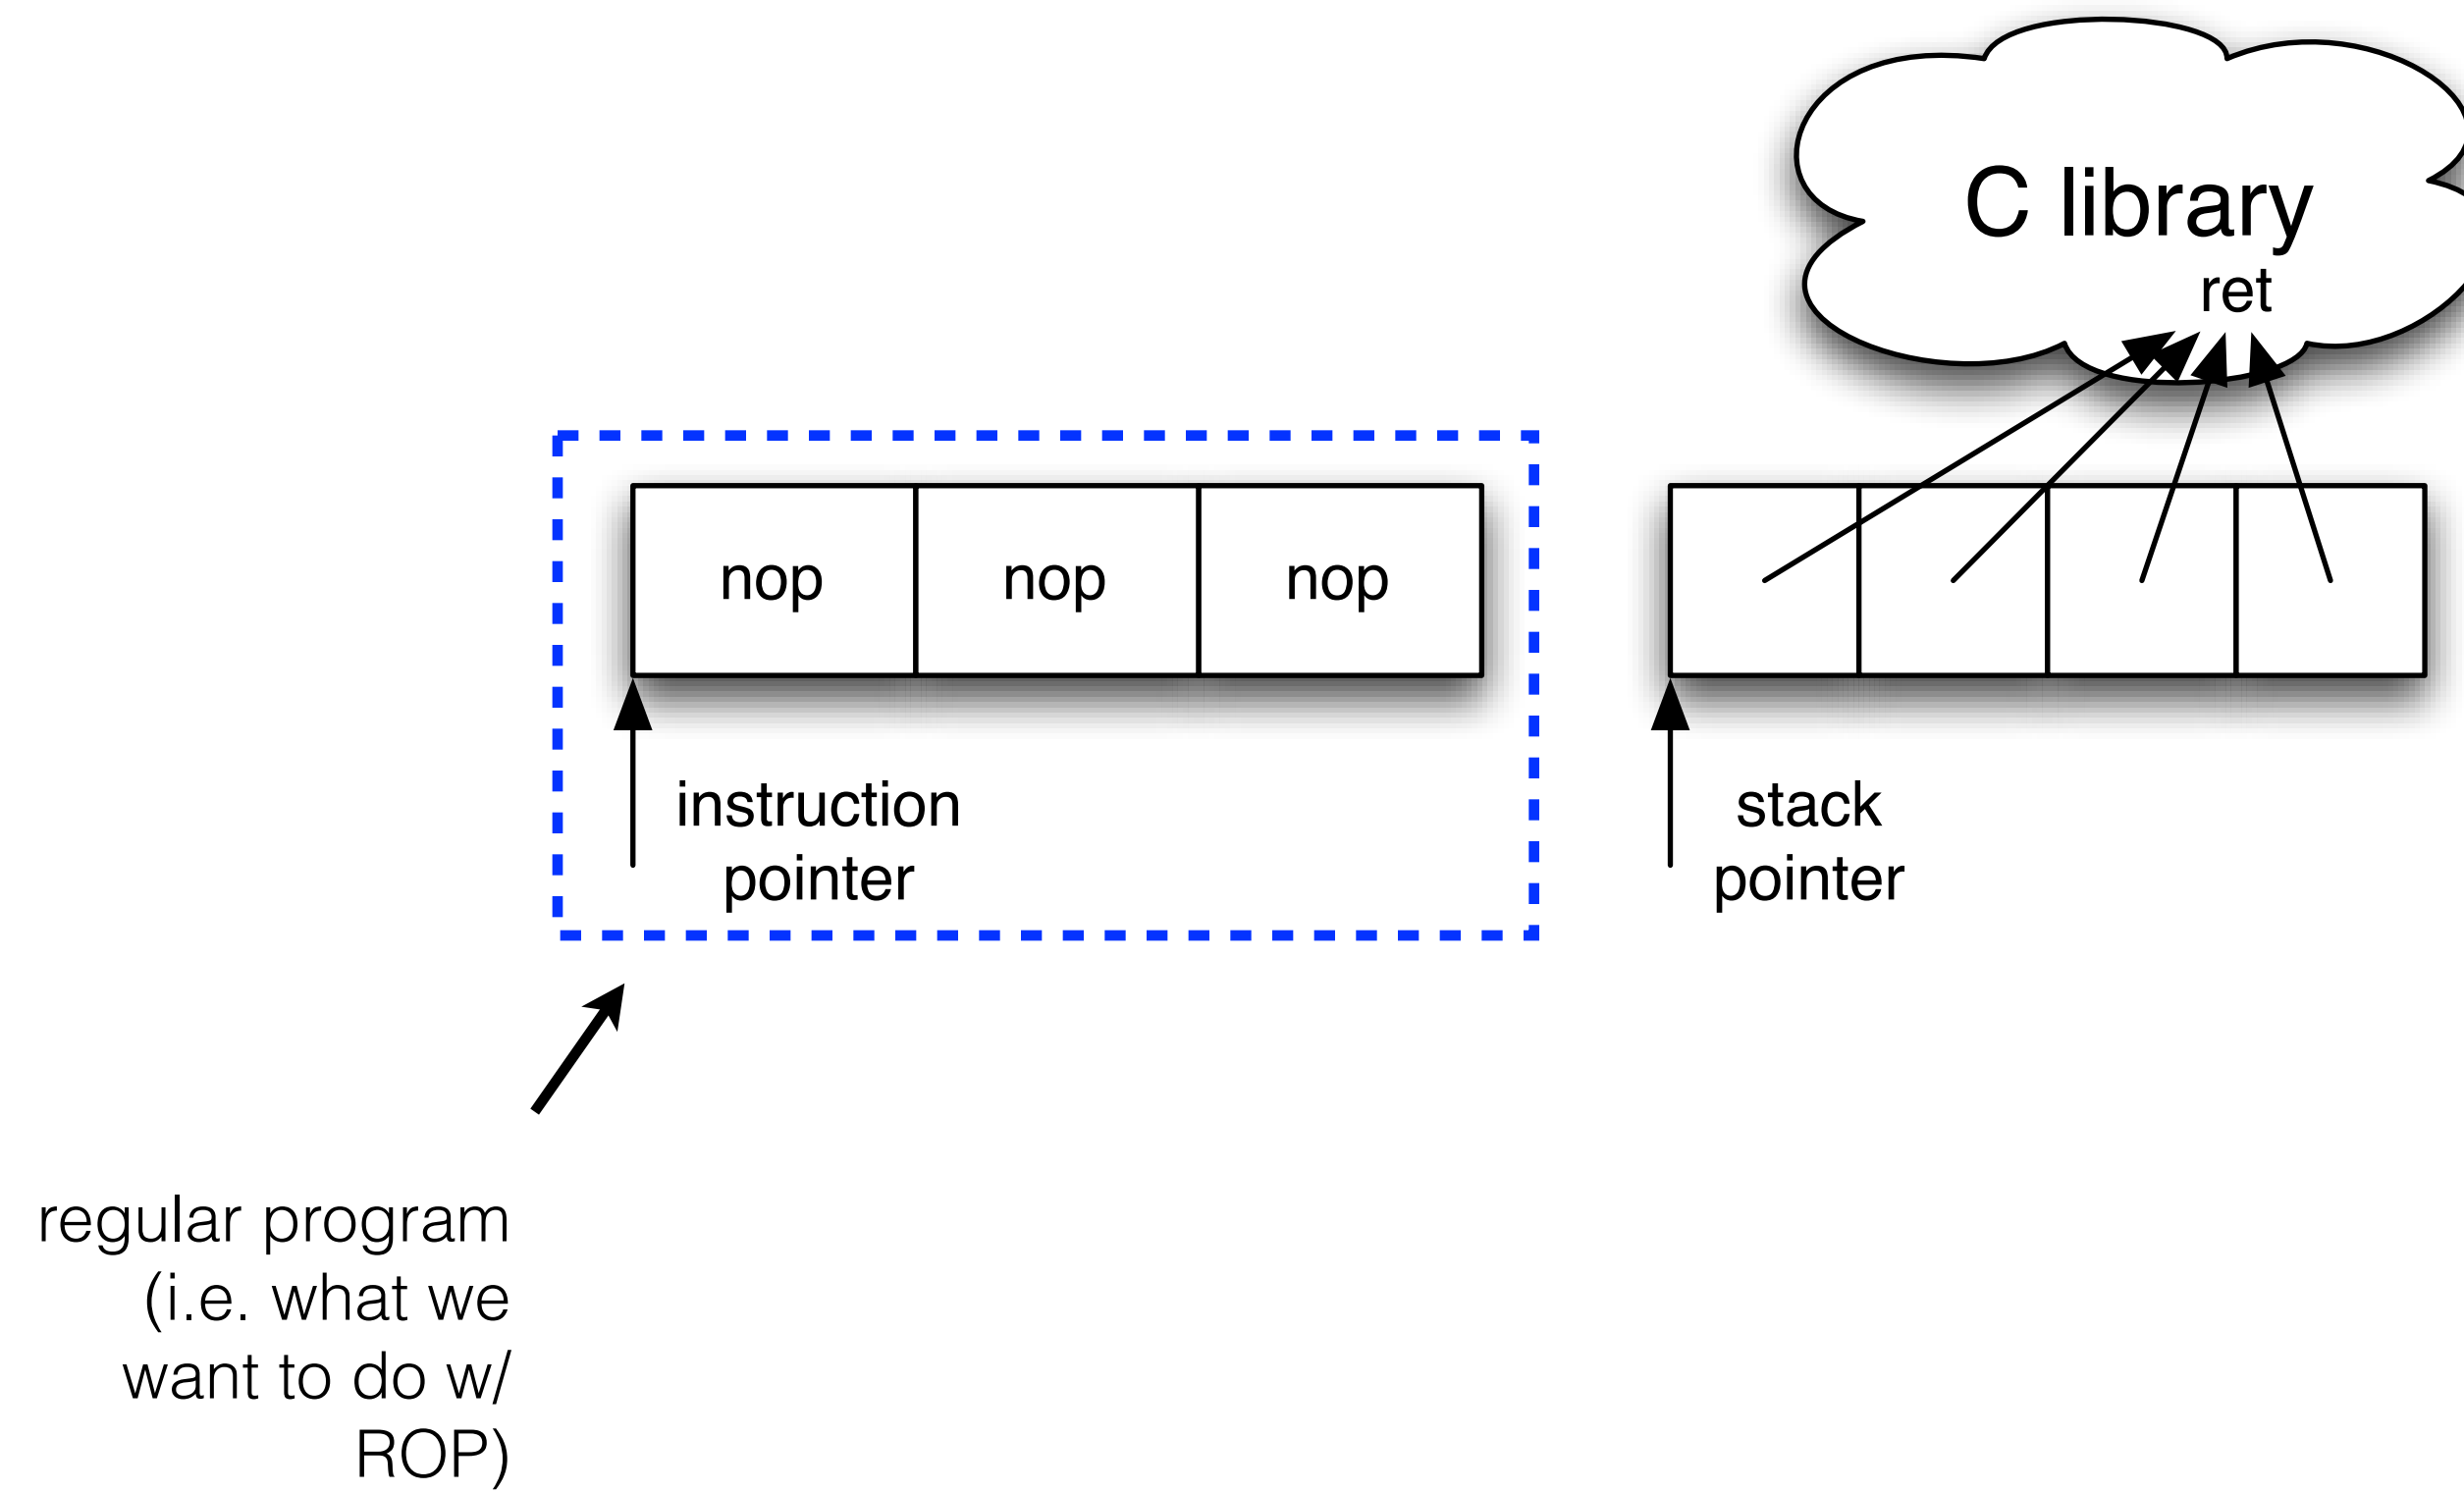
ROP NOP:



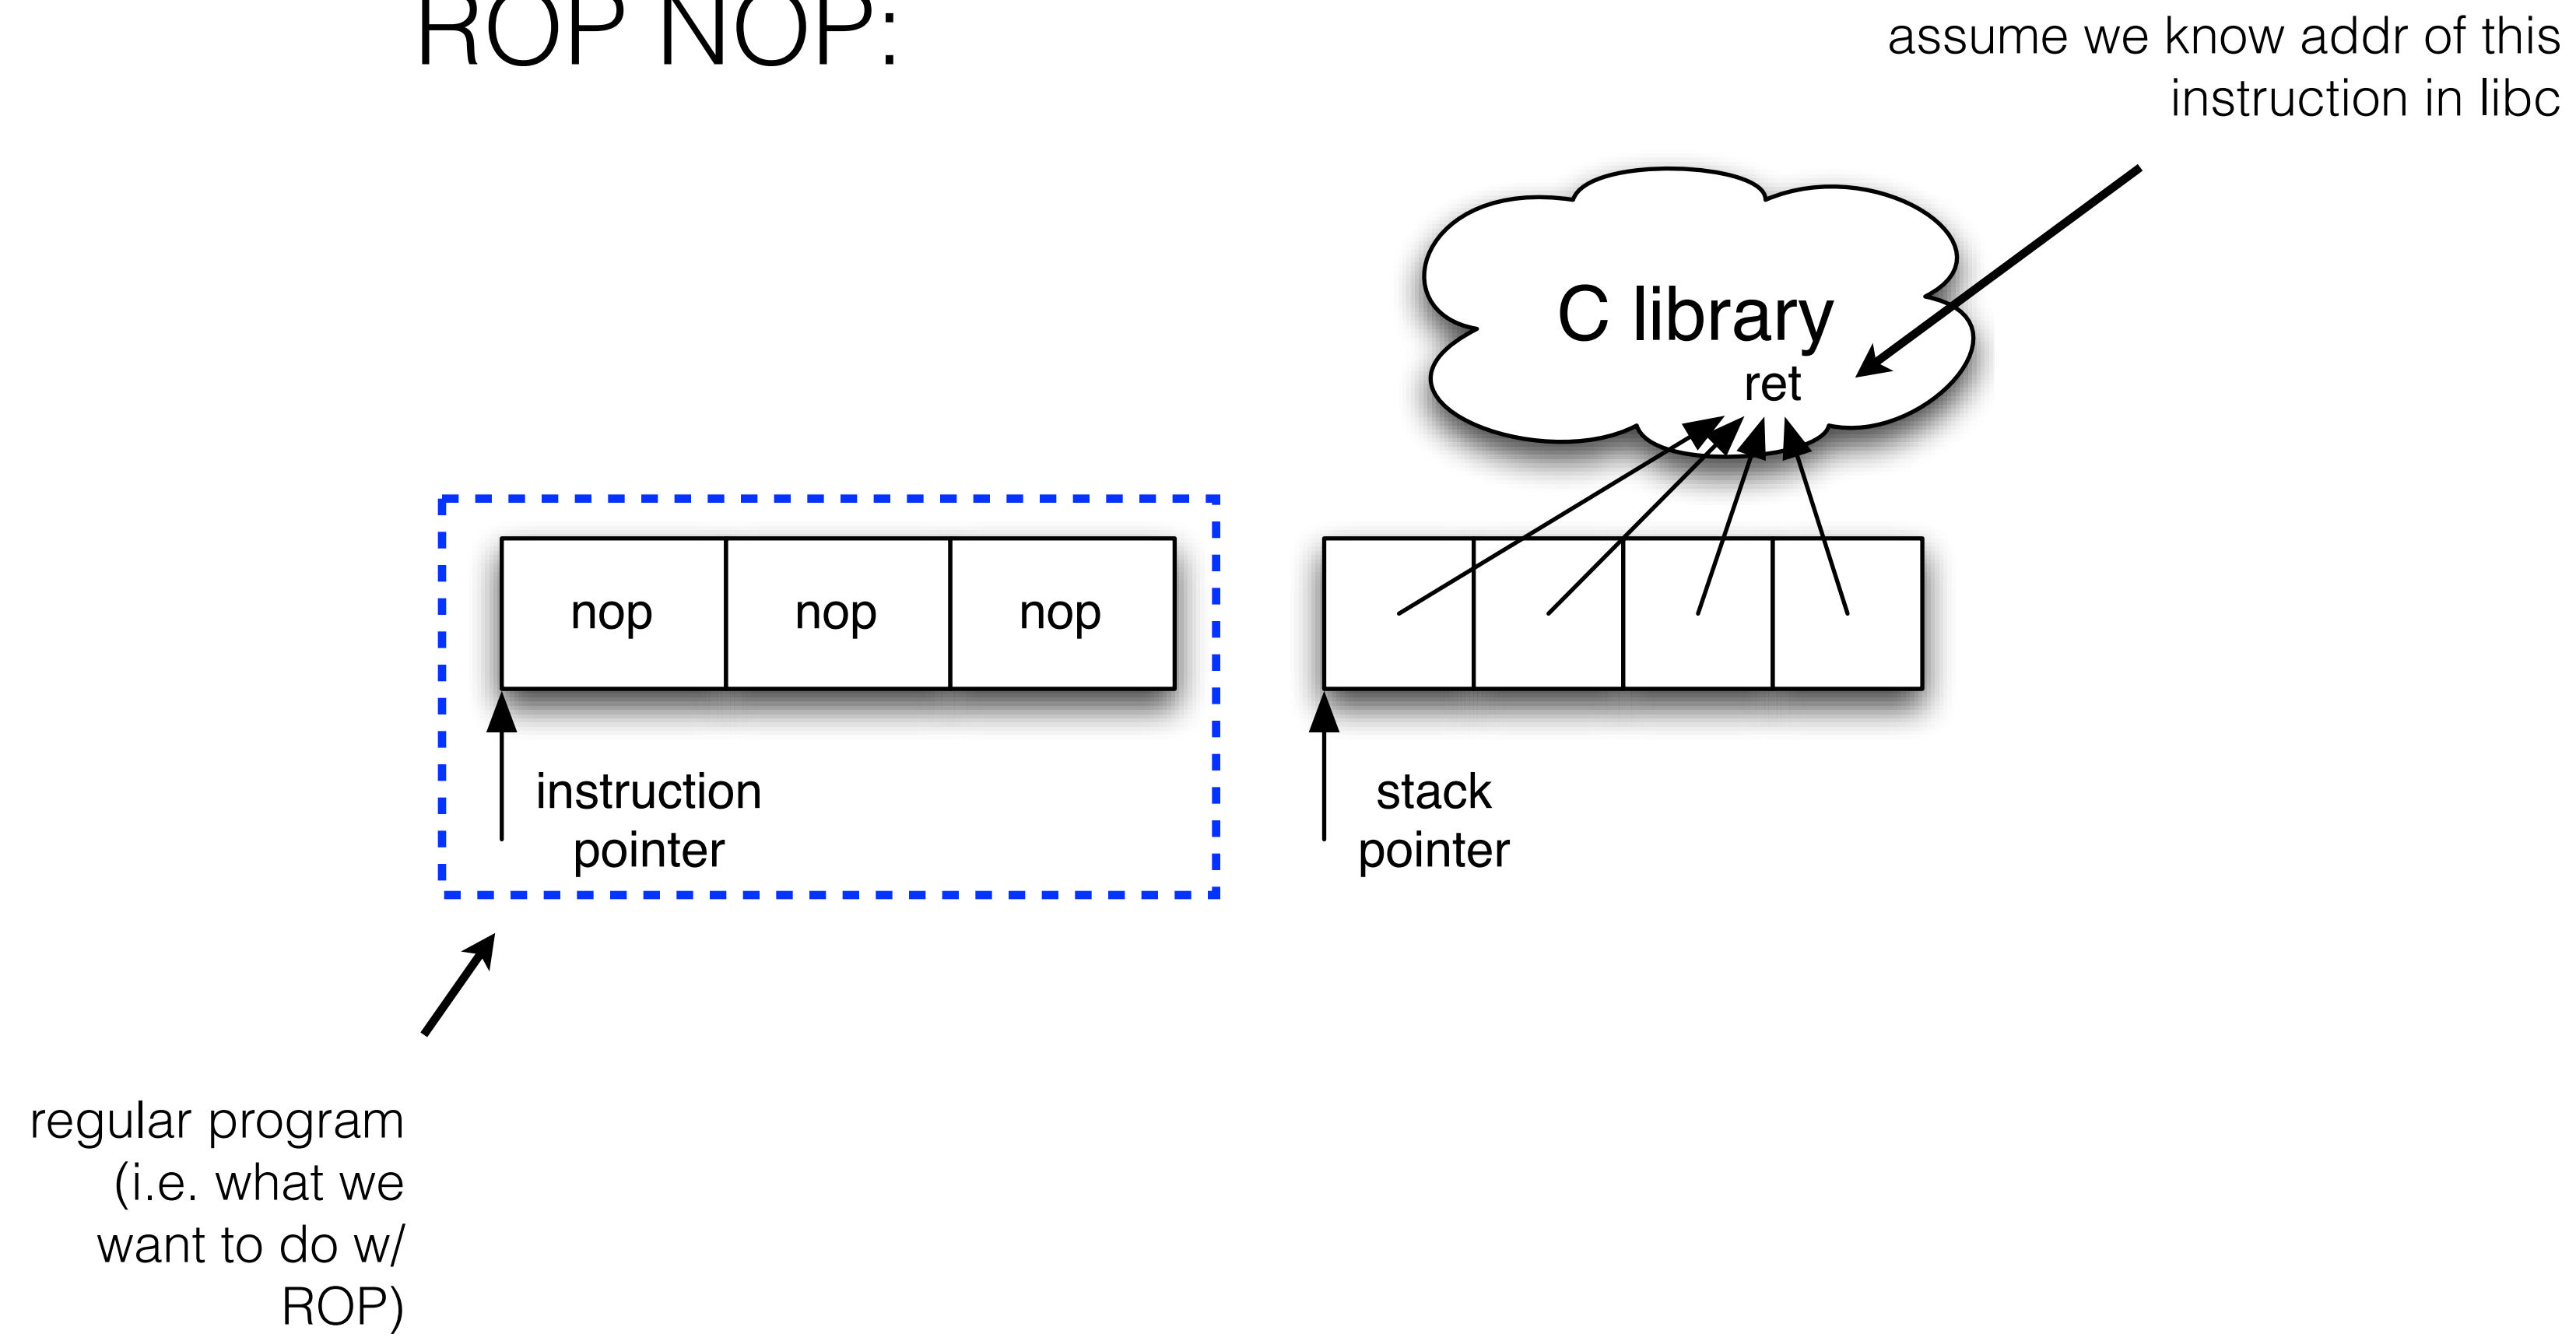
ROP NOP:



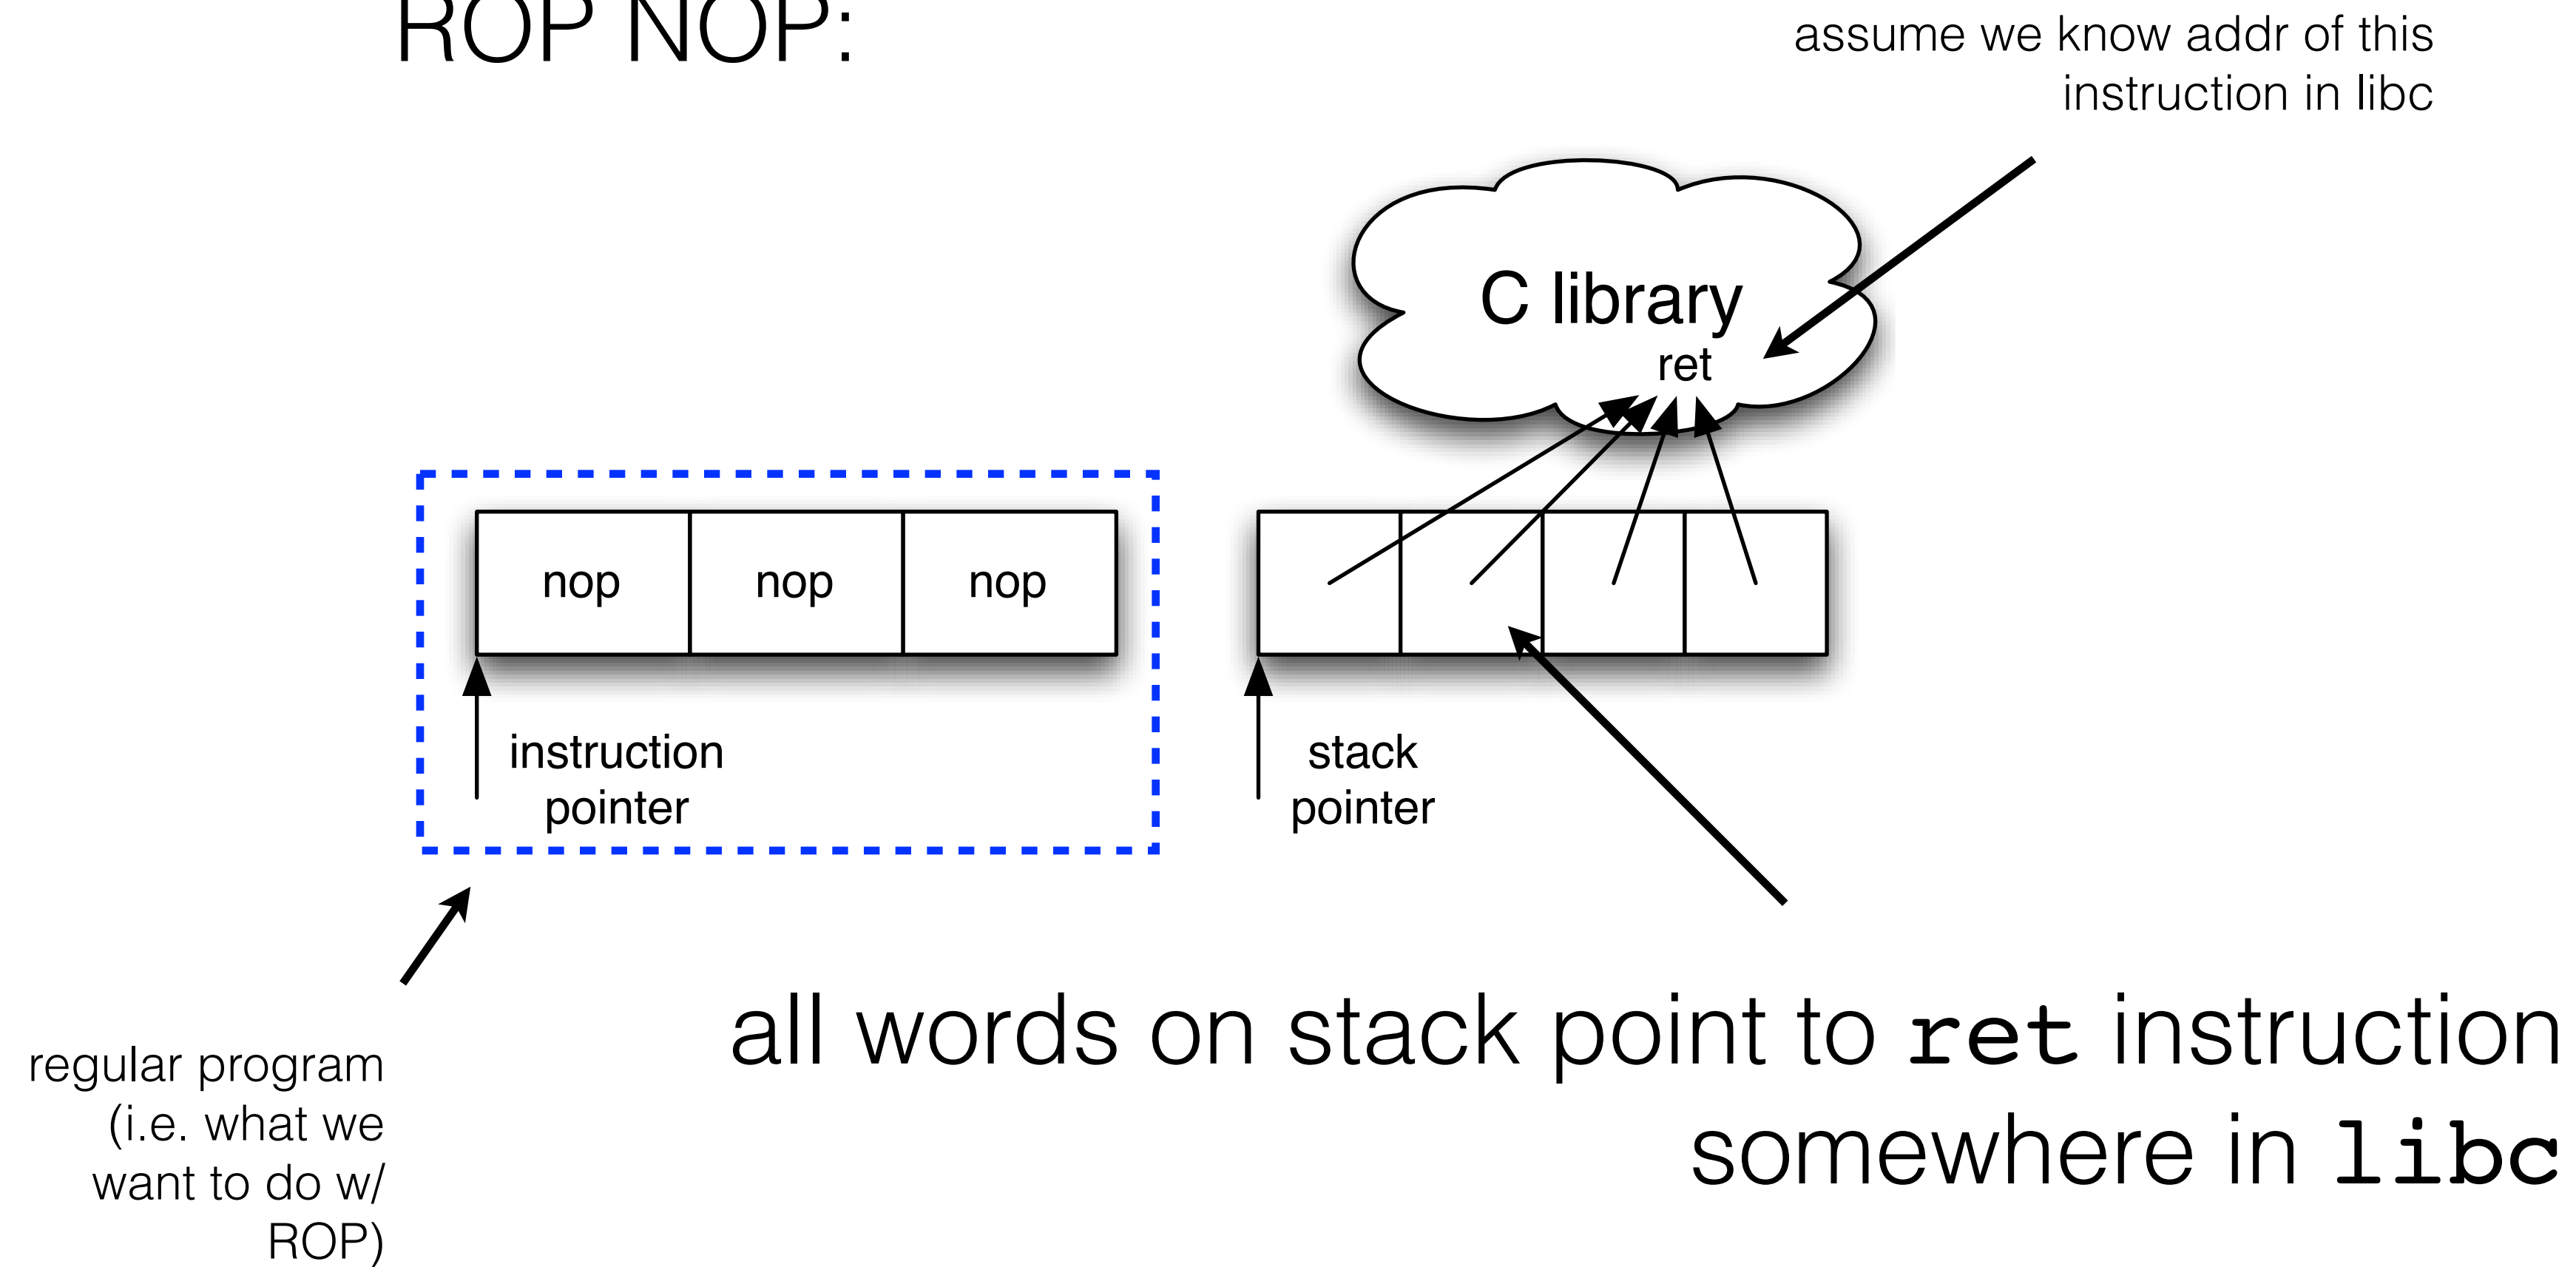
ROP NOP:



ROP NOP:

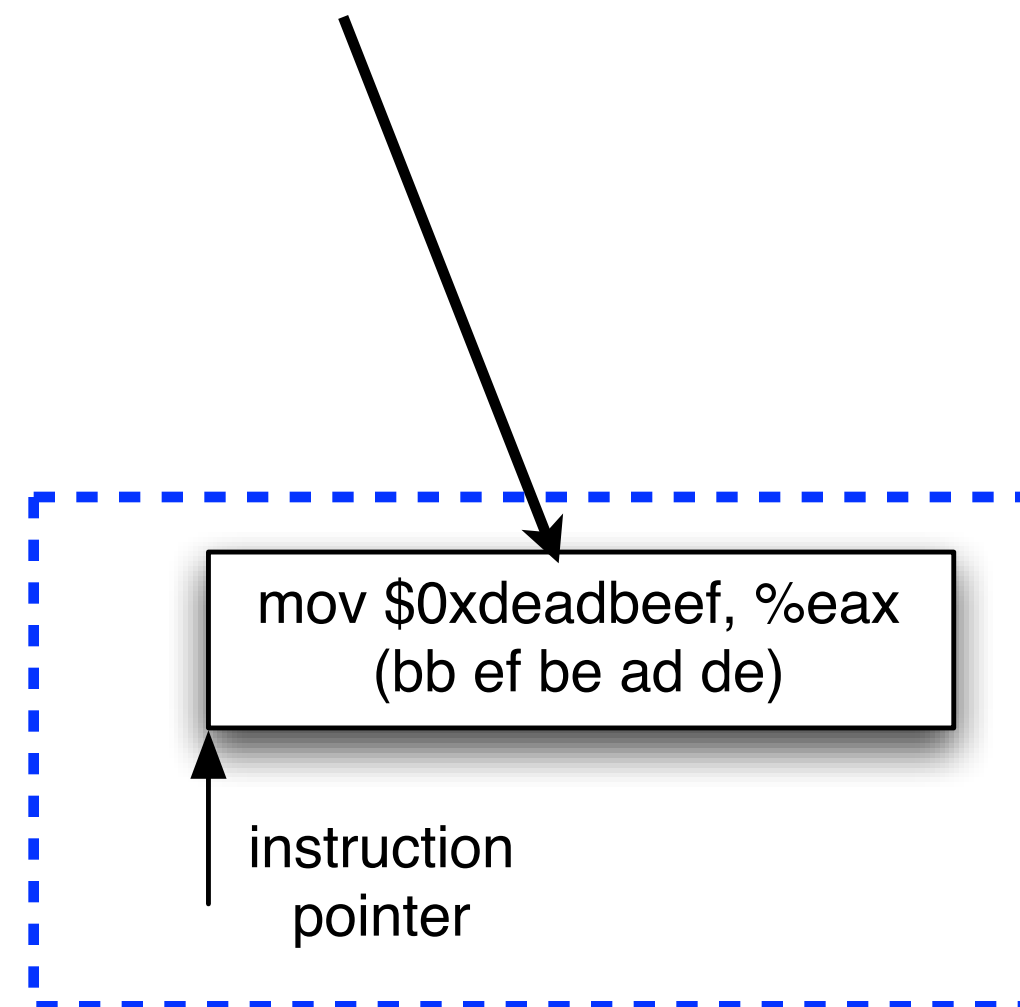


ROP NOP:



ROP immediate load

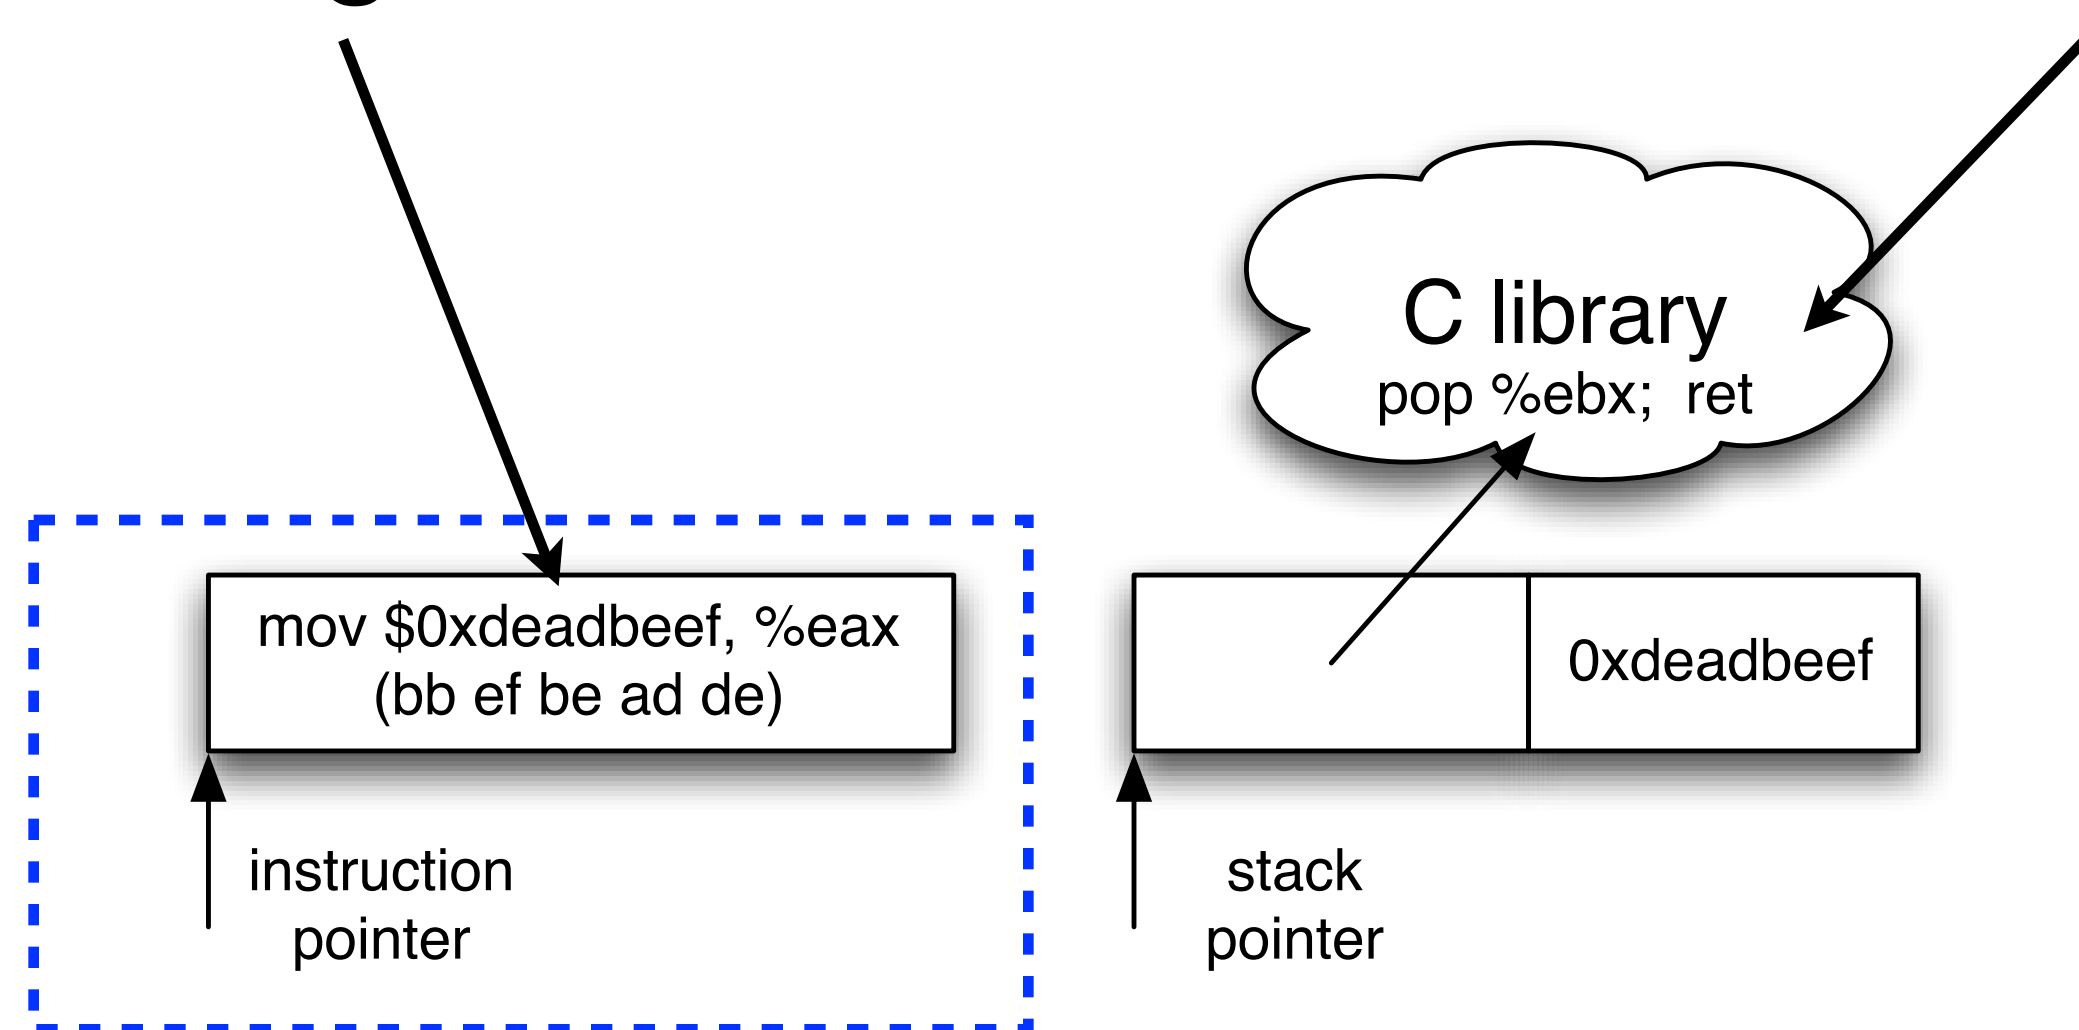
value to load into
register



ROP immediate load

value to load into
register

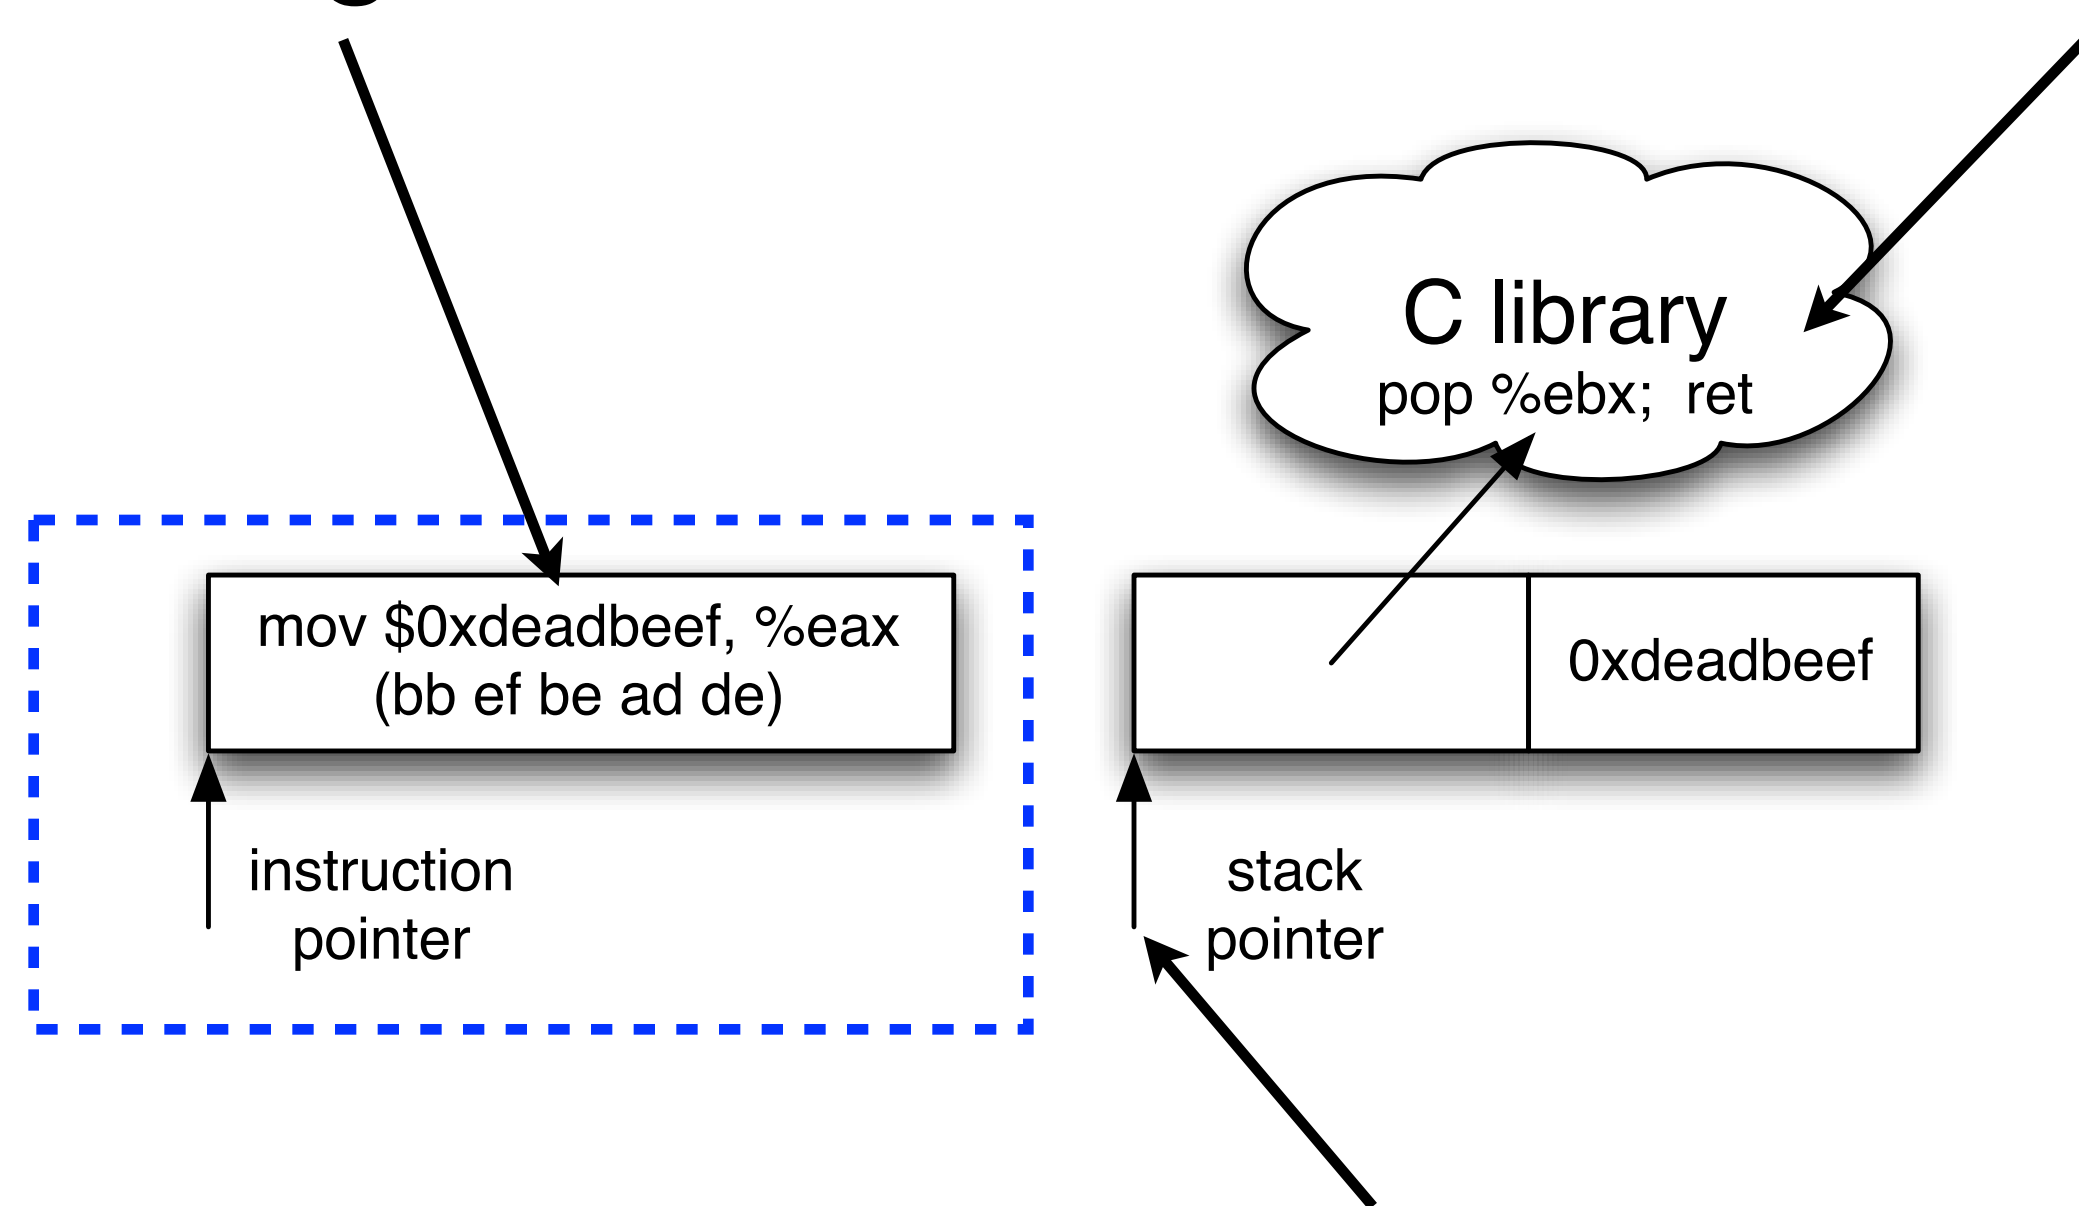
assume we know addr of this instruction
sequence in libc



ROP immediate load

value to load into
register

assume we know addr of this instruction
sequence in libc

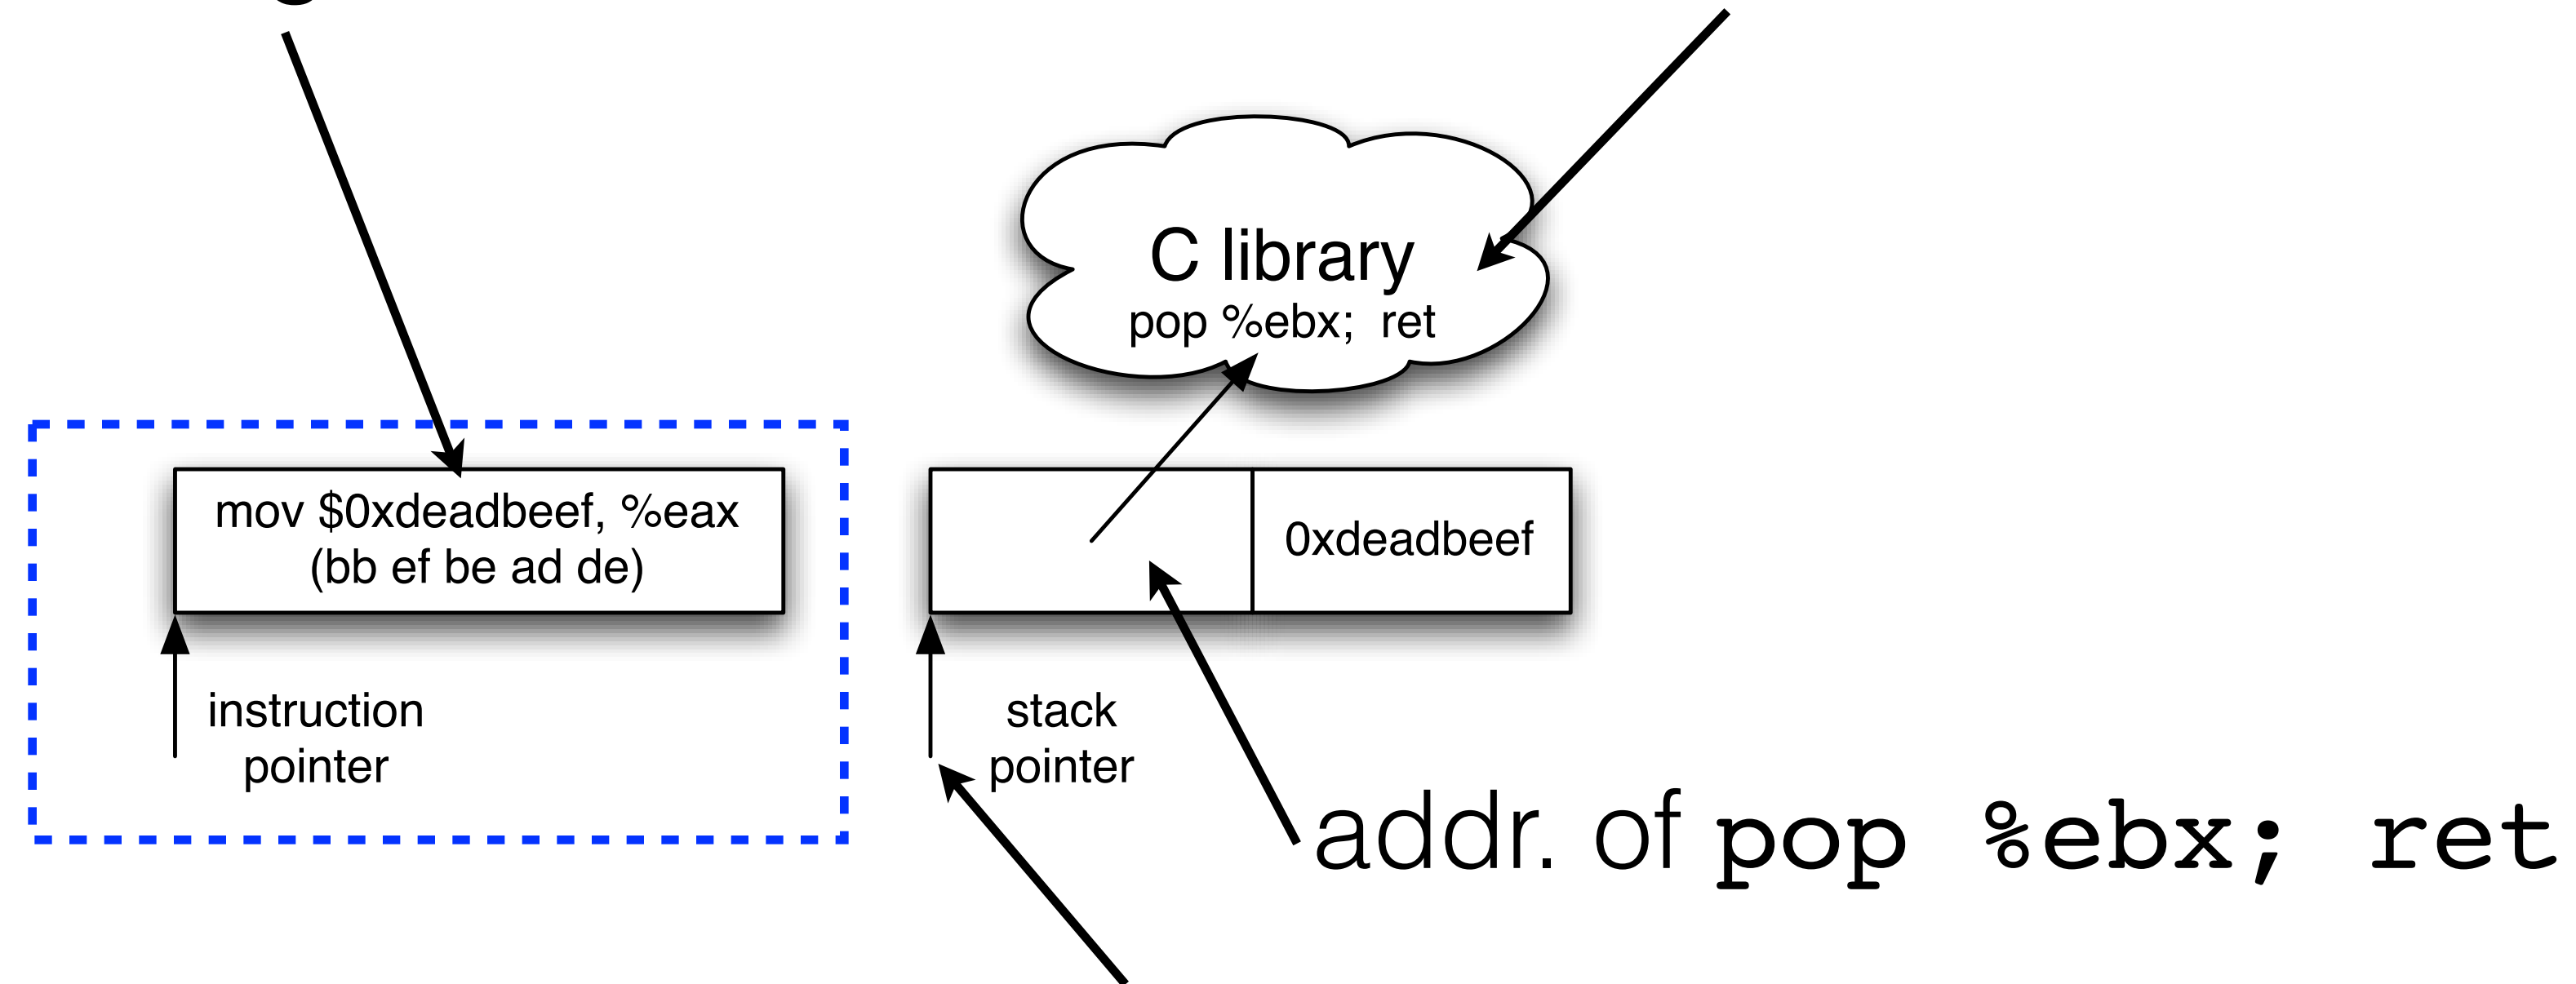


1. `ret` issued with SP here

ROP immediate load

value to load into register

assume we know addr of this instruction sequence in libc

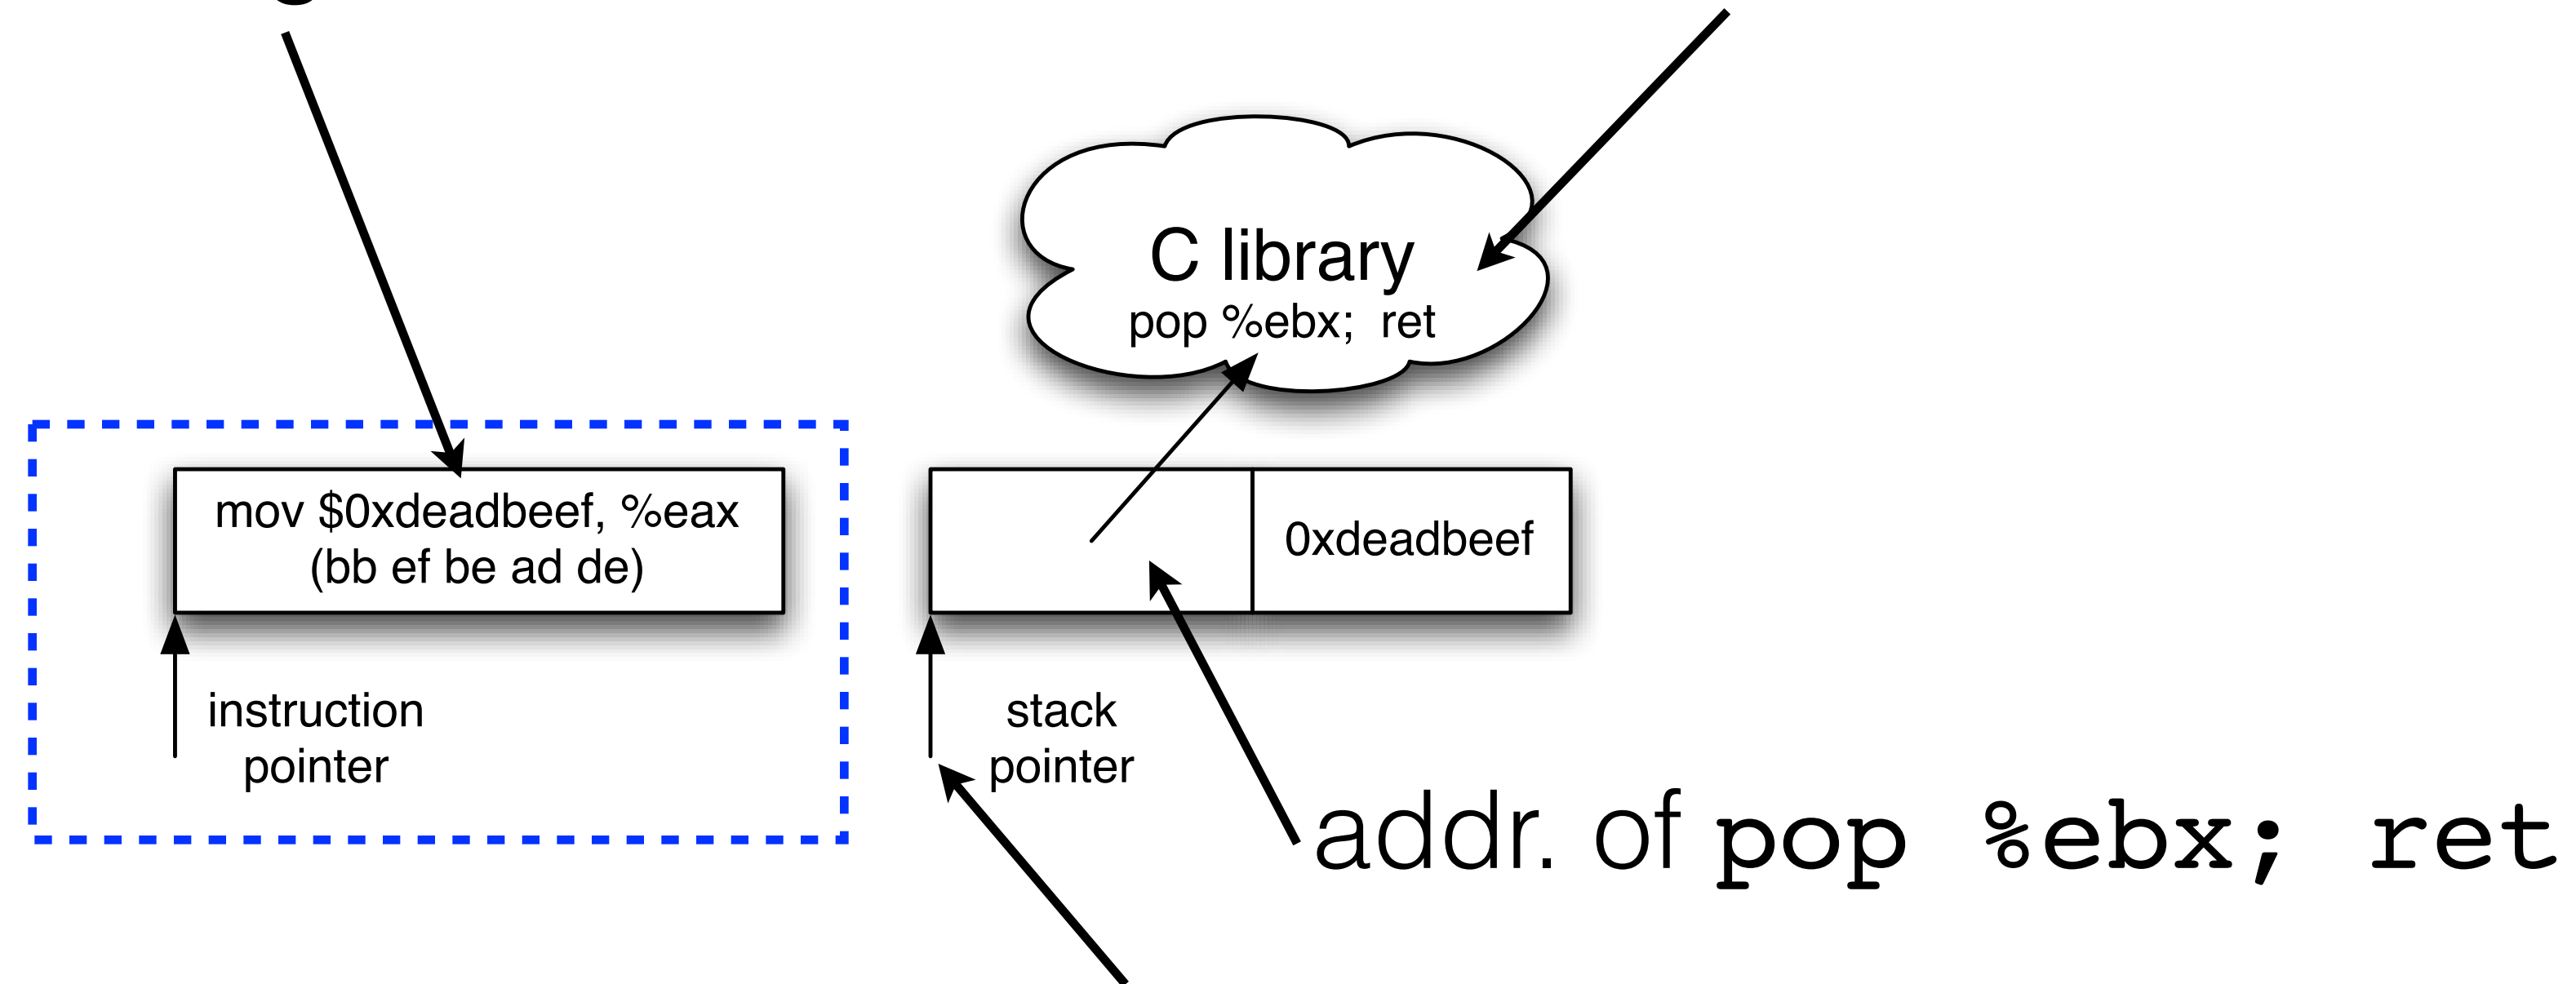


1. `ret` issued with SP here
2. IP pointed at addr of `pop...`

ROP immediate load

value to load into register

assume we know addr of this instruction sequence in libc

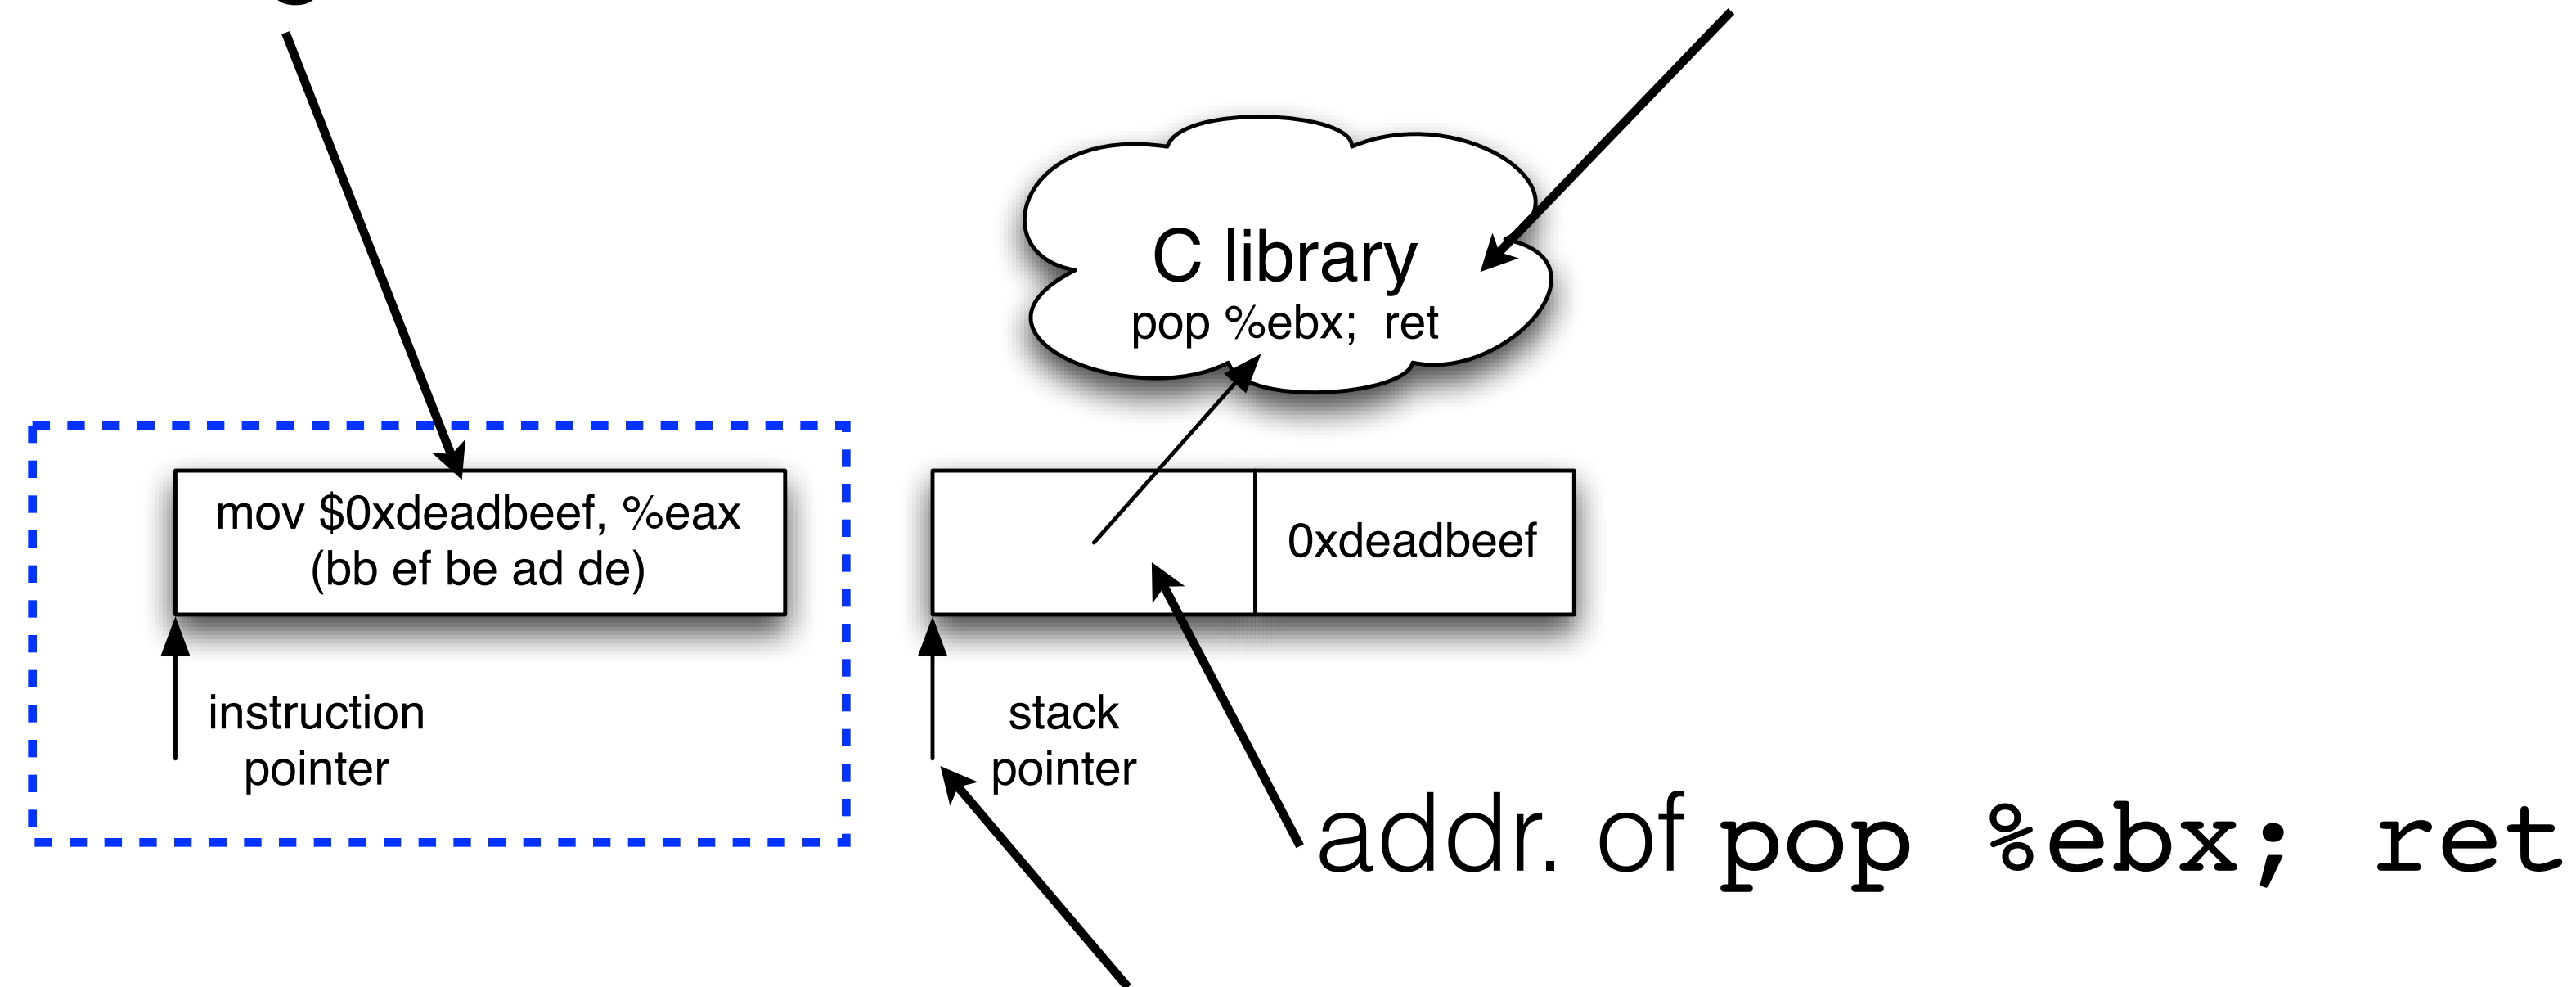


1. `ret` issued with SP here
2. IP pointed at addr of `pop...`
3. SP moved right one word

ROP immediate load

value to load into register

assume we know addr of this instruction sequence in libc

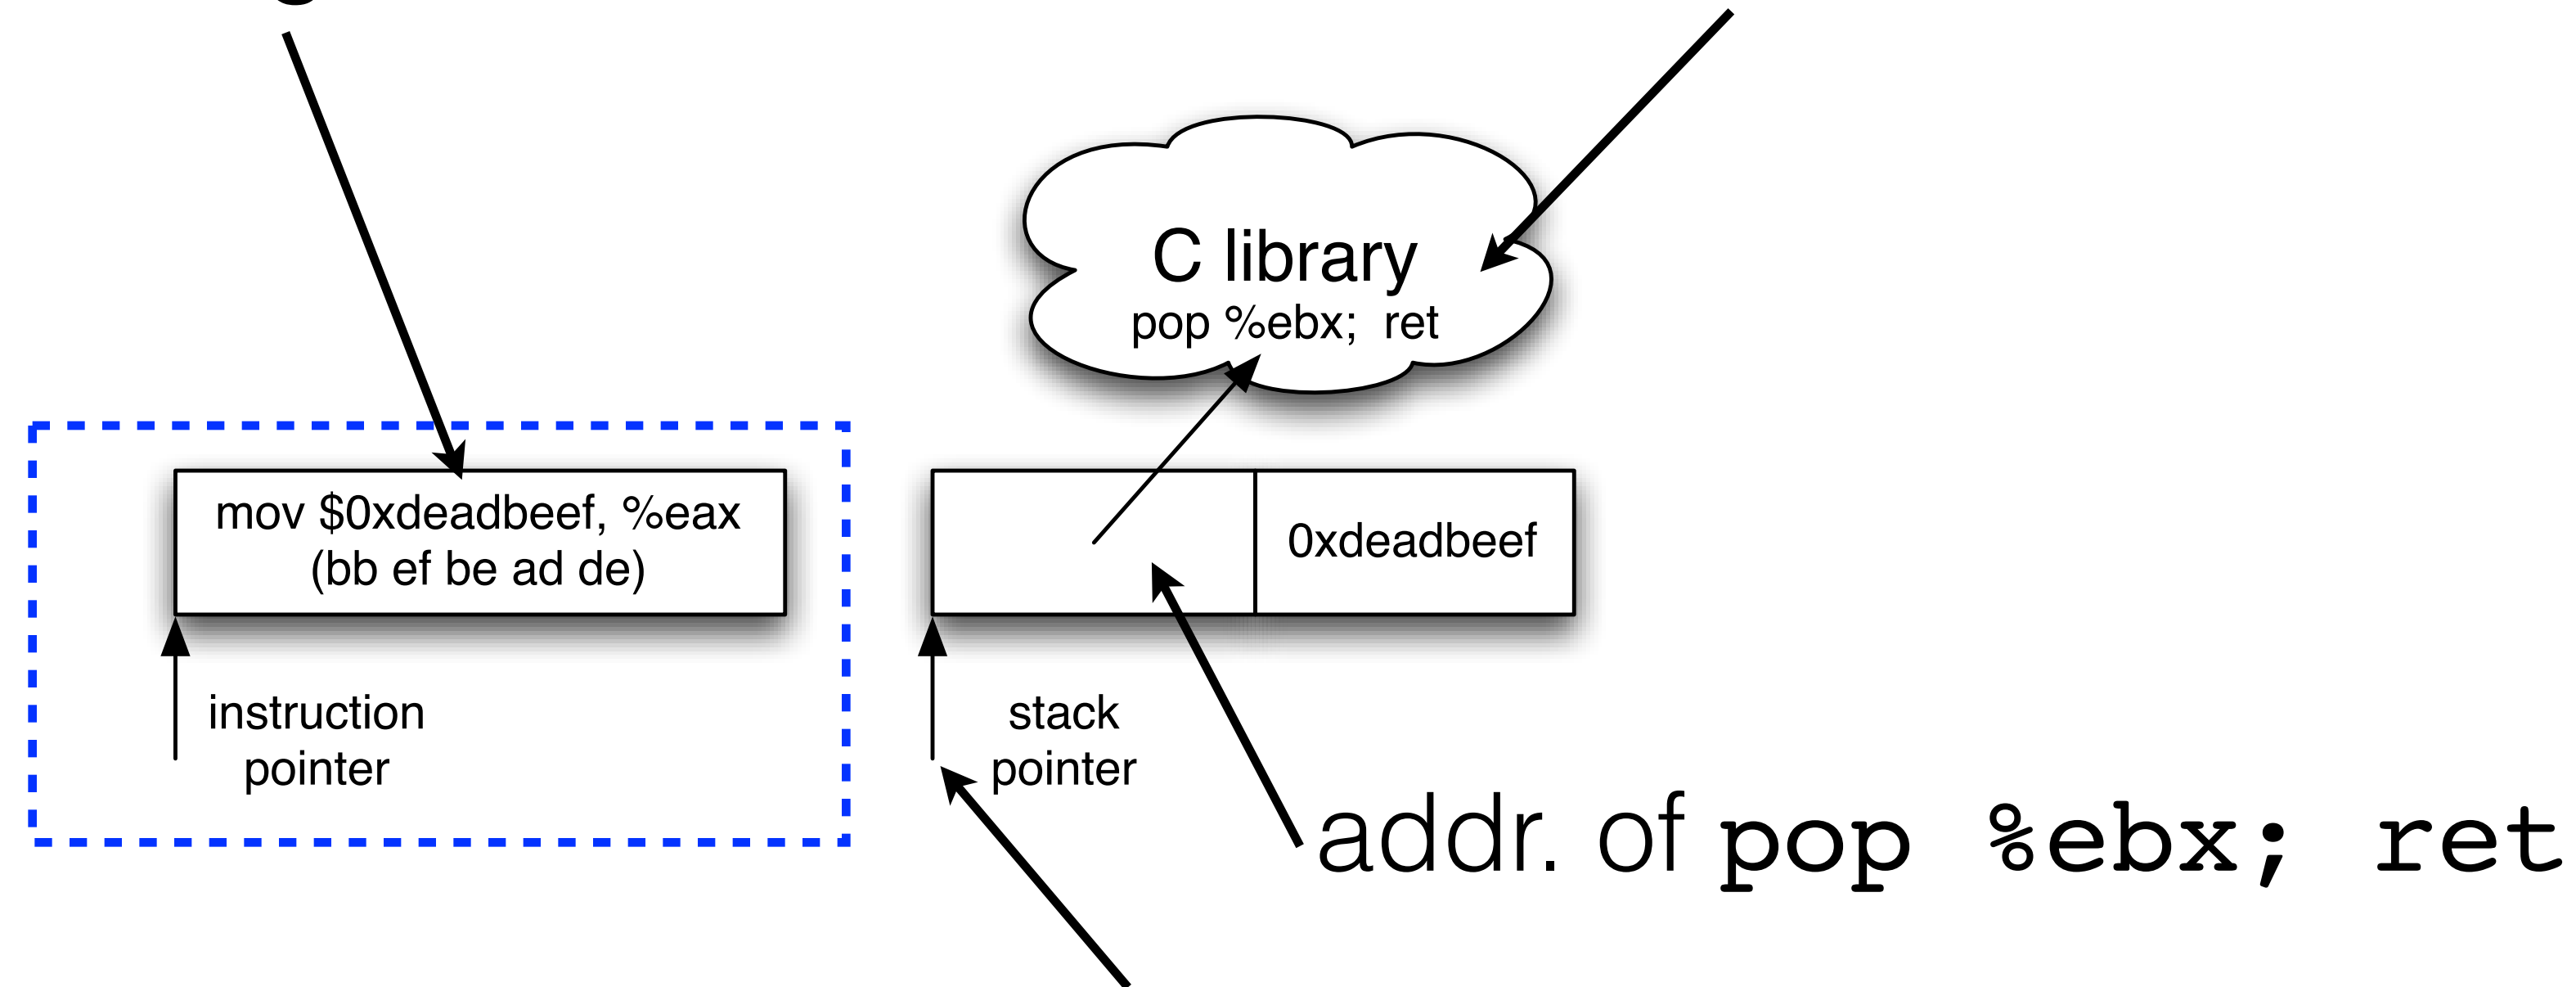


1. `ret` issued with SP here
2. IP pointed at addr of `pop...`
3. SP moved right one word
4. `pop` puts `0xdeadbeef` into `ebx`

ROP immediate load

value to load into register

assume we know addr of this instruction sequence in libc

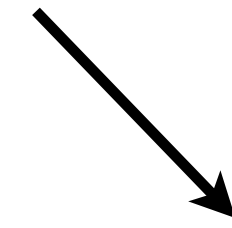


1. `ret` issued with SP here
2. IP pointed at addr of `pop...`
3. SP moved right one word
4. `pop` puts `0xdeadbeef` into `ebx`
5. SP moved right one word

ROP word load

dot represents constructed stack

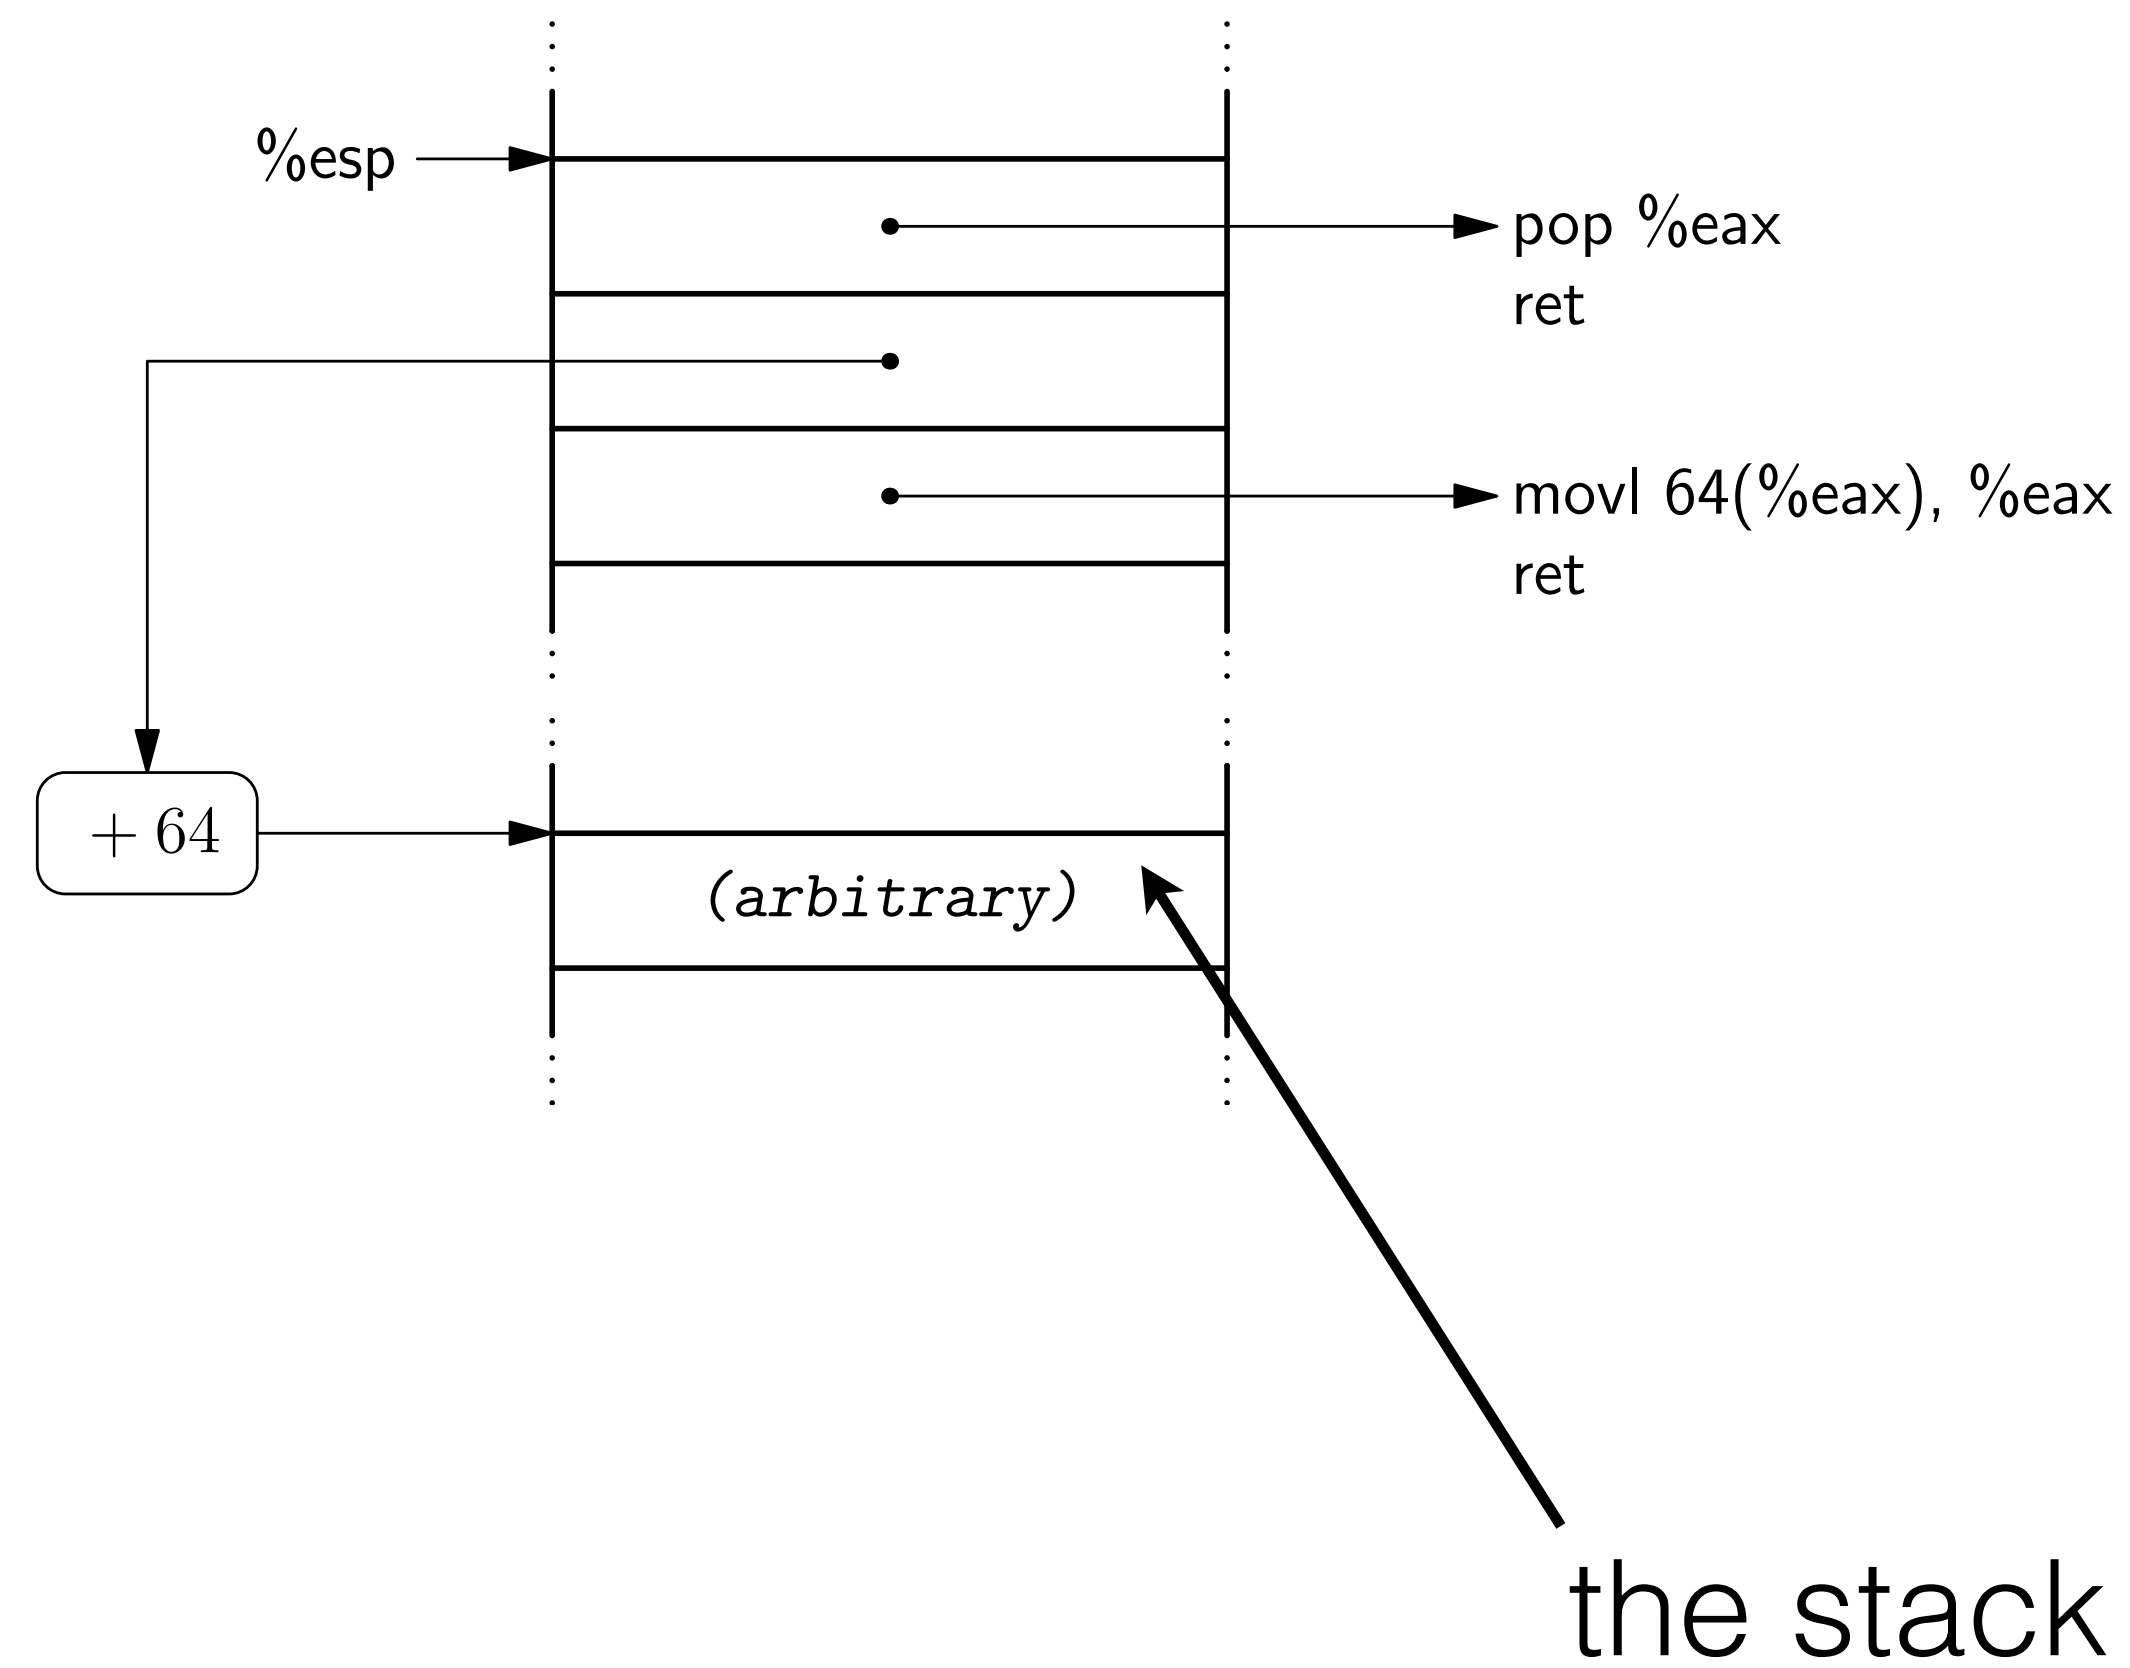
ROP word load



word into %eax

ROP word load

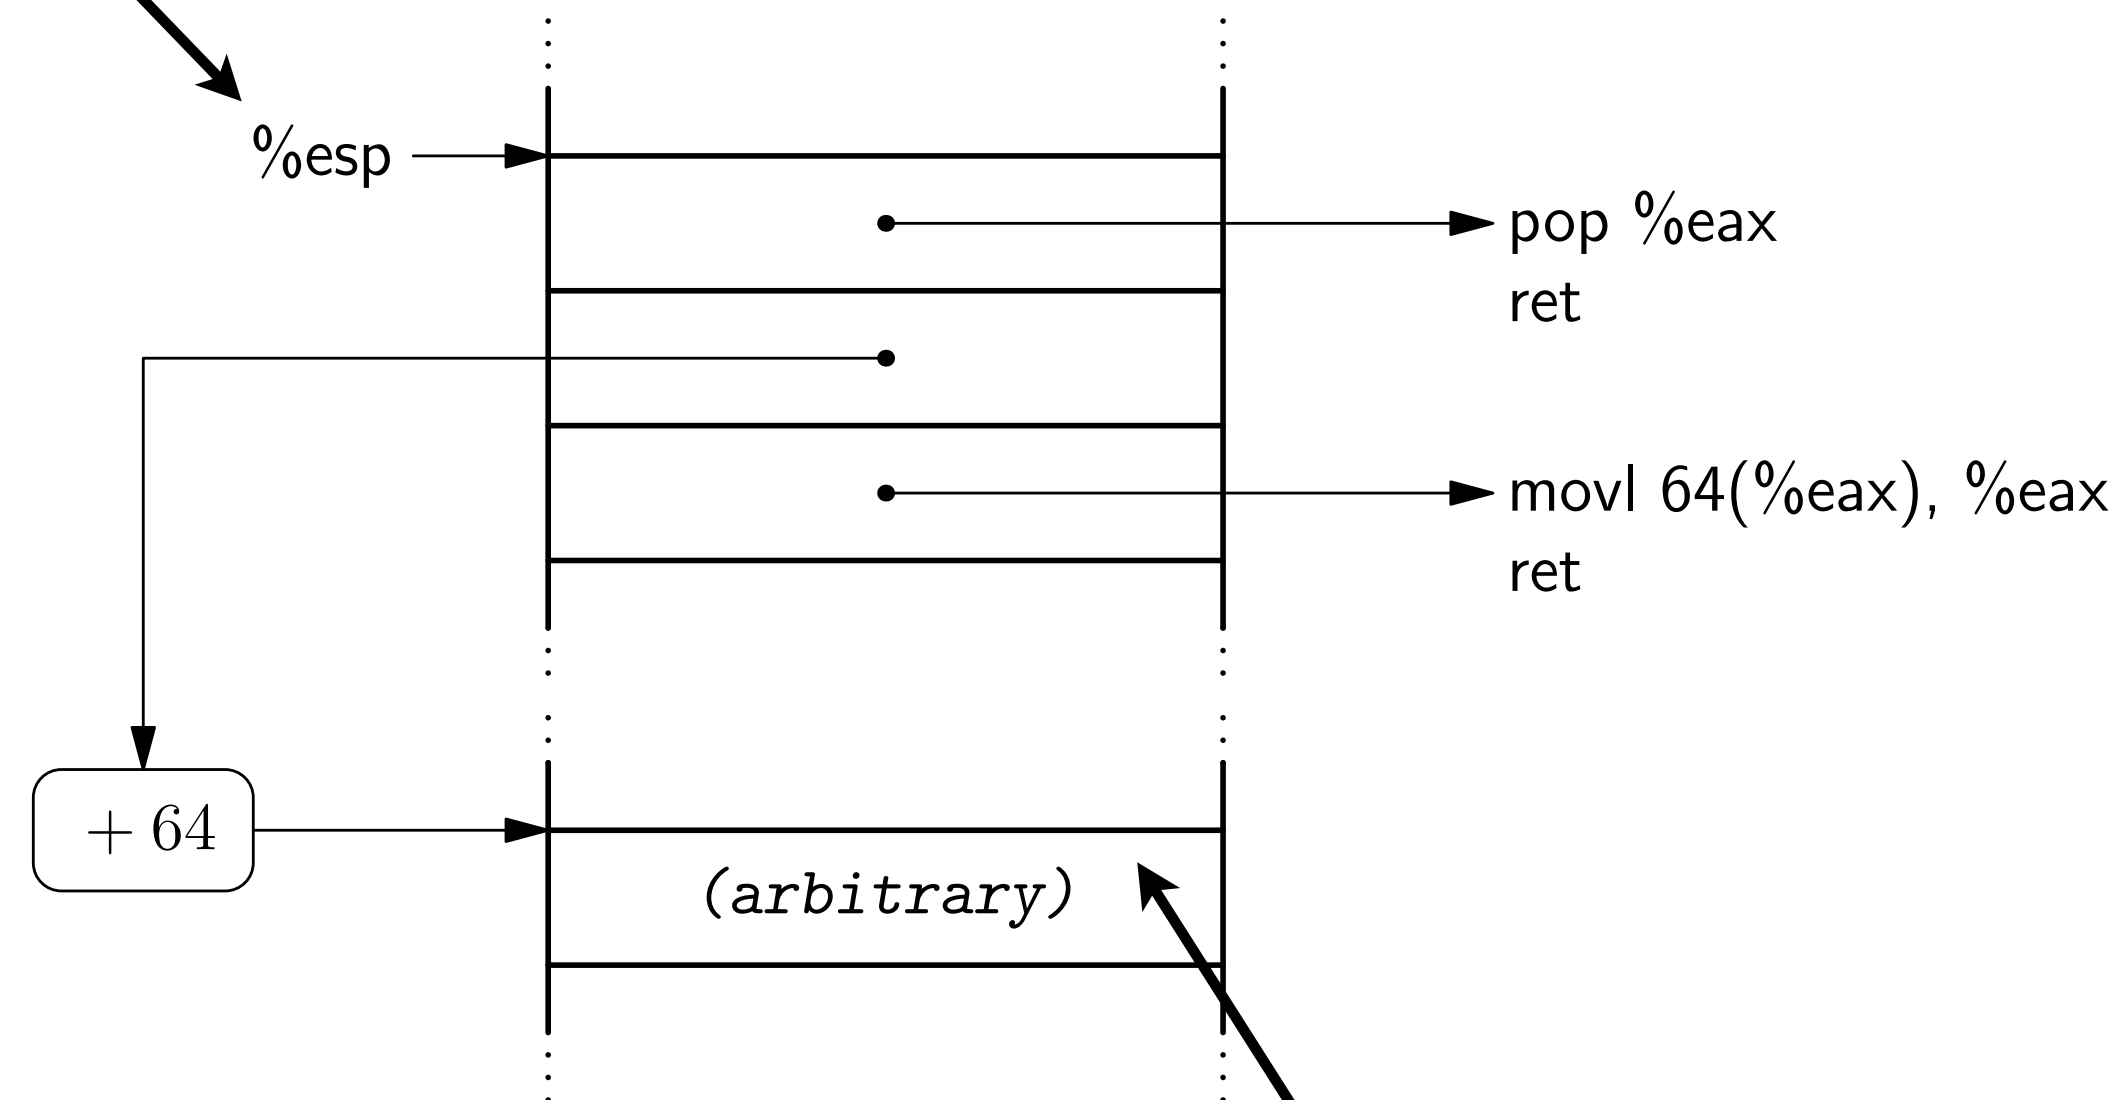
word into %eax



ROP word load

word into %eax

ret issued when SP here



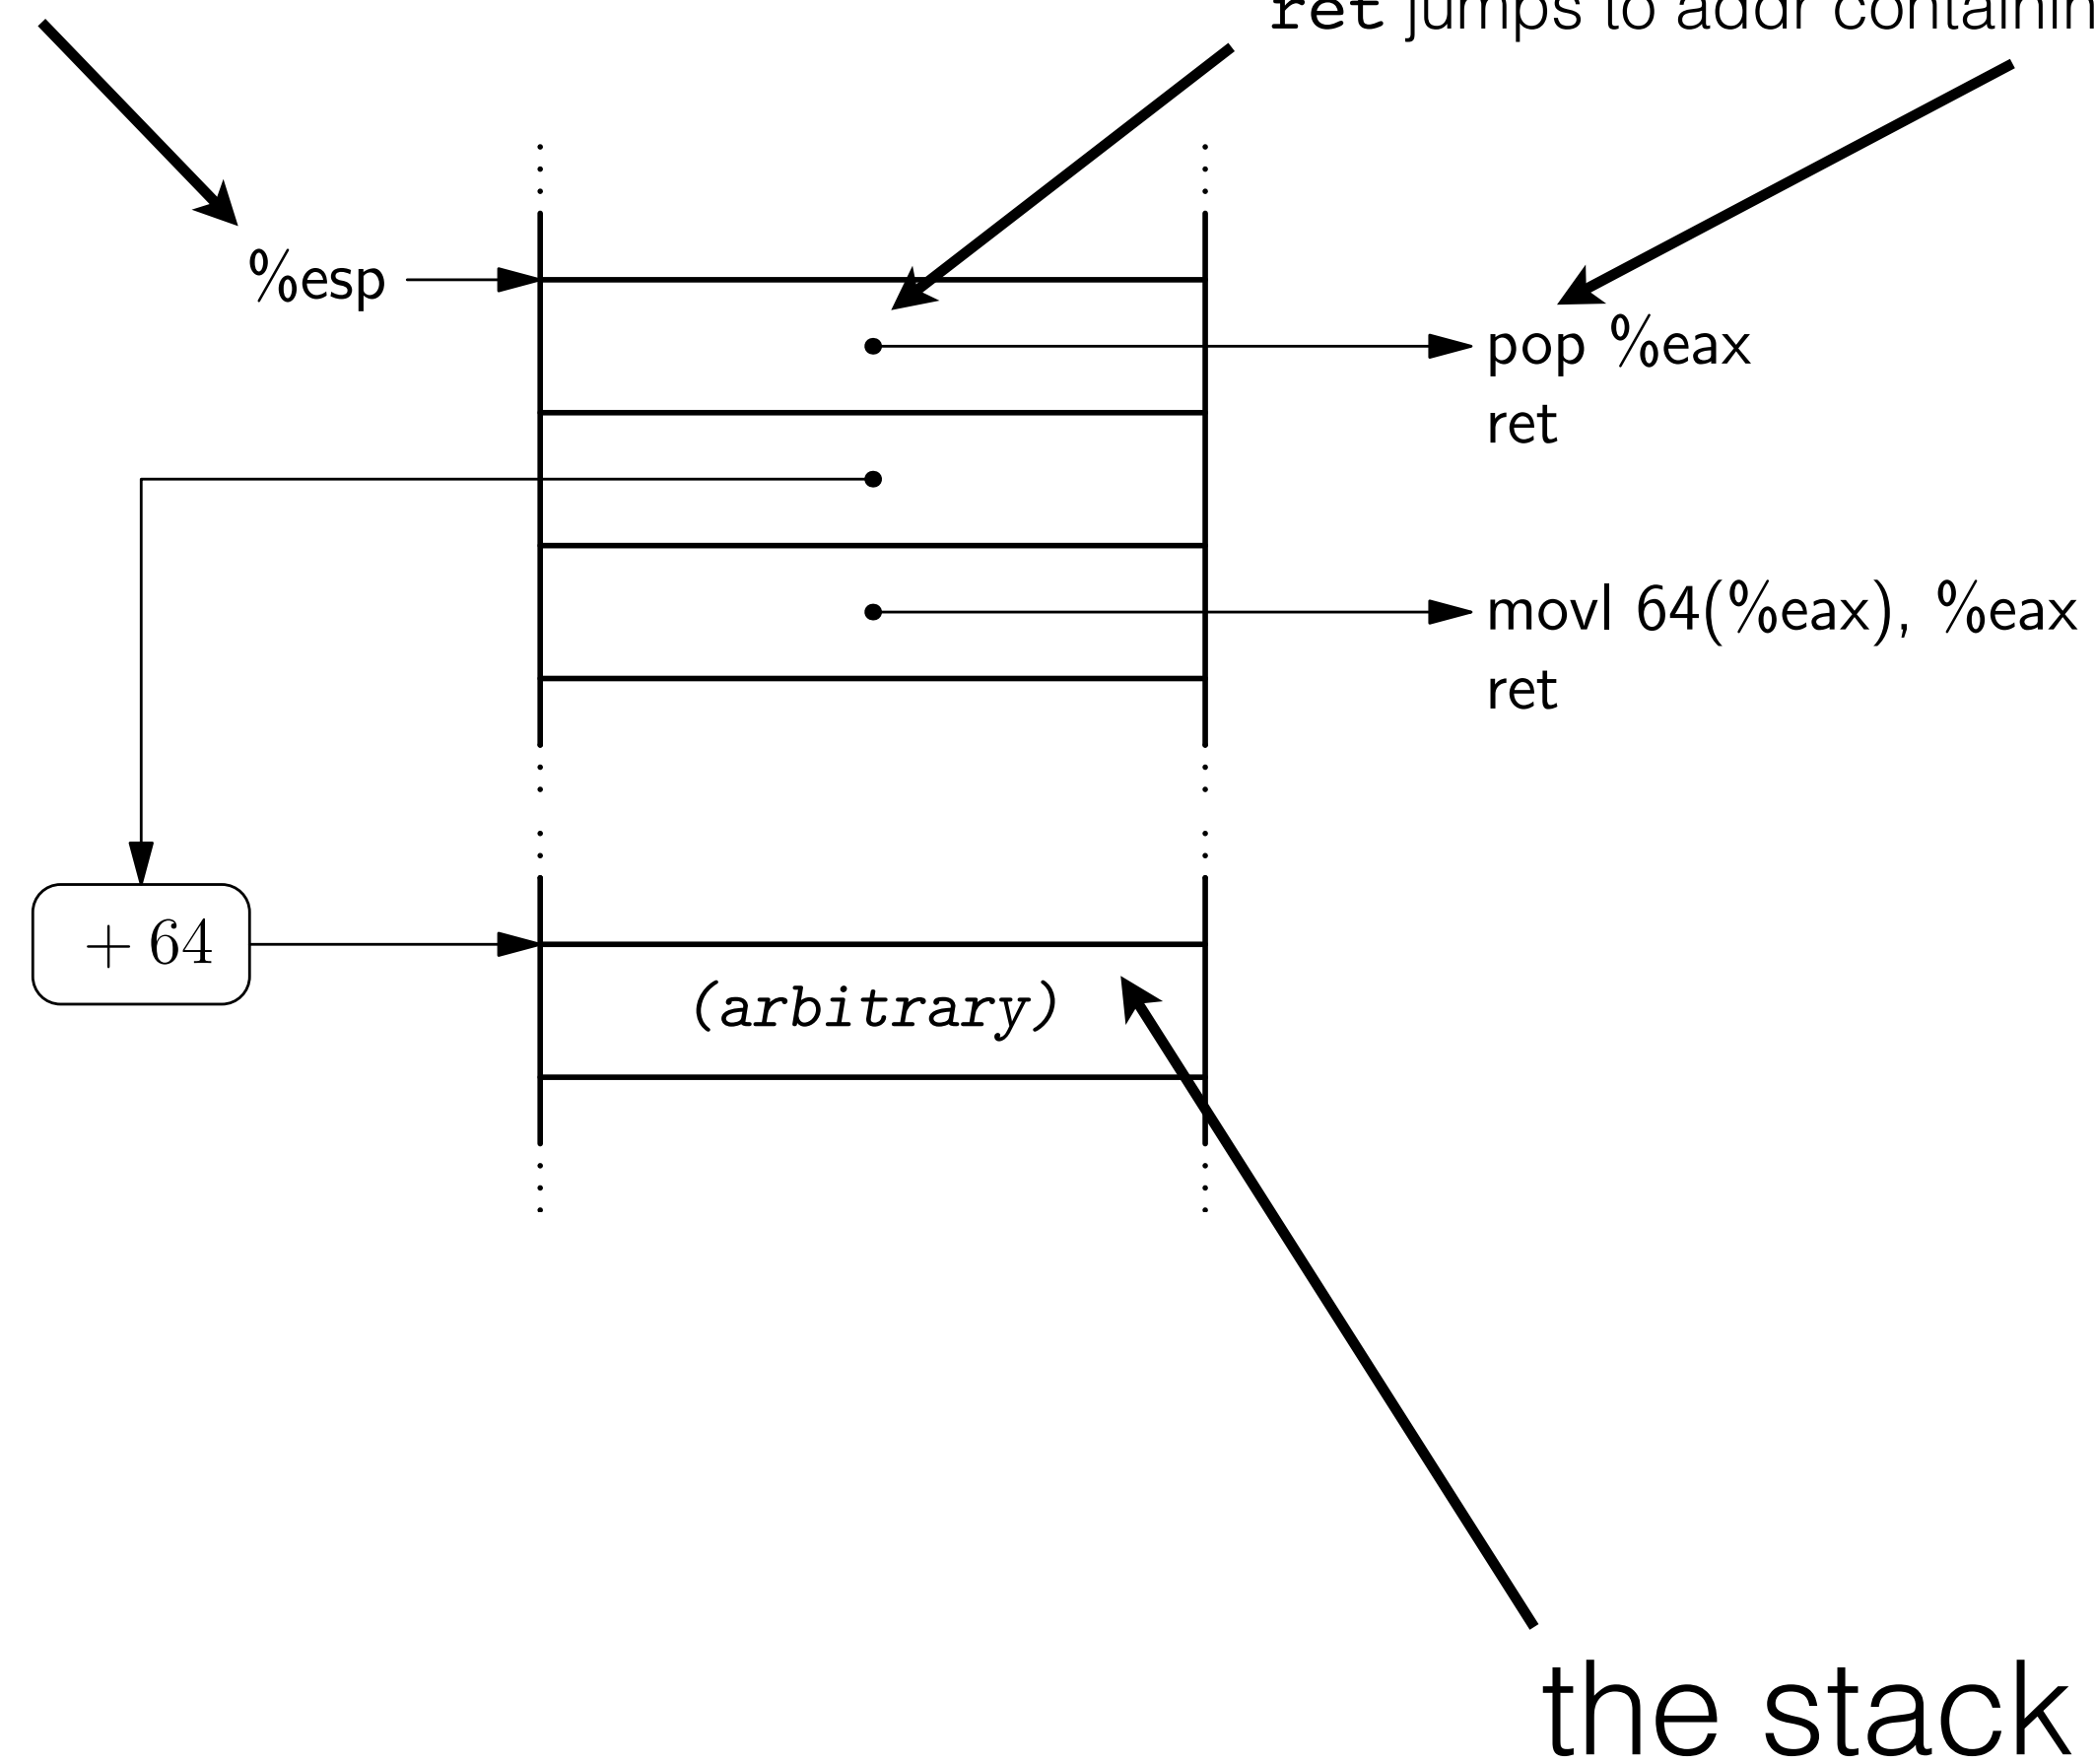
the stack

ROP word load

word into %eax

ret issued when SP here

ret jumps to addr containing



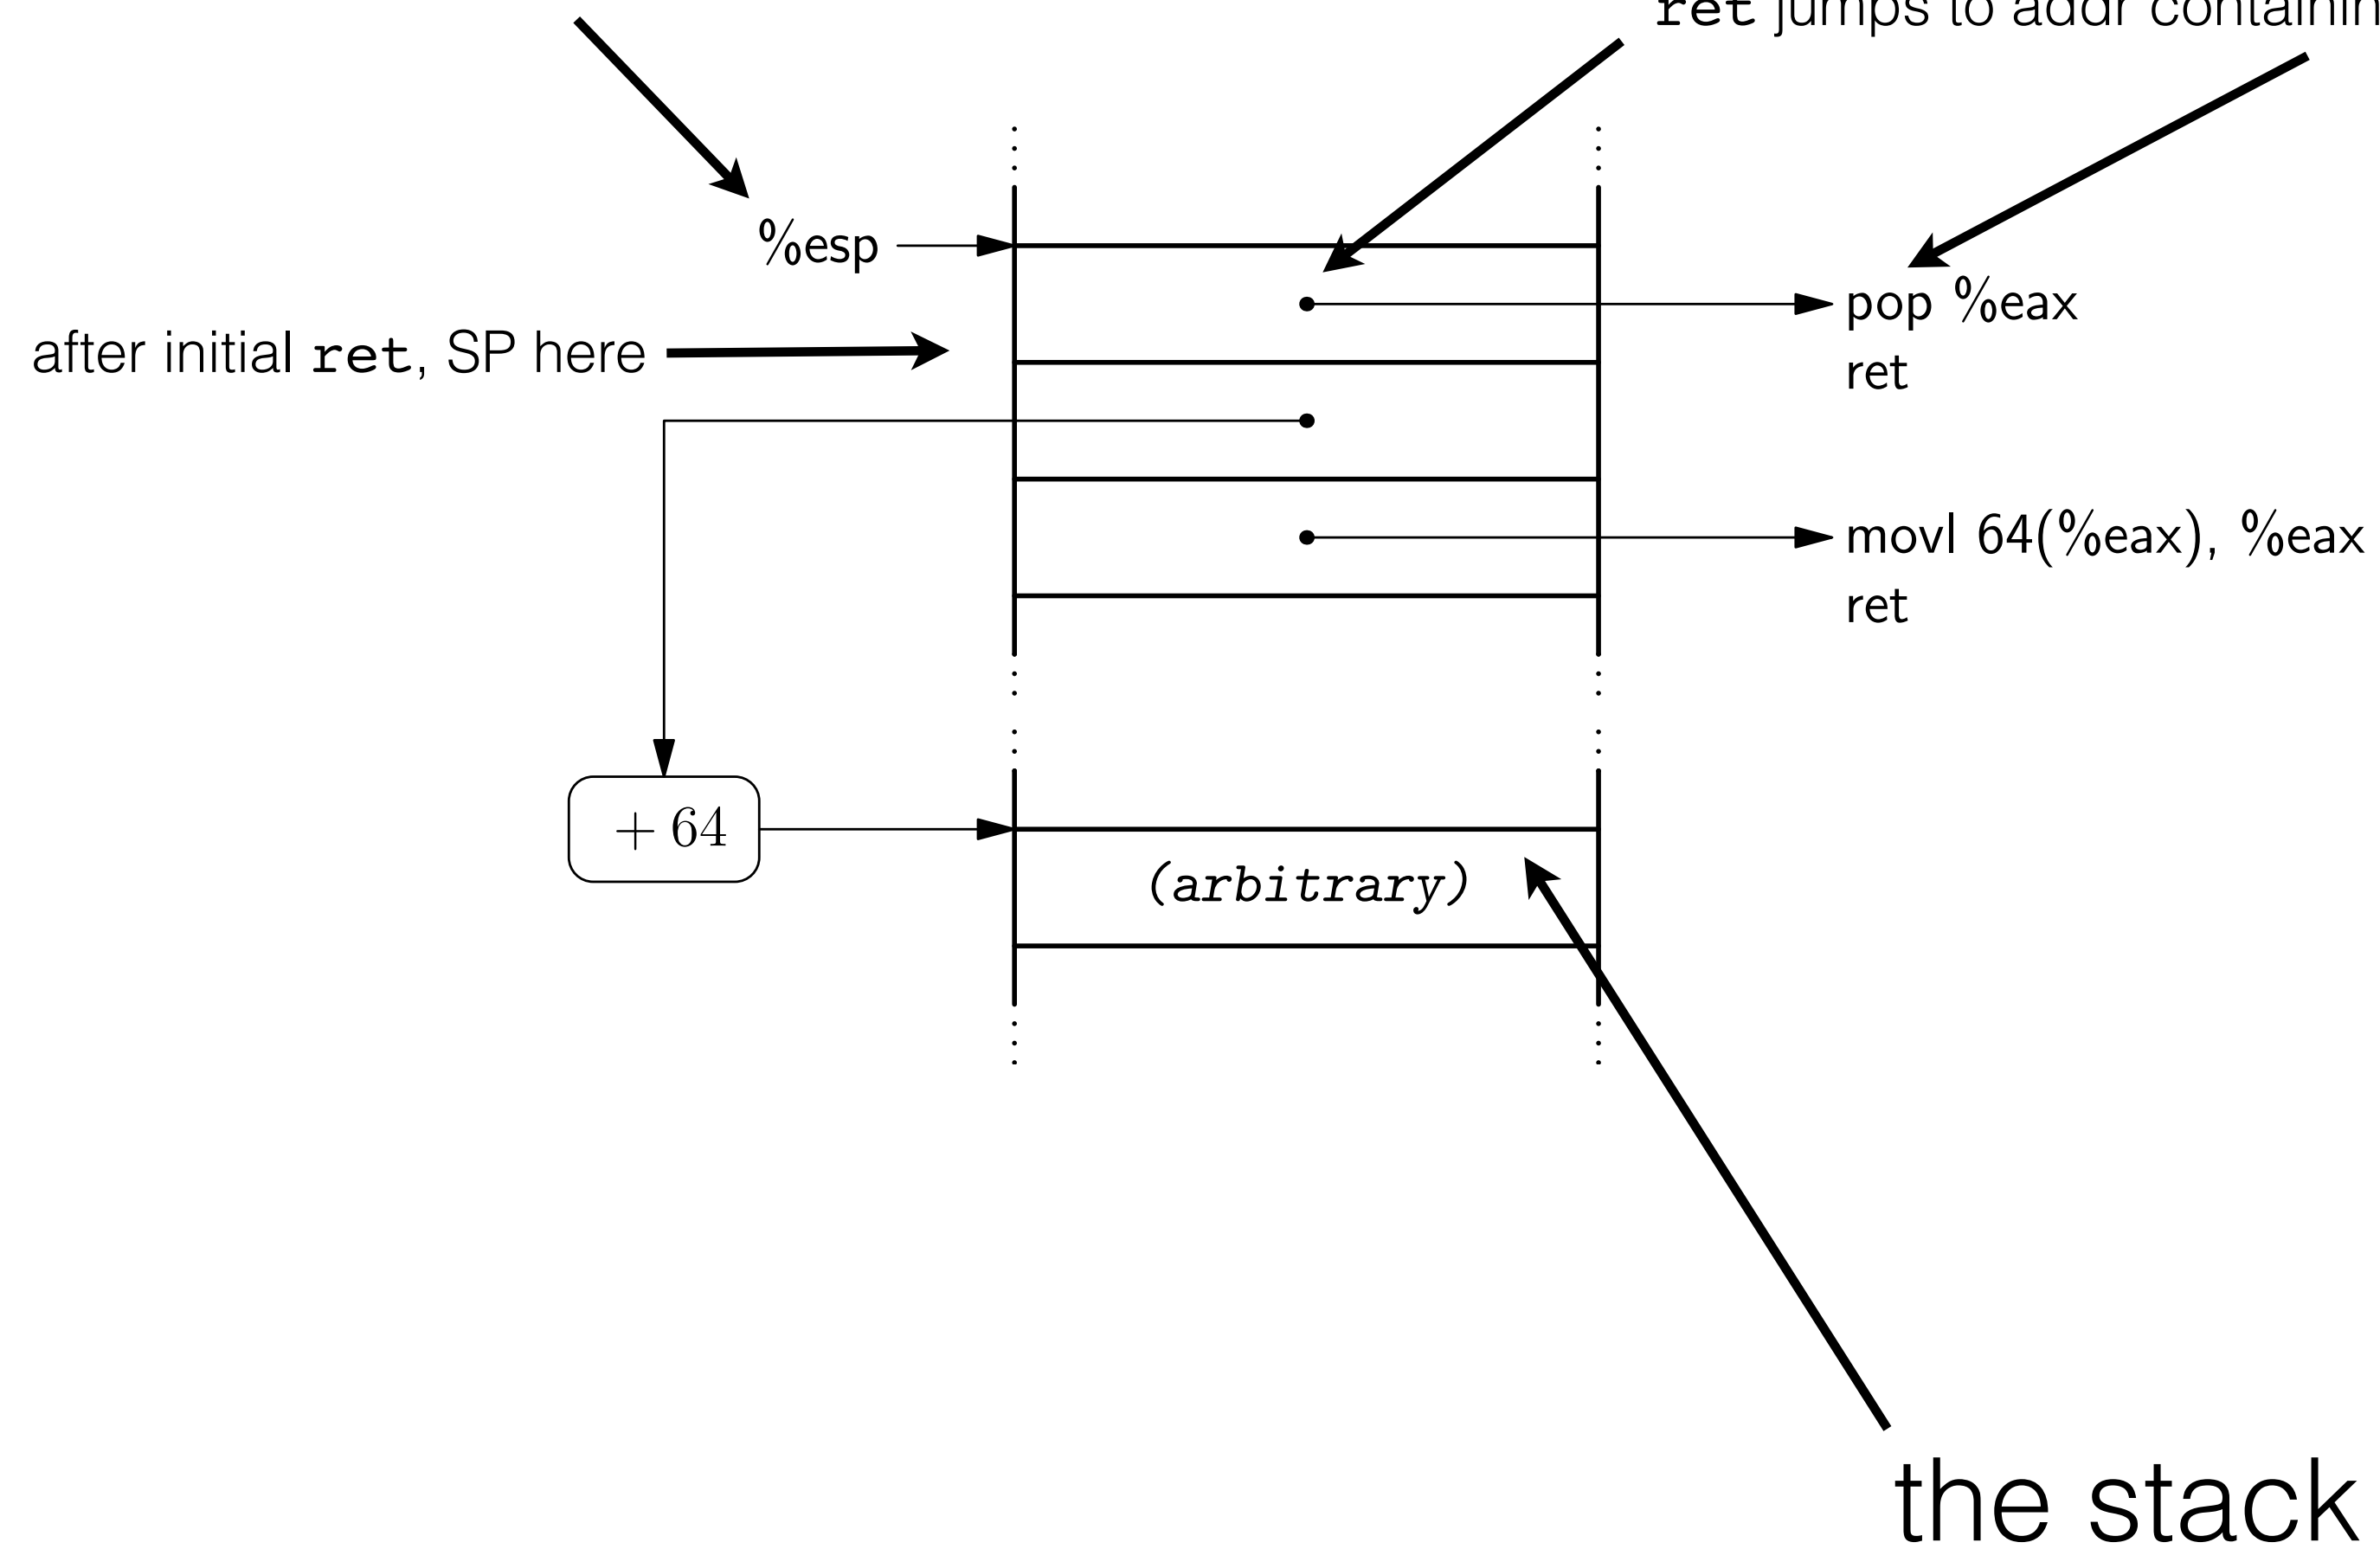
the stack

ROP word load

word into %eax

ret issued when SP here

ret jumps to addr containing



dot represents constructed stack

ROP word load

word into `%eax`

`ret` issued when SP here

`ret` jumps to addr containing

after initial `ret`, SP here

`%esp`

`pop %eax`
`ret`

what is put in `%eax`:
addr of word to load, less 64

`movl 64(%eax), %eax`
`ret`

+ 64

(arbitrary)

the stack

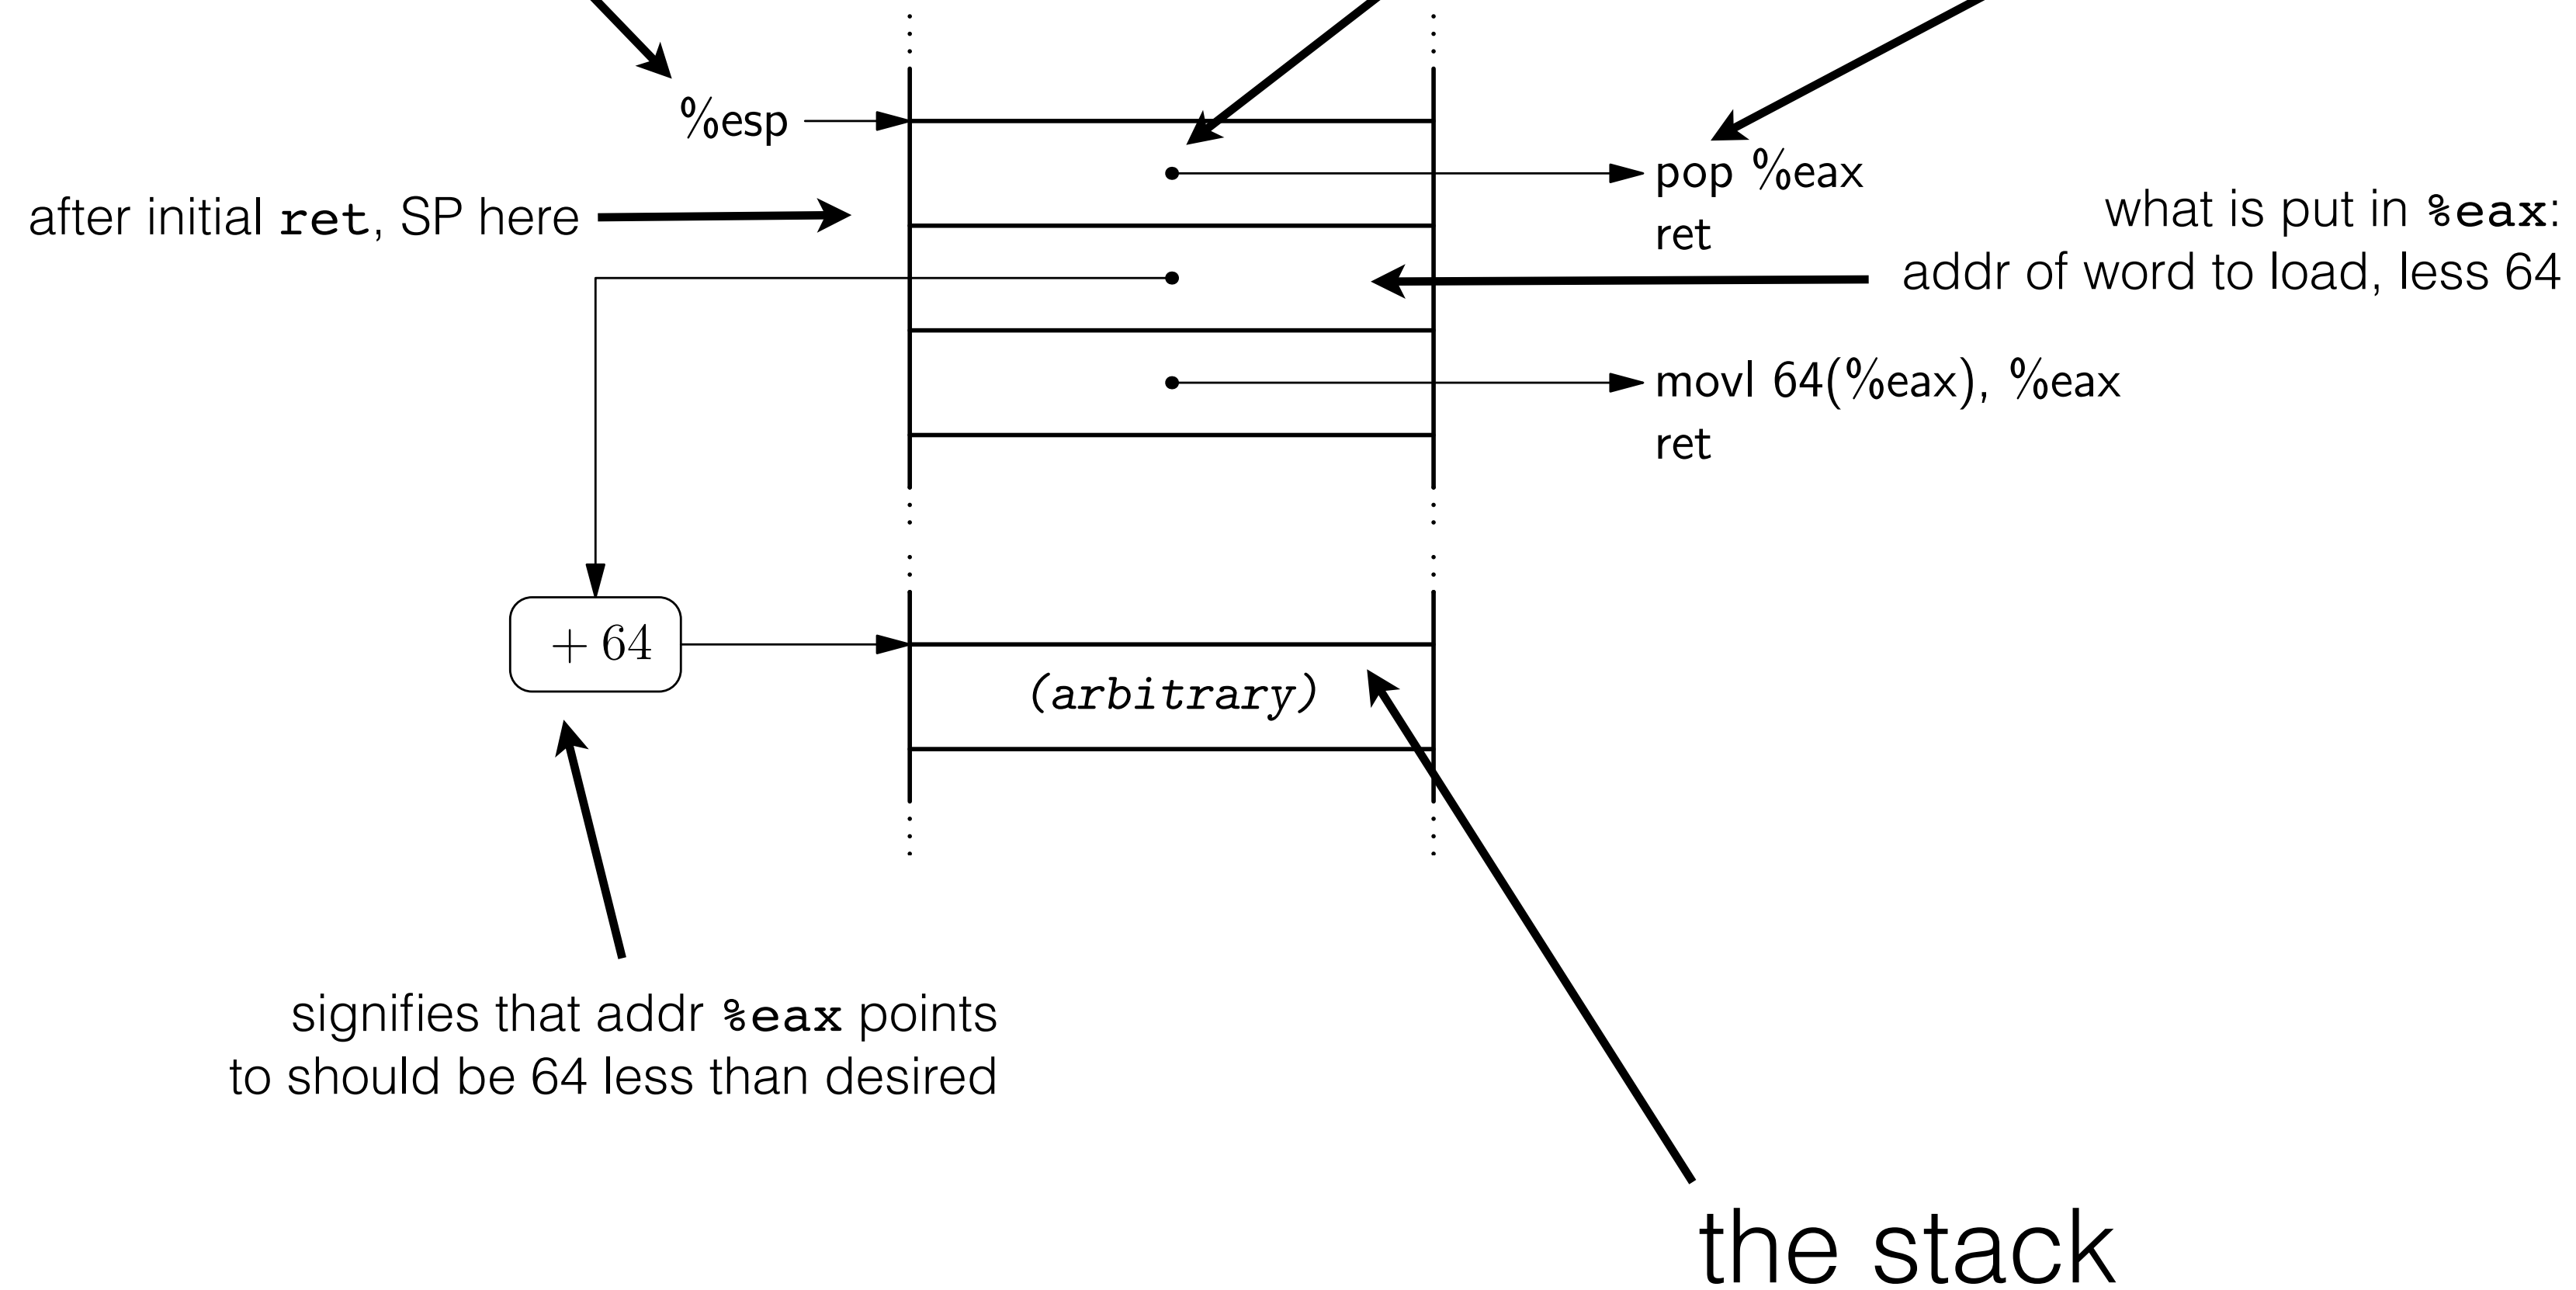
dot represents constructed stack

ROP word load

word into `%eax`

`ret` issued when SP here

`ret` jumps to addr containing

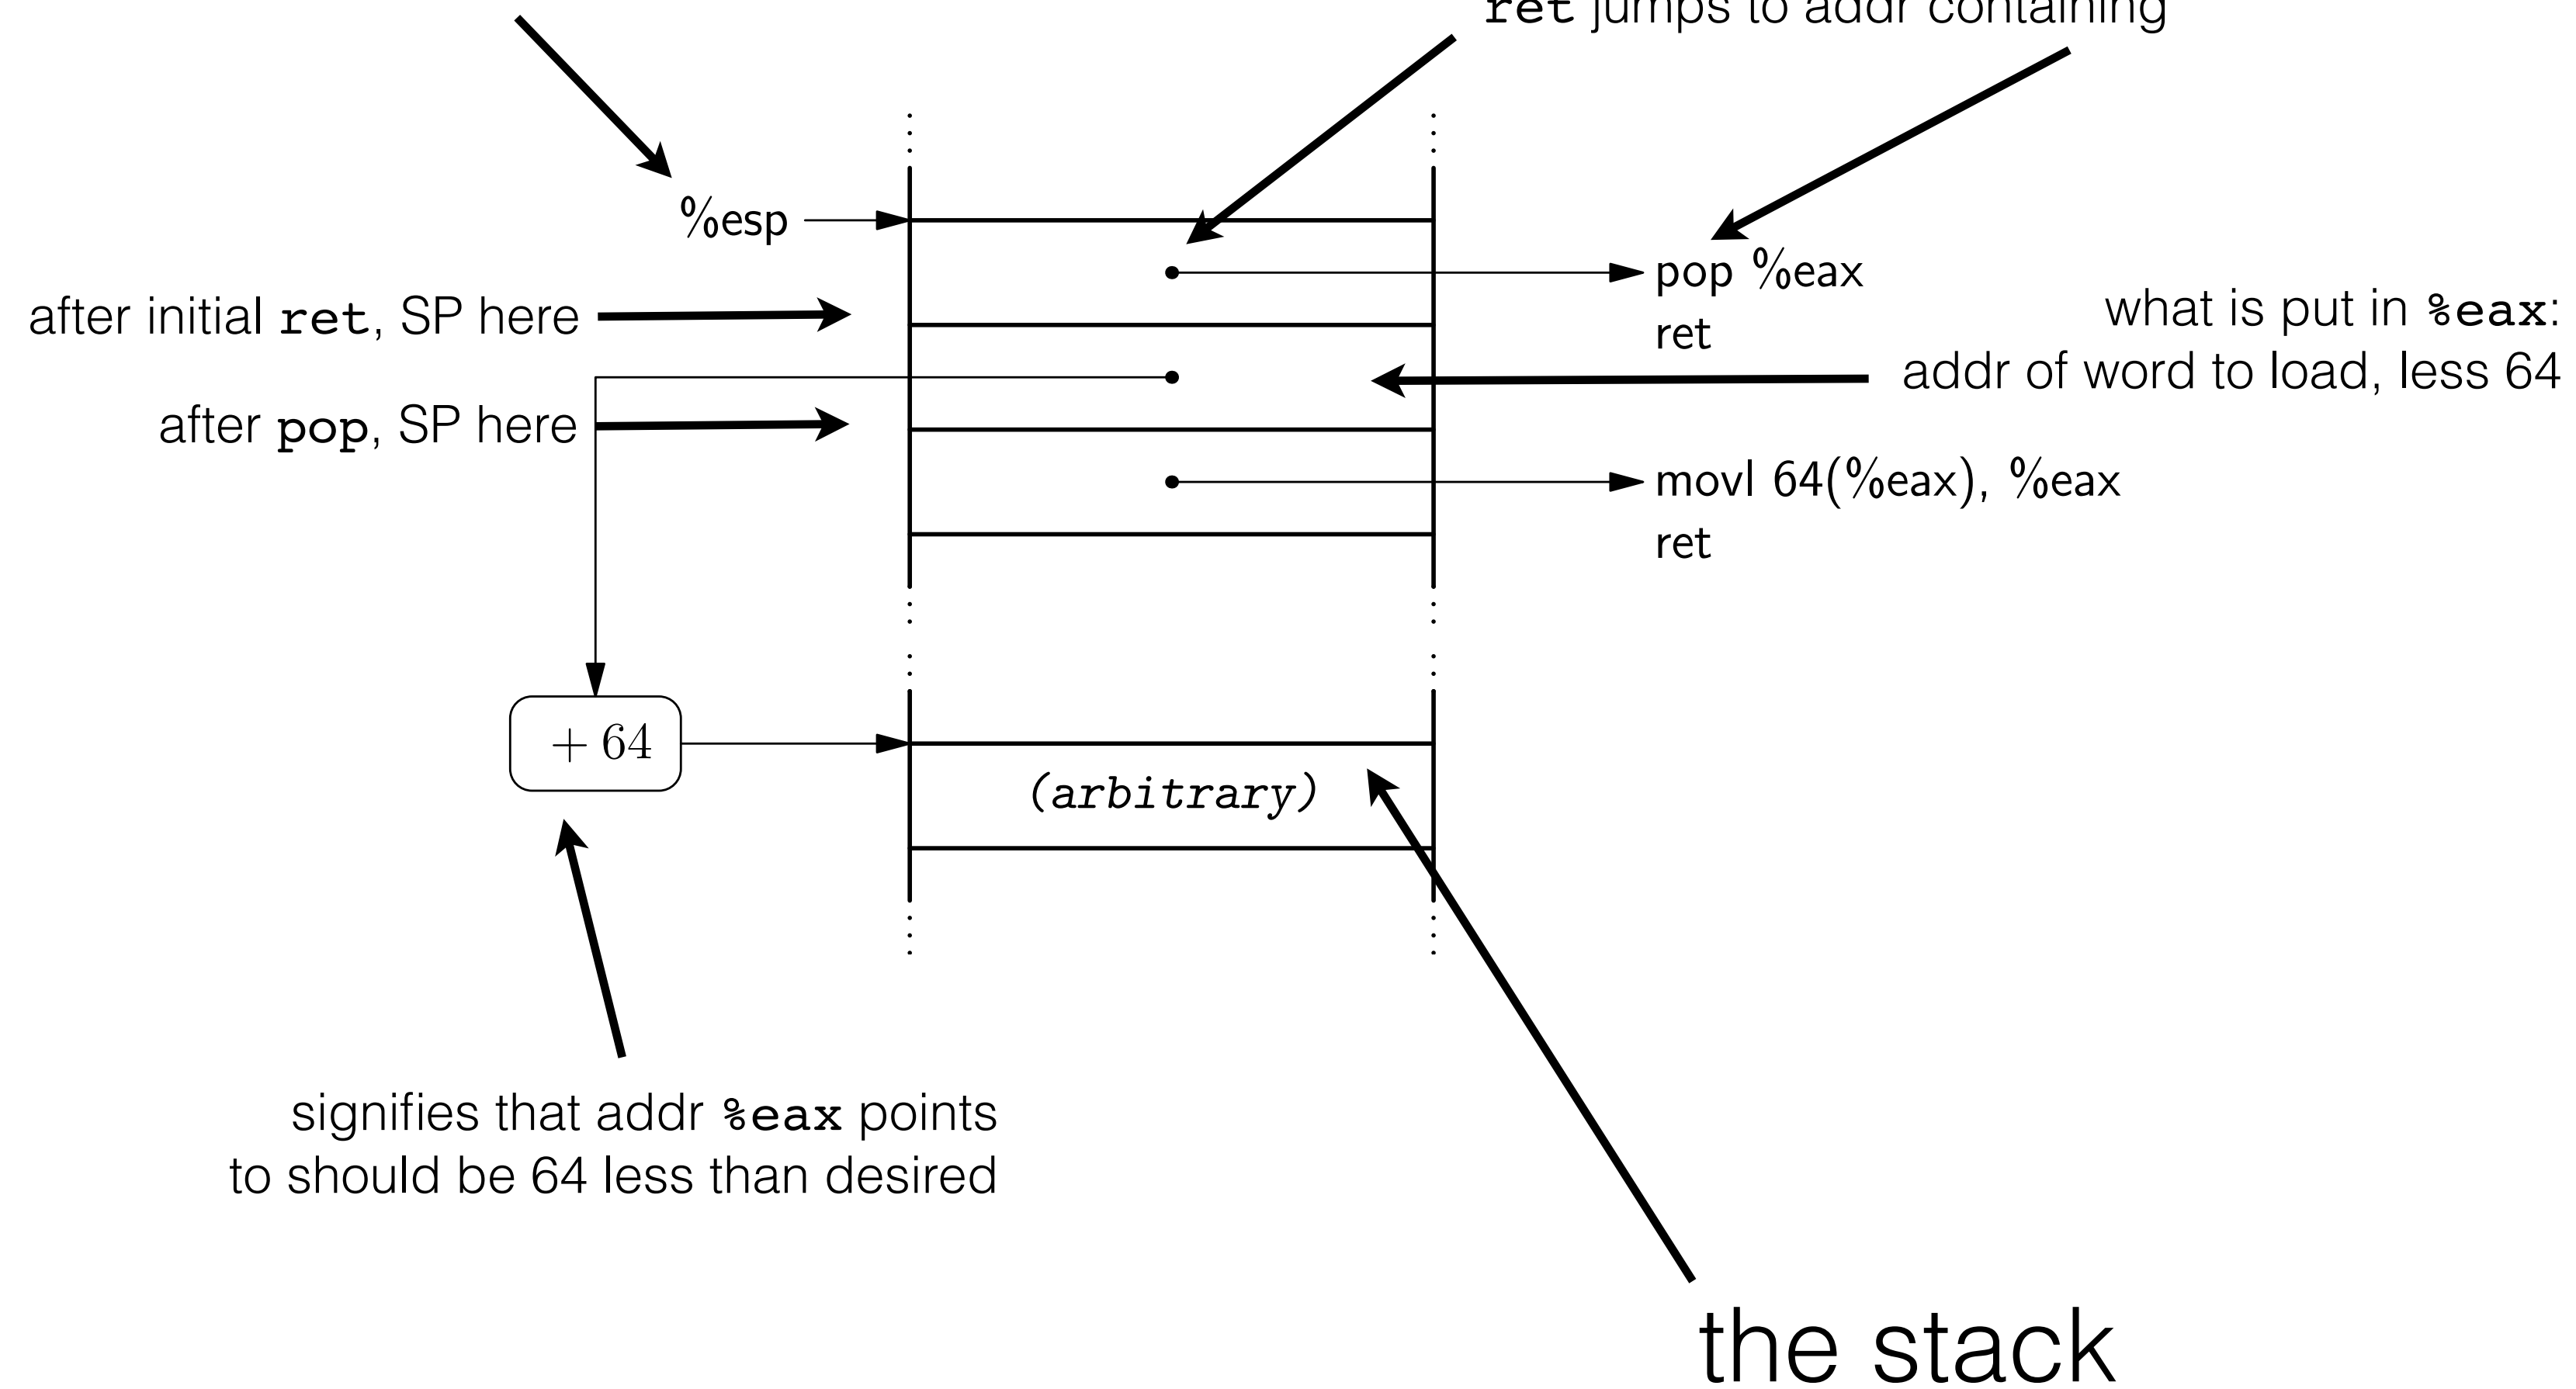


ROP word load

word into `%eax`

`ret` issued when SP here

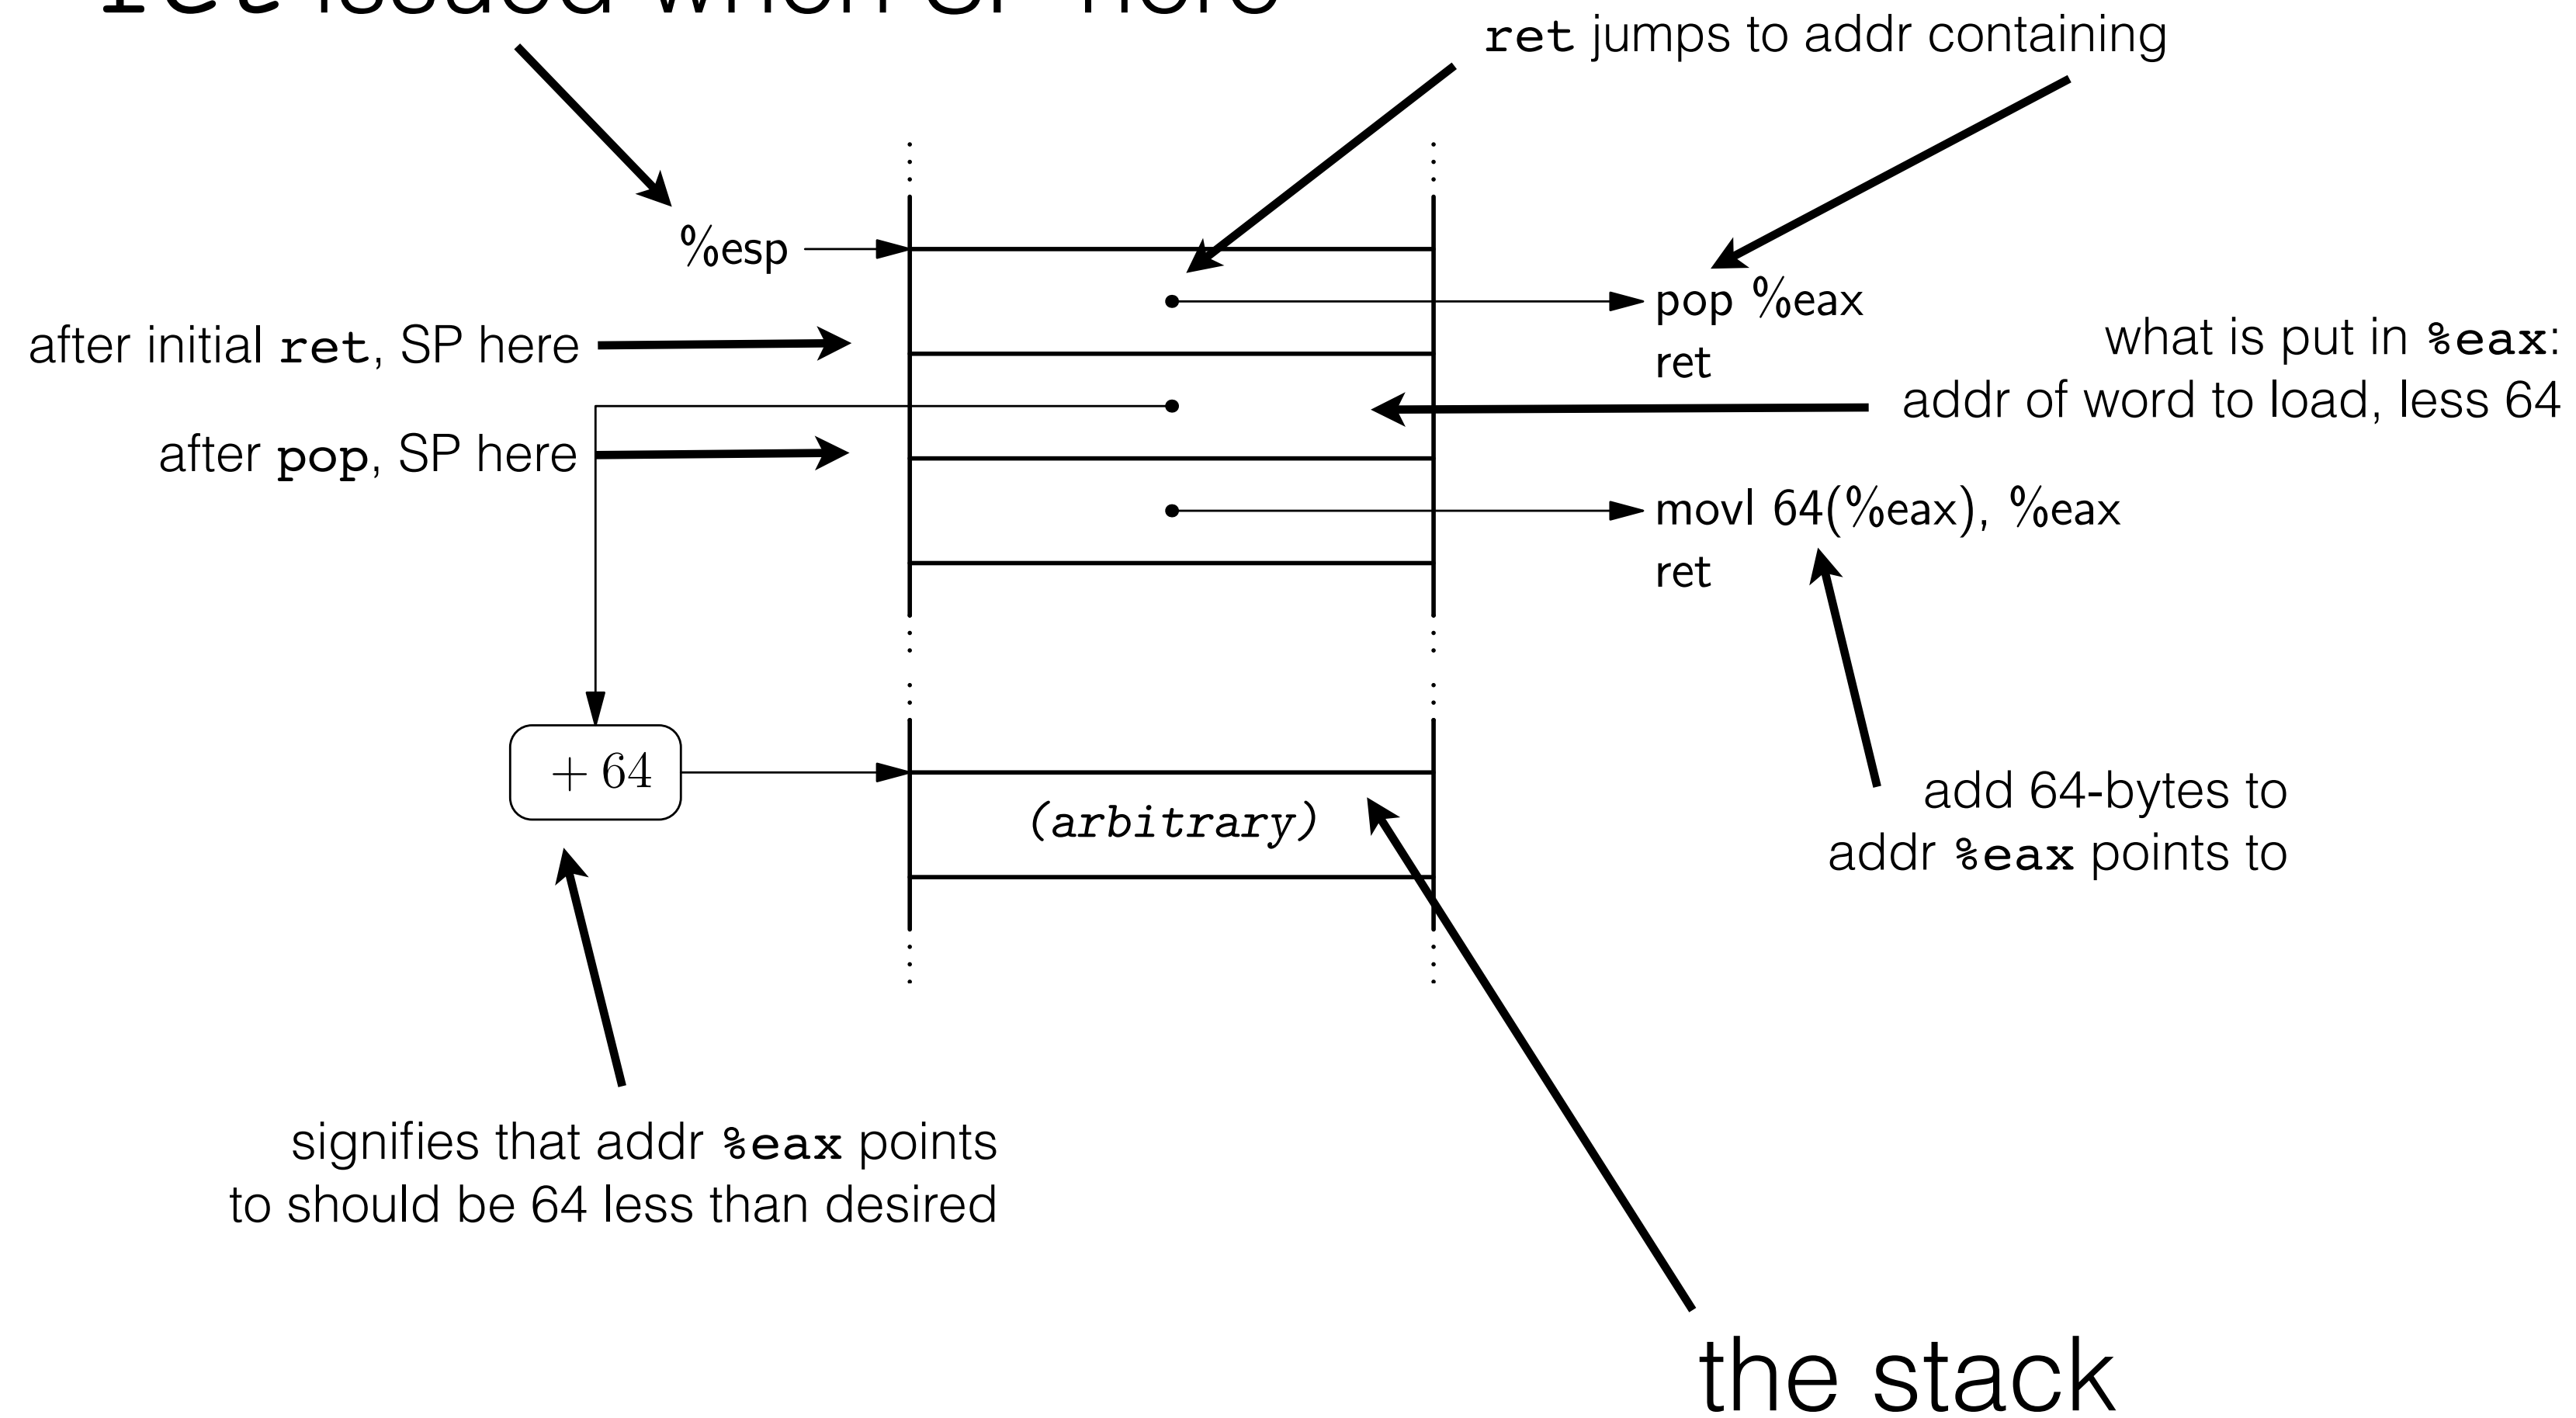
`ret` jumps to addr containing



ROP word load

word into `%eax`

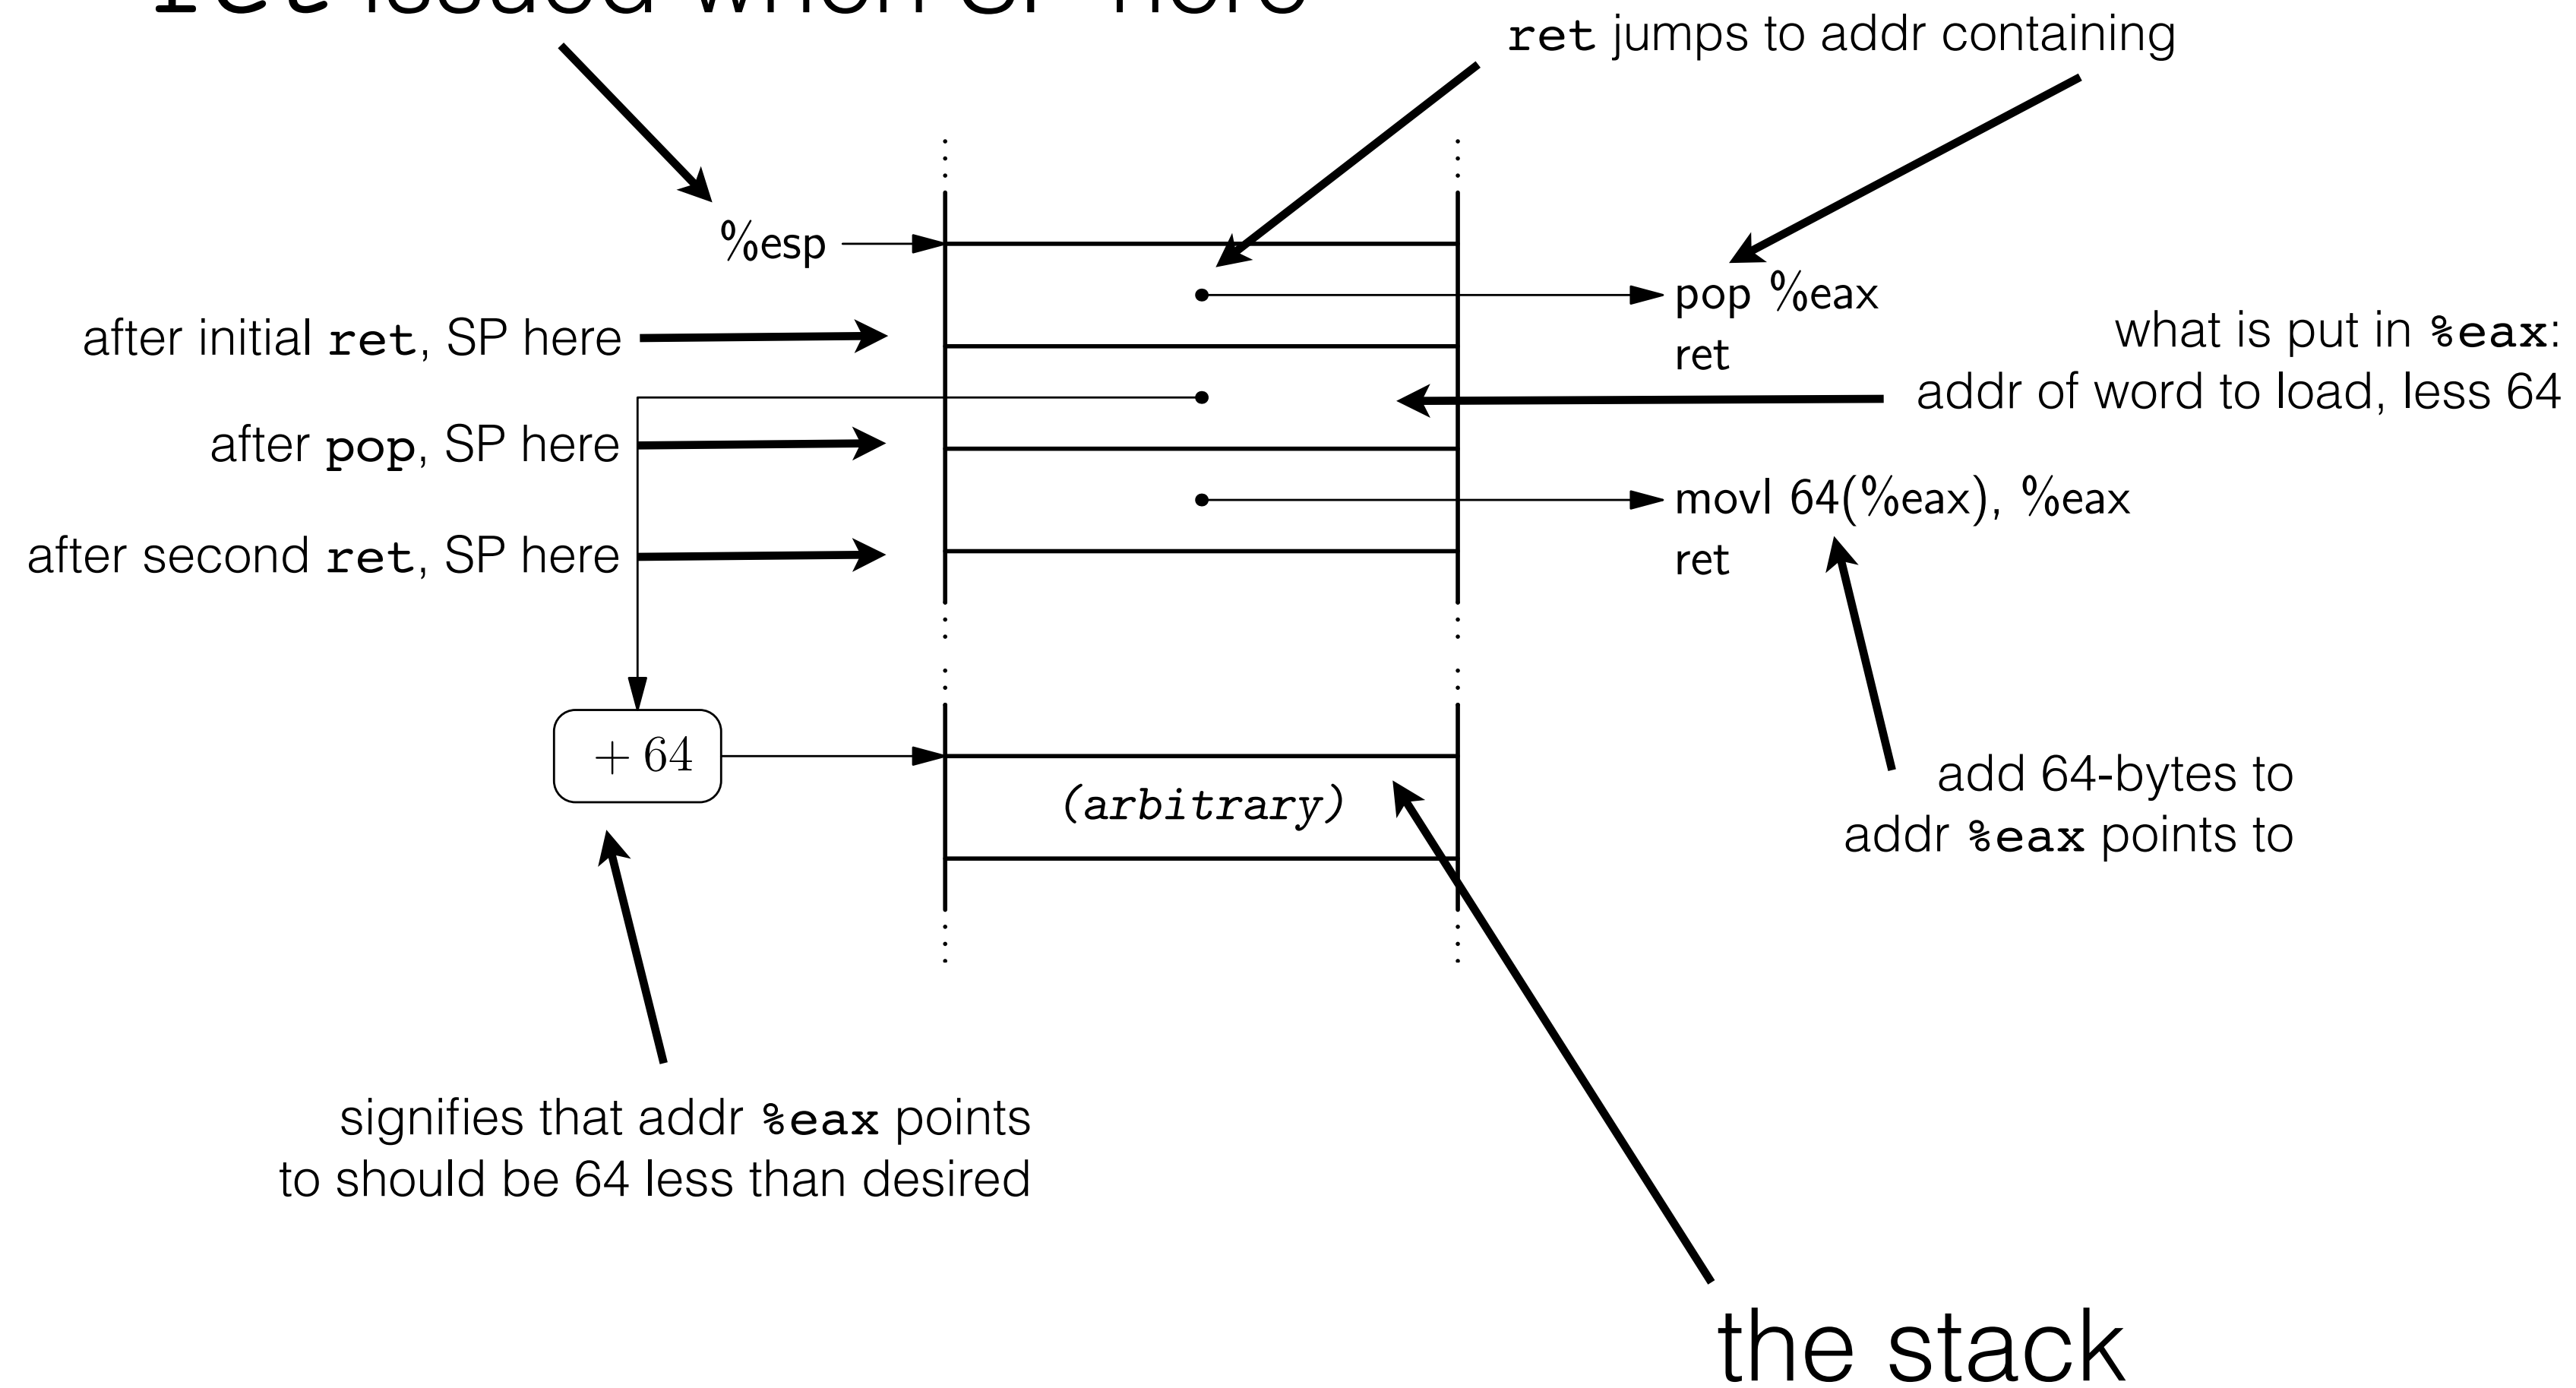
`ret` issued when SP here



ROP word load

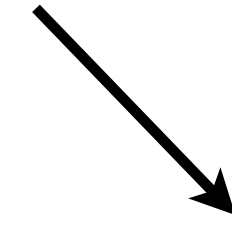
word into `%eax`

`ret` issued when SP here



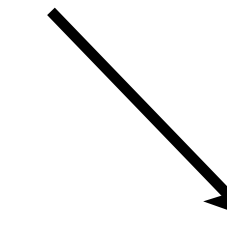
ROP word store

ROP word store

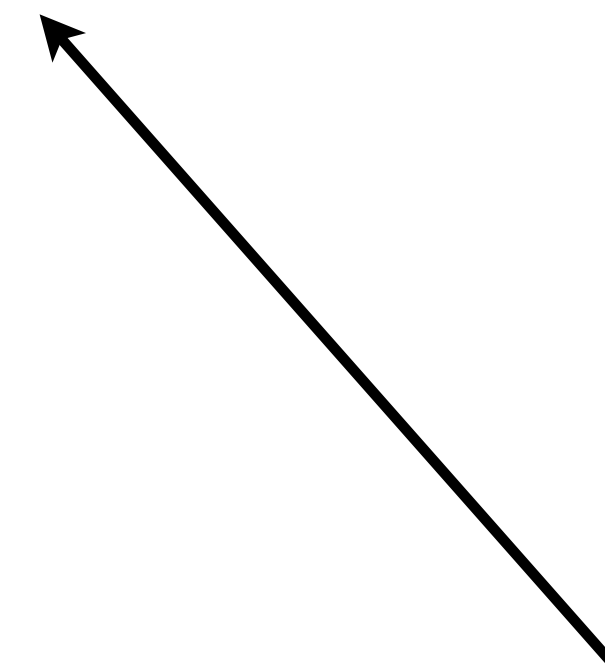
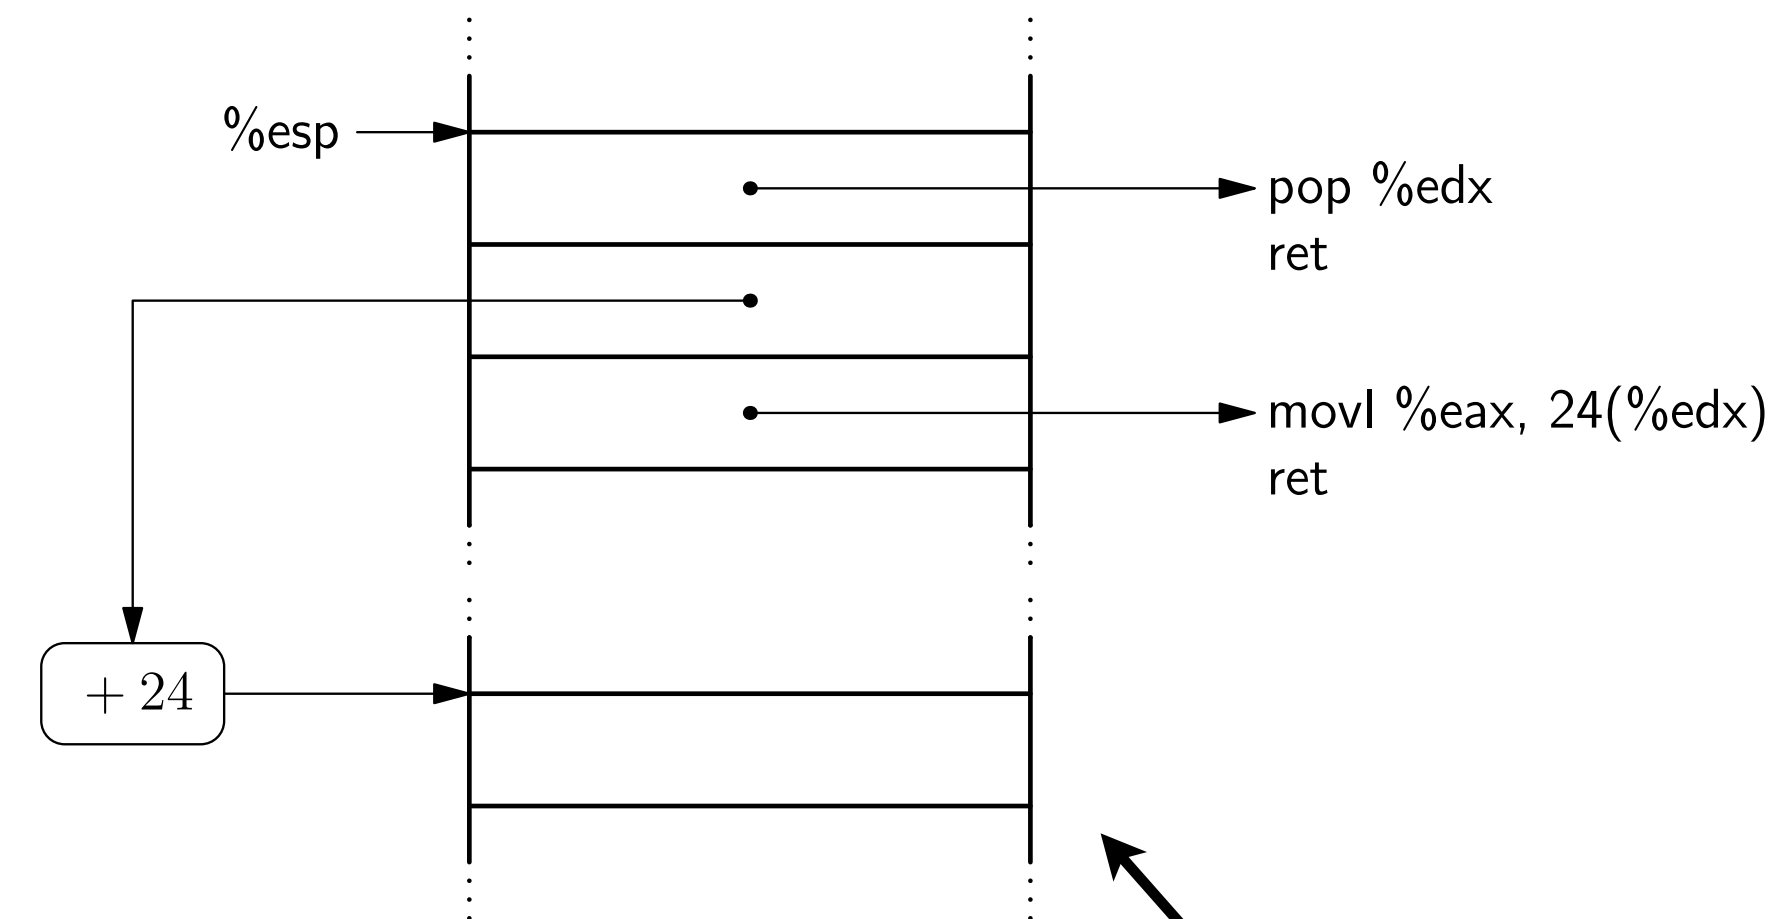


word in `%eax` to
memory

ROP word store

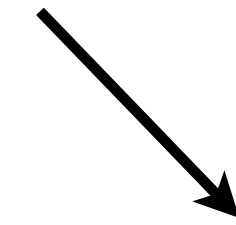


word in `%eax` to
memory



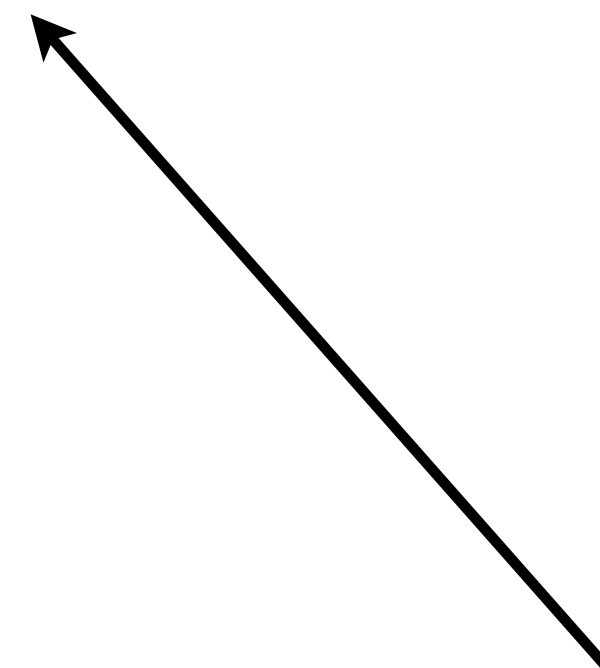
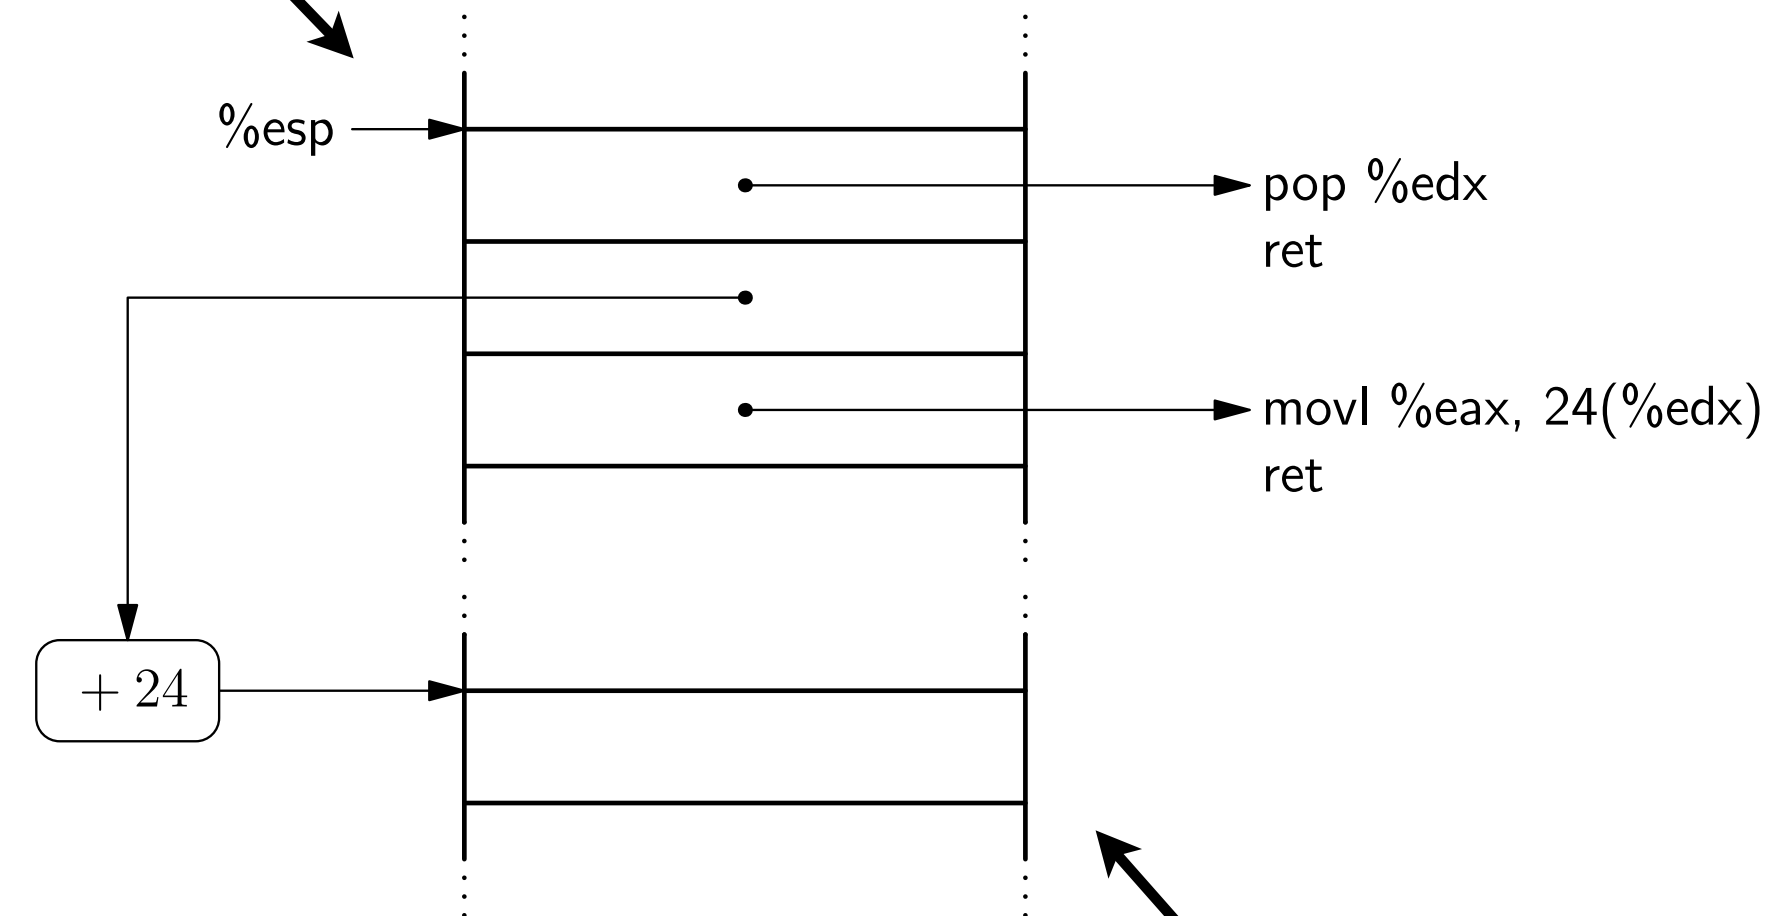
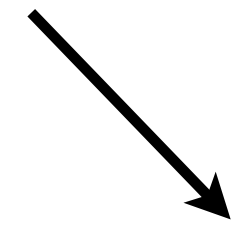
the stack

ROP word store



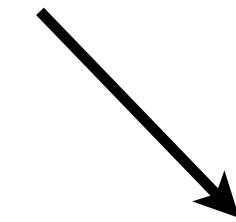
word in `%eax` to
memory

`ret` issued when SP here



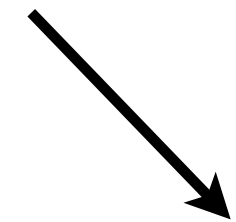
the stack

ROP word store

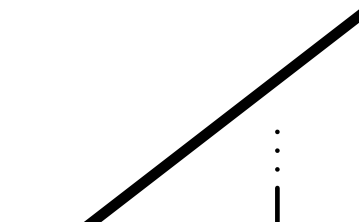


word in `%eax` to
memory

`ret` issued when SP here



`ret` jumps to addr containing



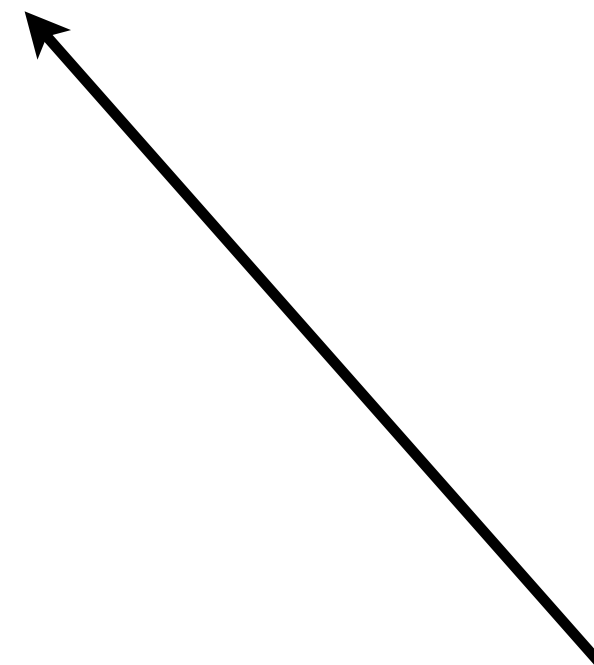
`pop %edx`
`ret`

`movl %eax, 24(%edx)`
`ret`

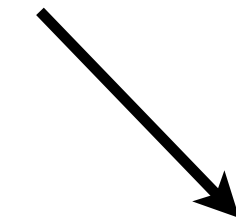
+ 24

`%esp`

the stack



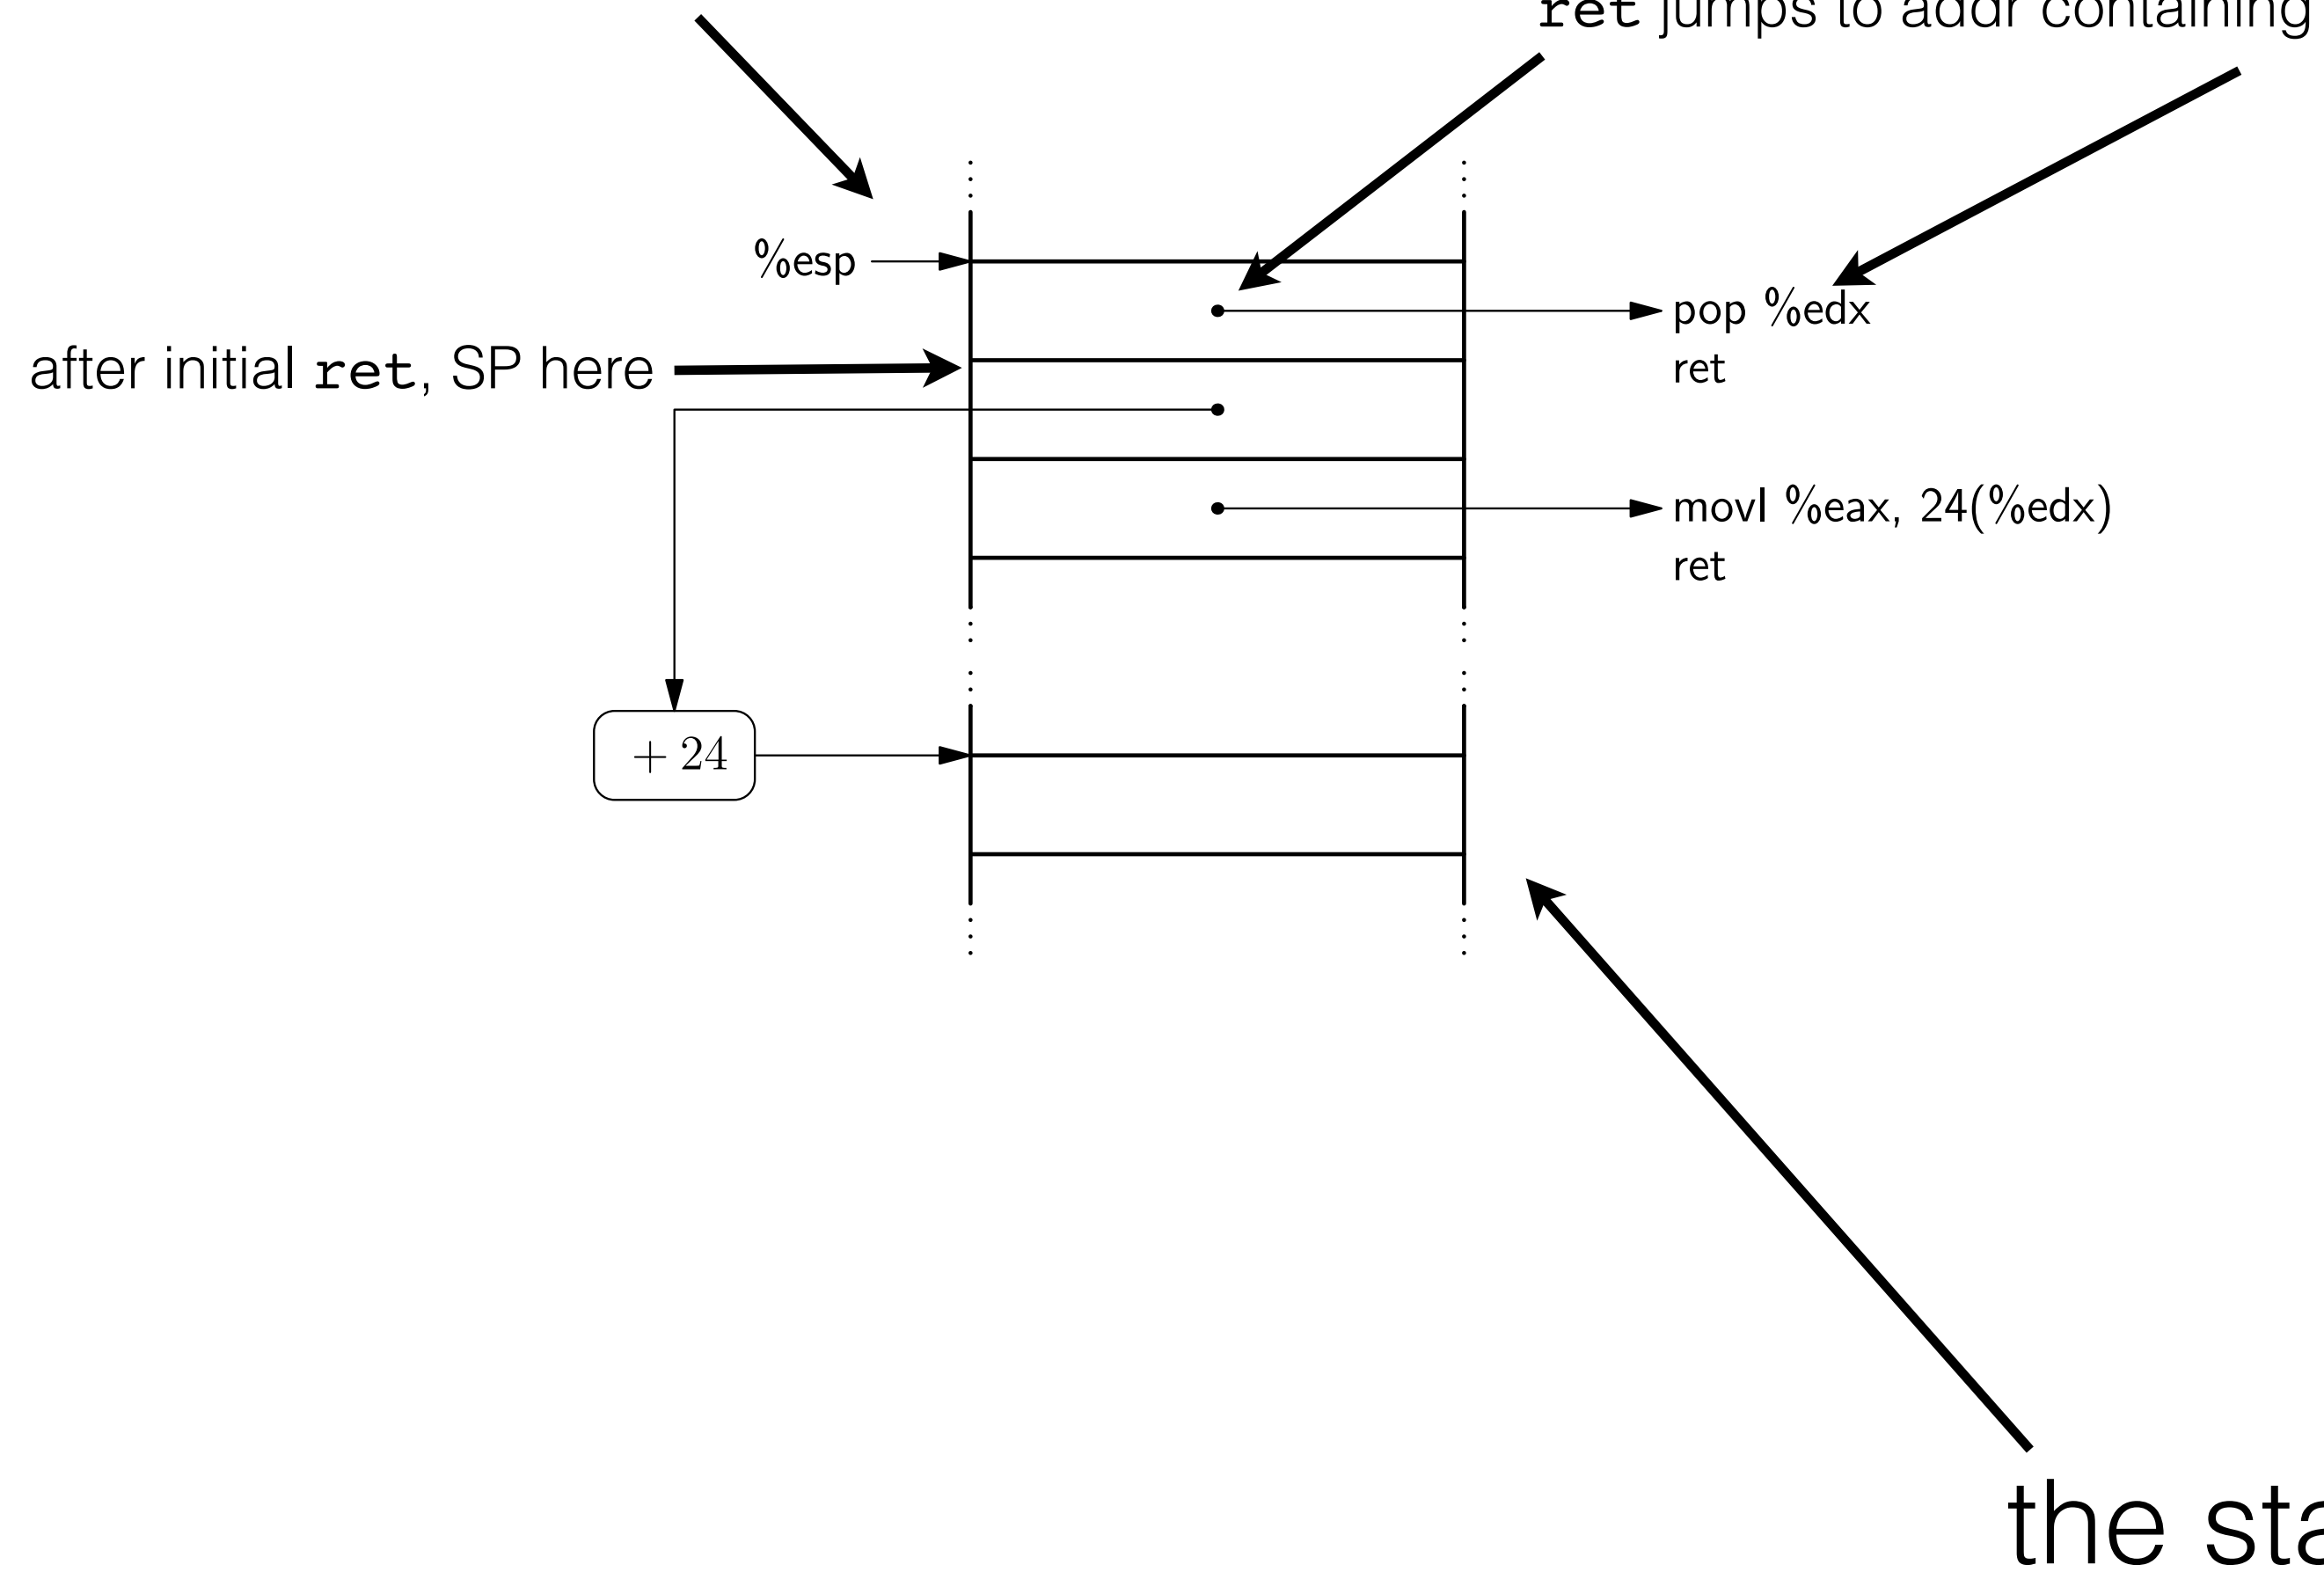
ROP word store



word in `%eax` to
memory

`ret` issued when SP here

`ret` jumps to addr containing

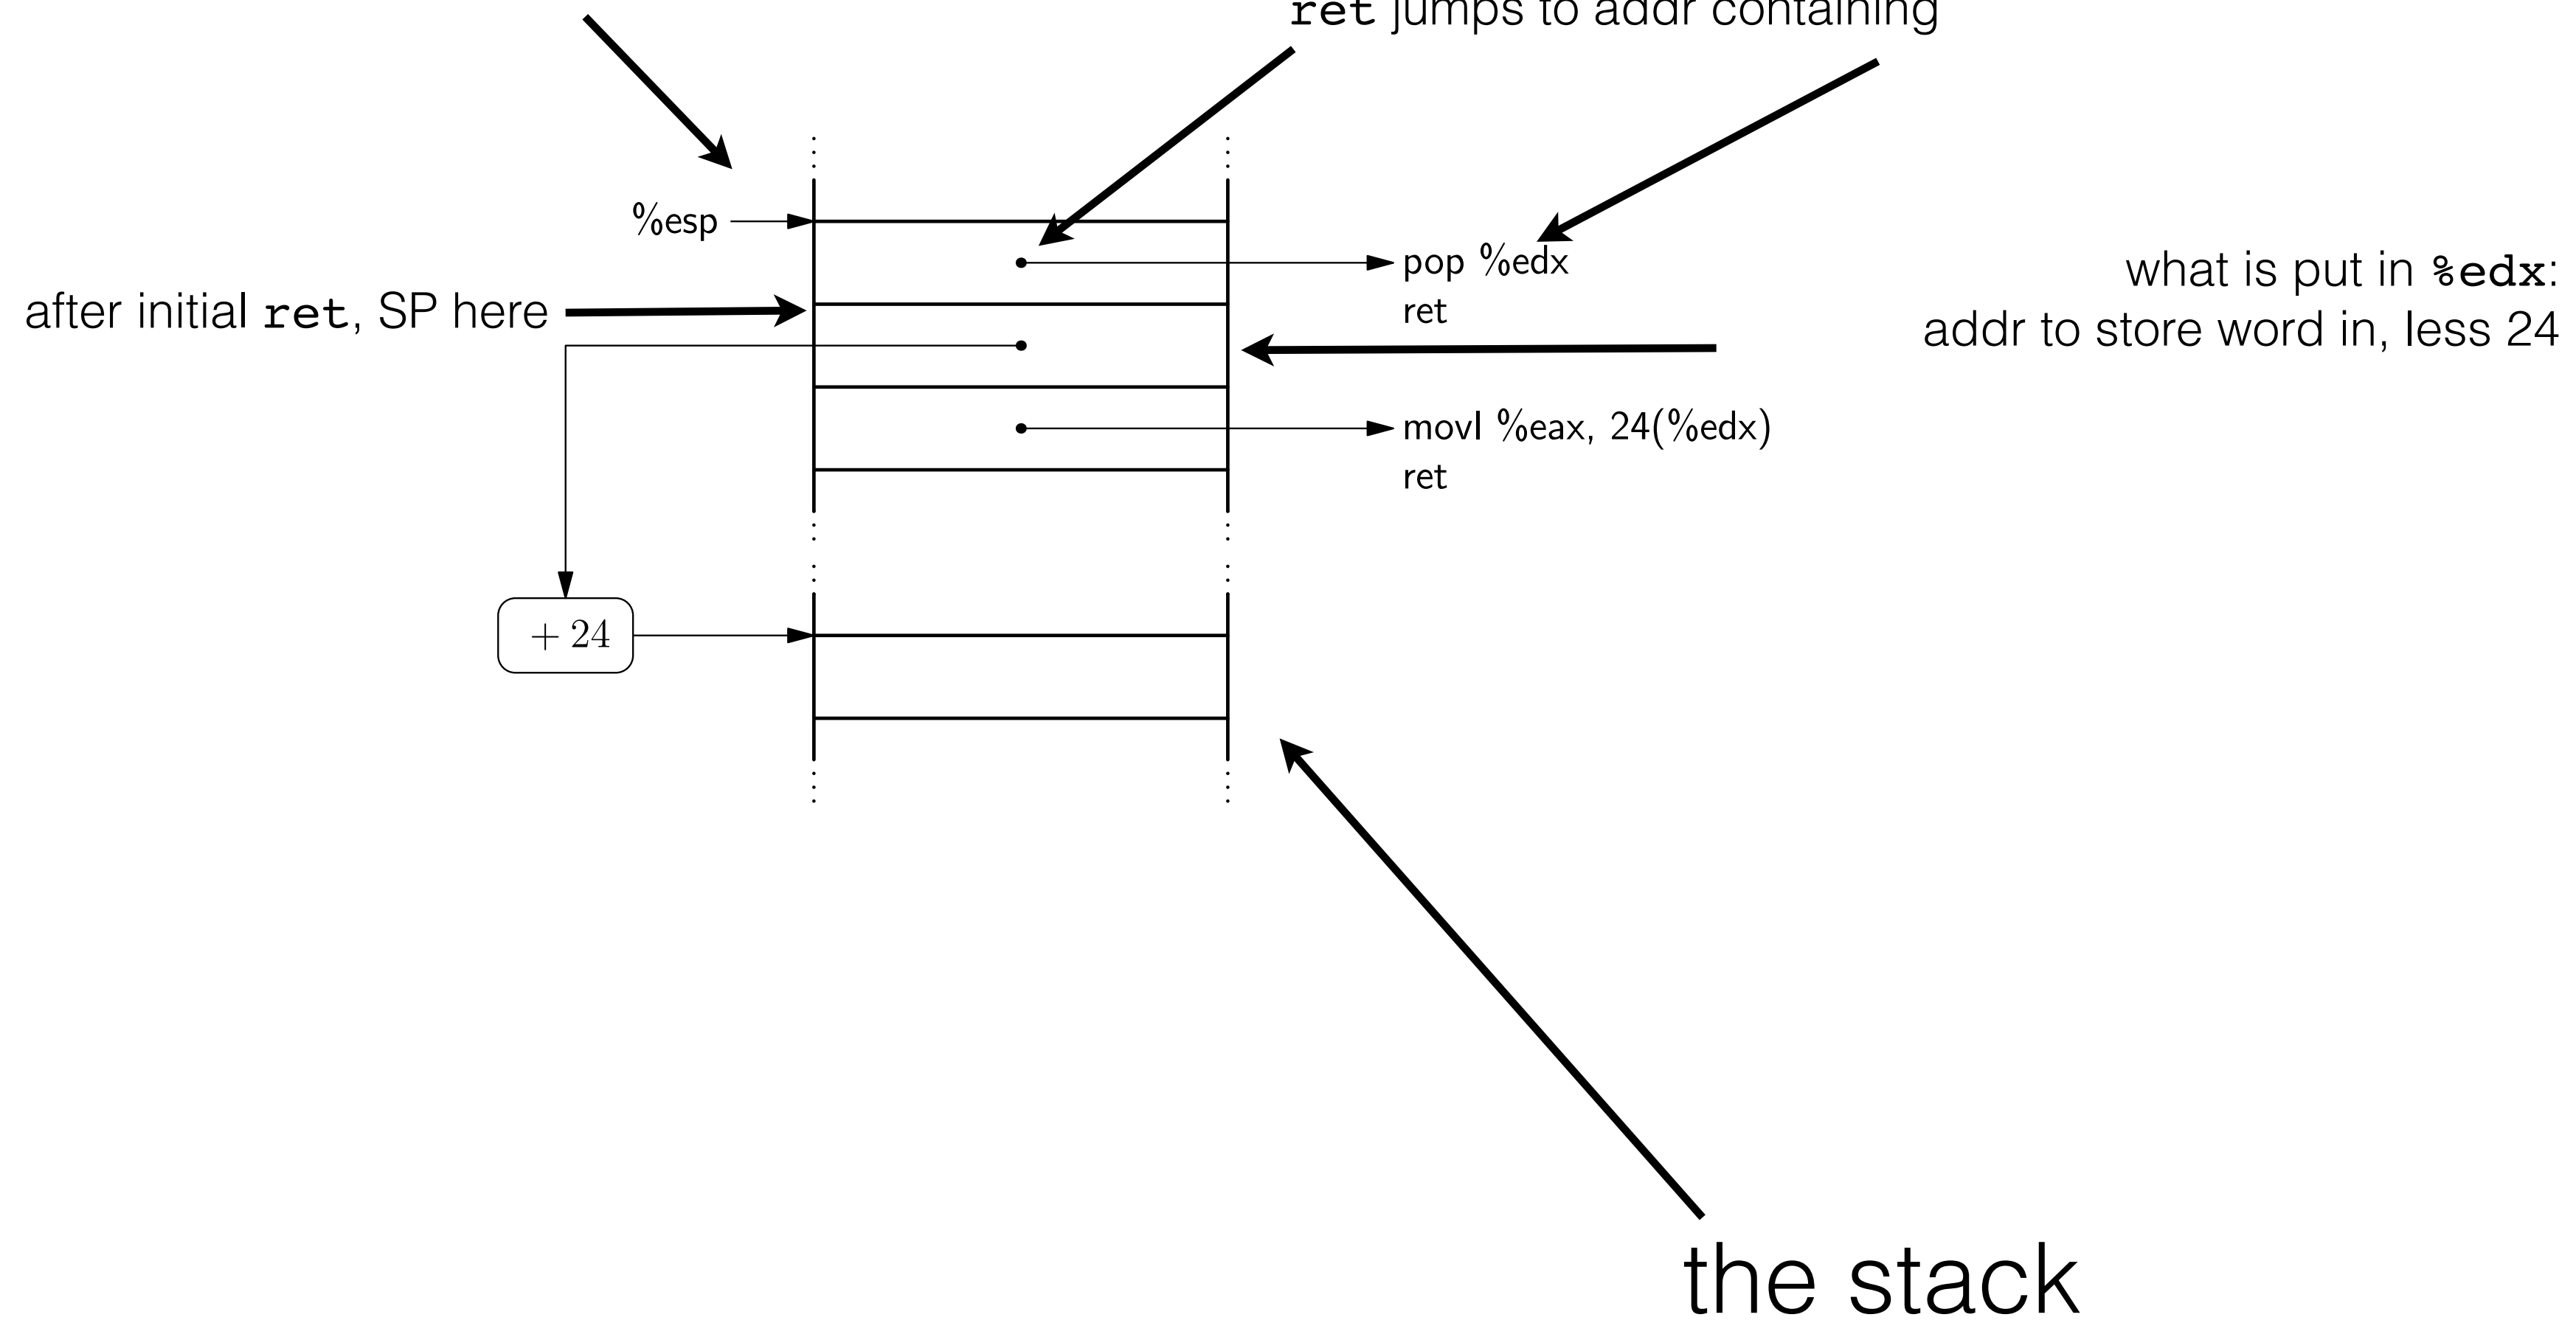


ROP word store

word in `%eax` to
memory

ret issued when SP here

ret jumps to addr containing

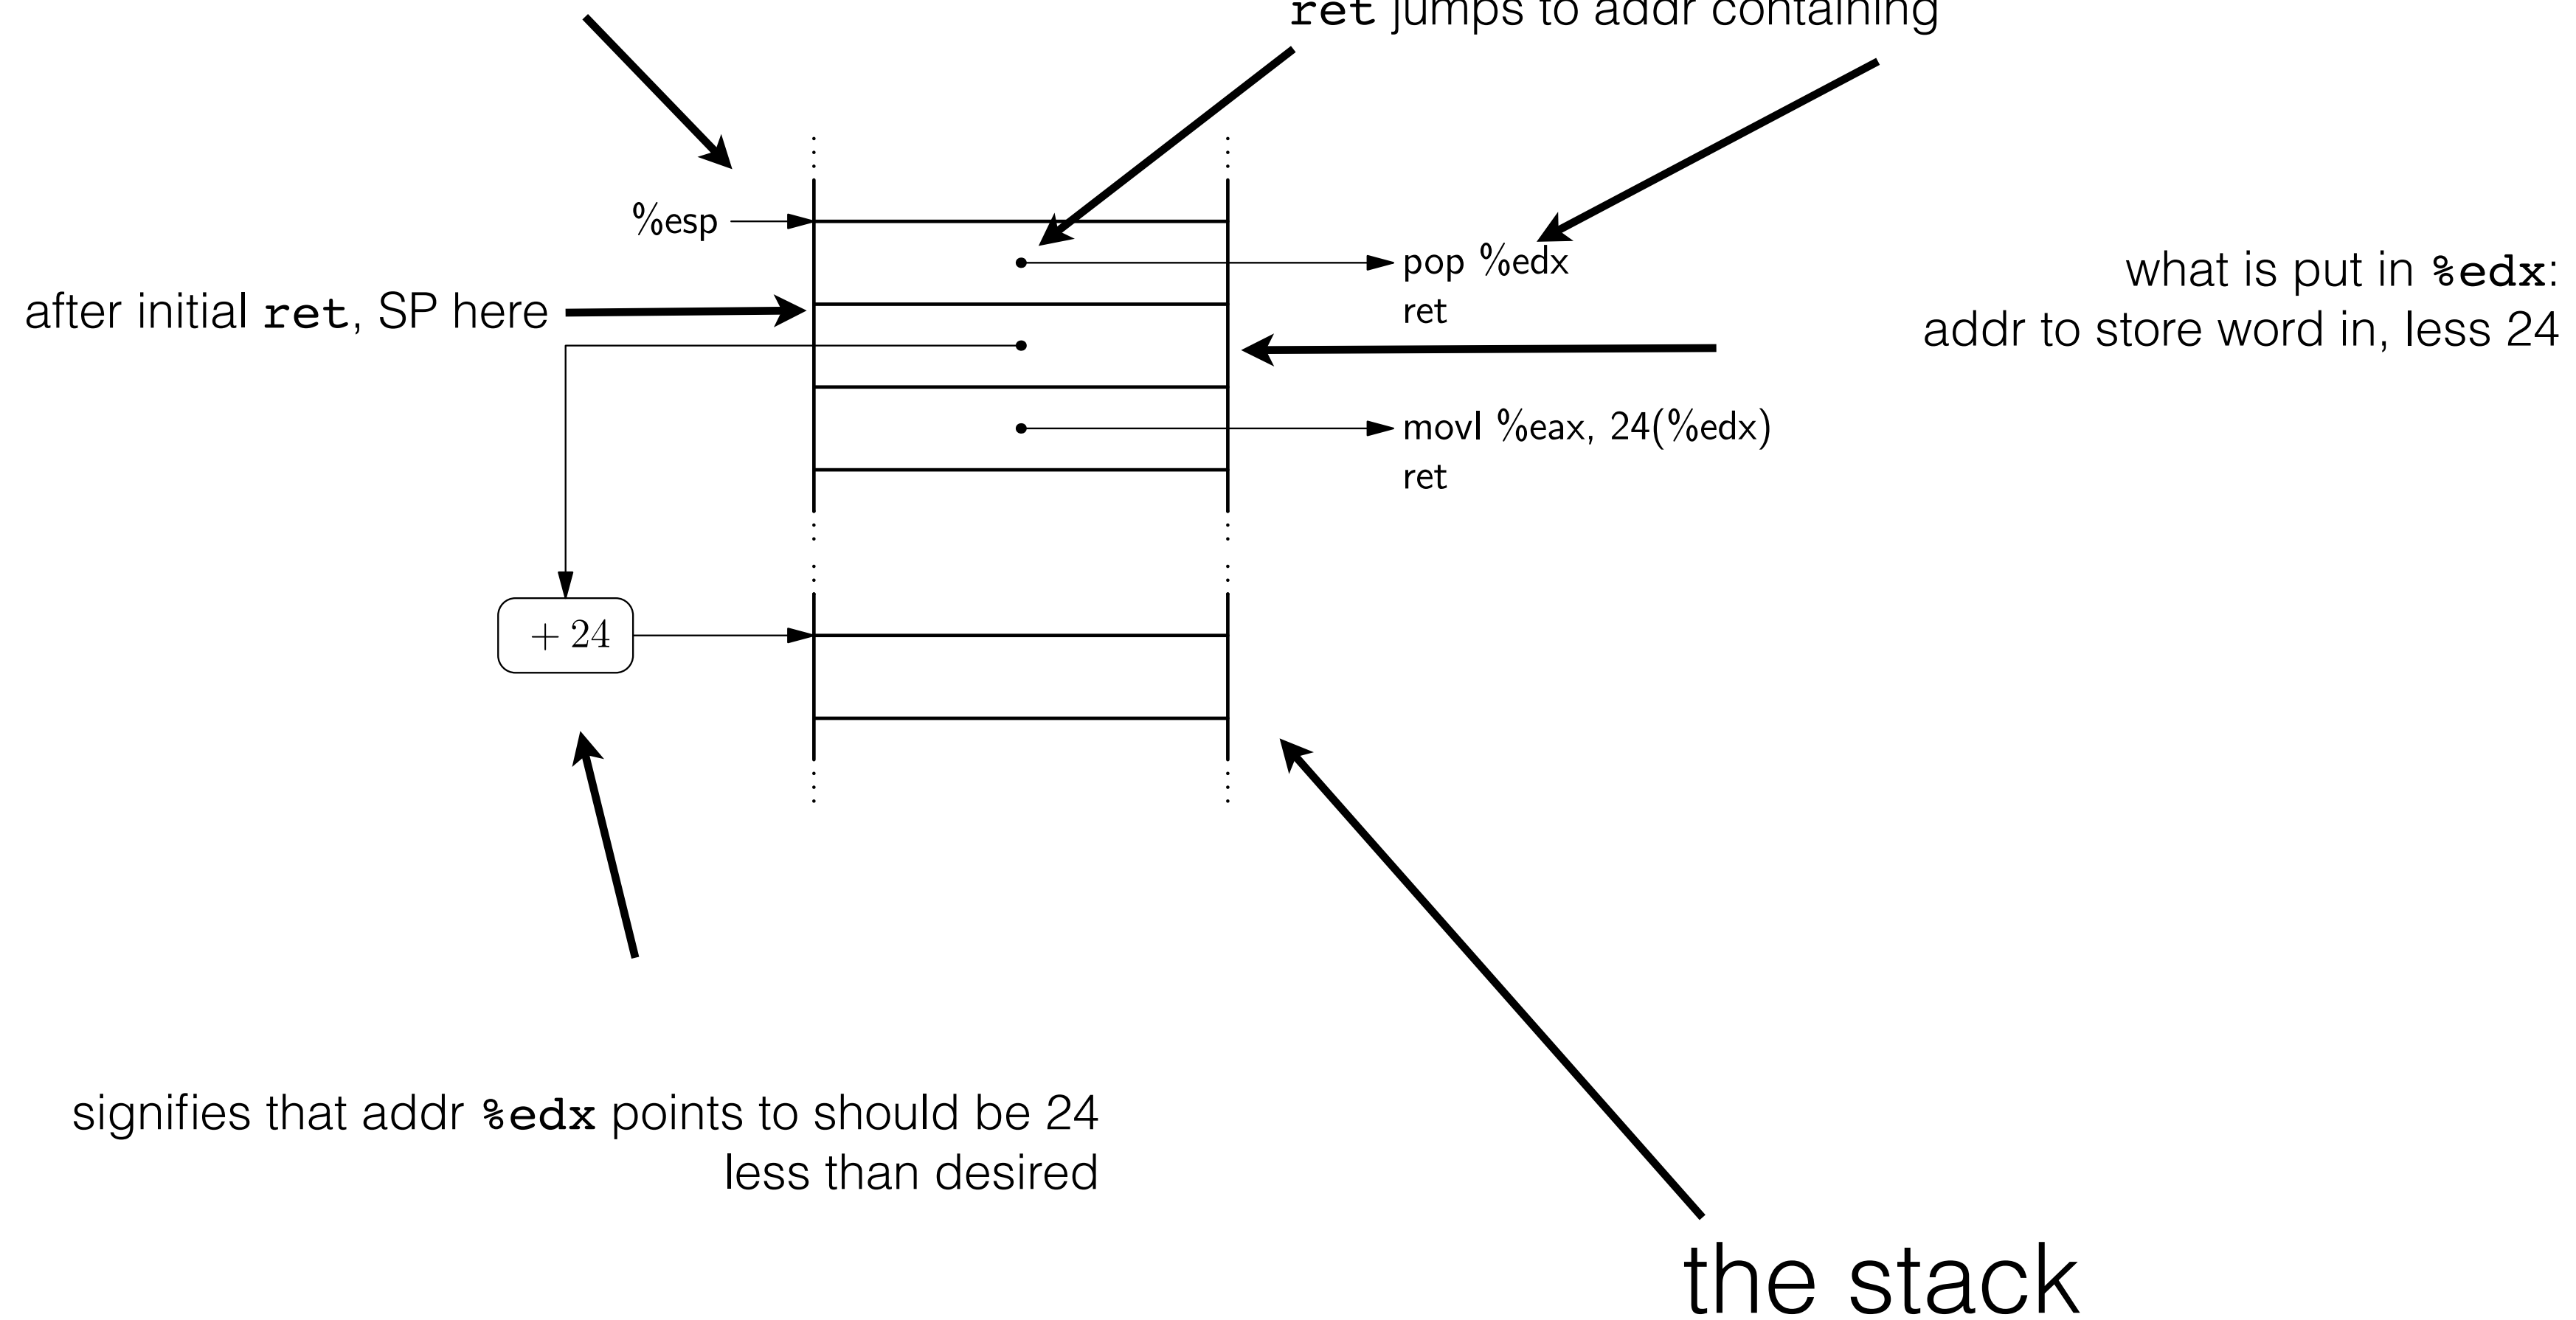


ROP word store

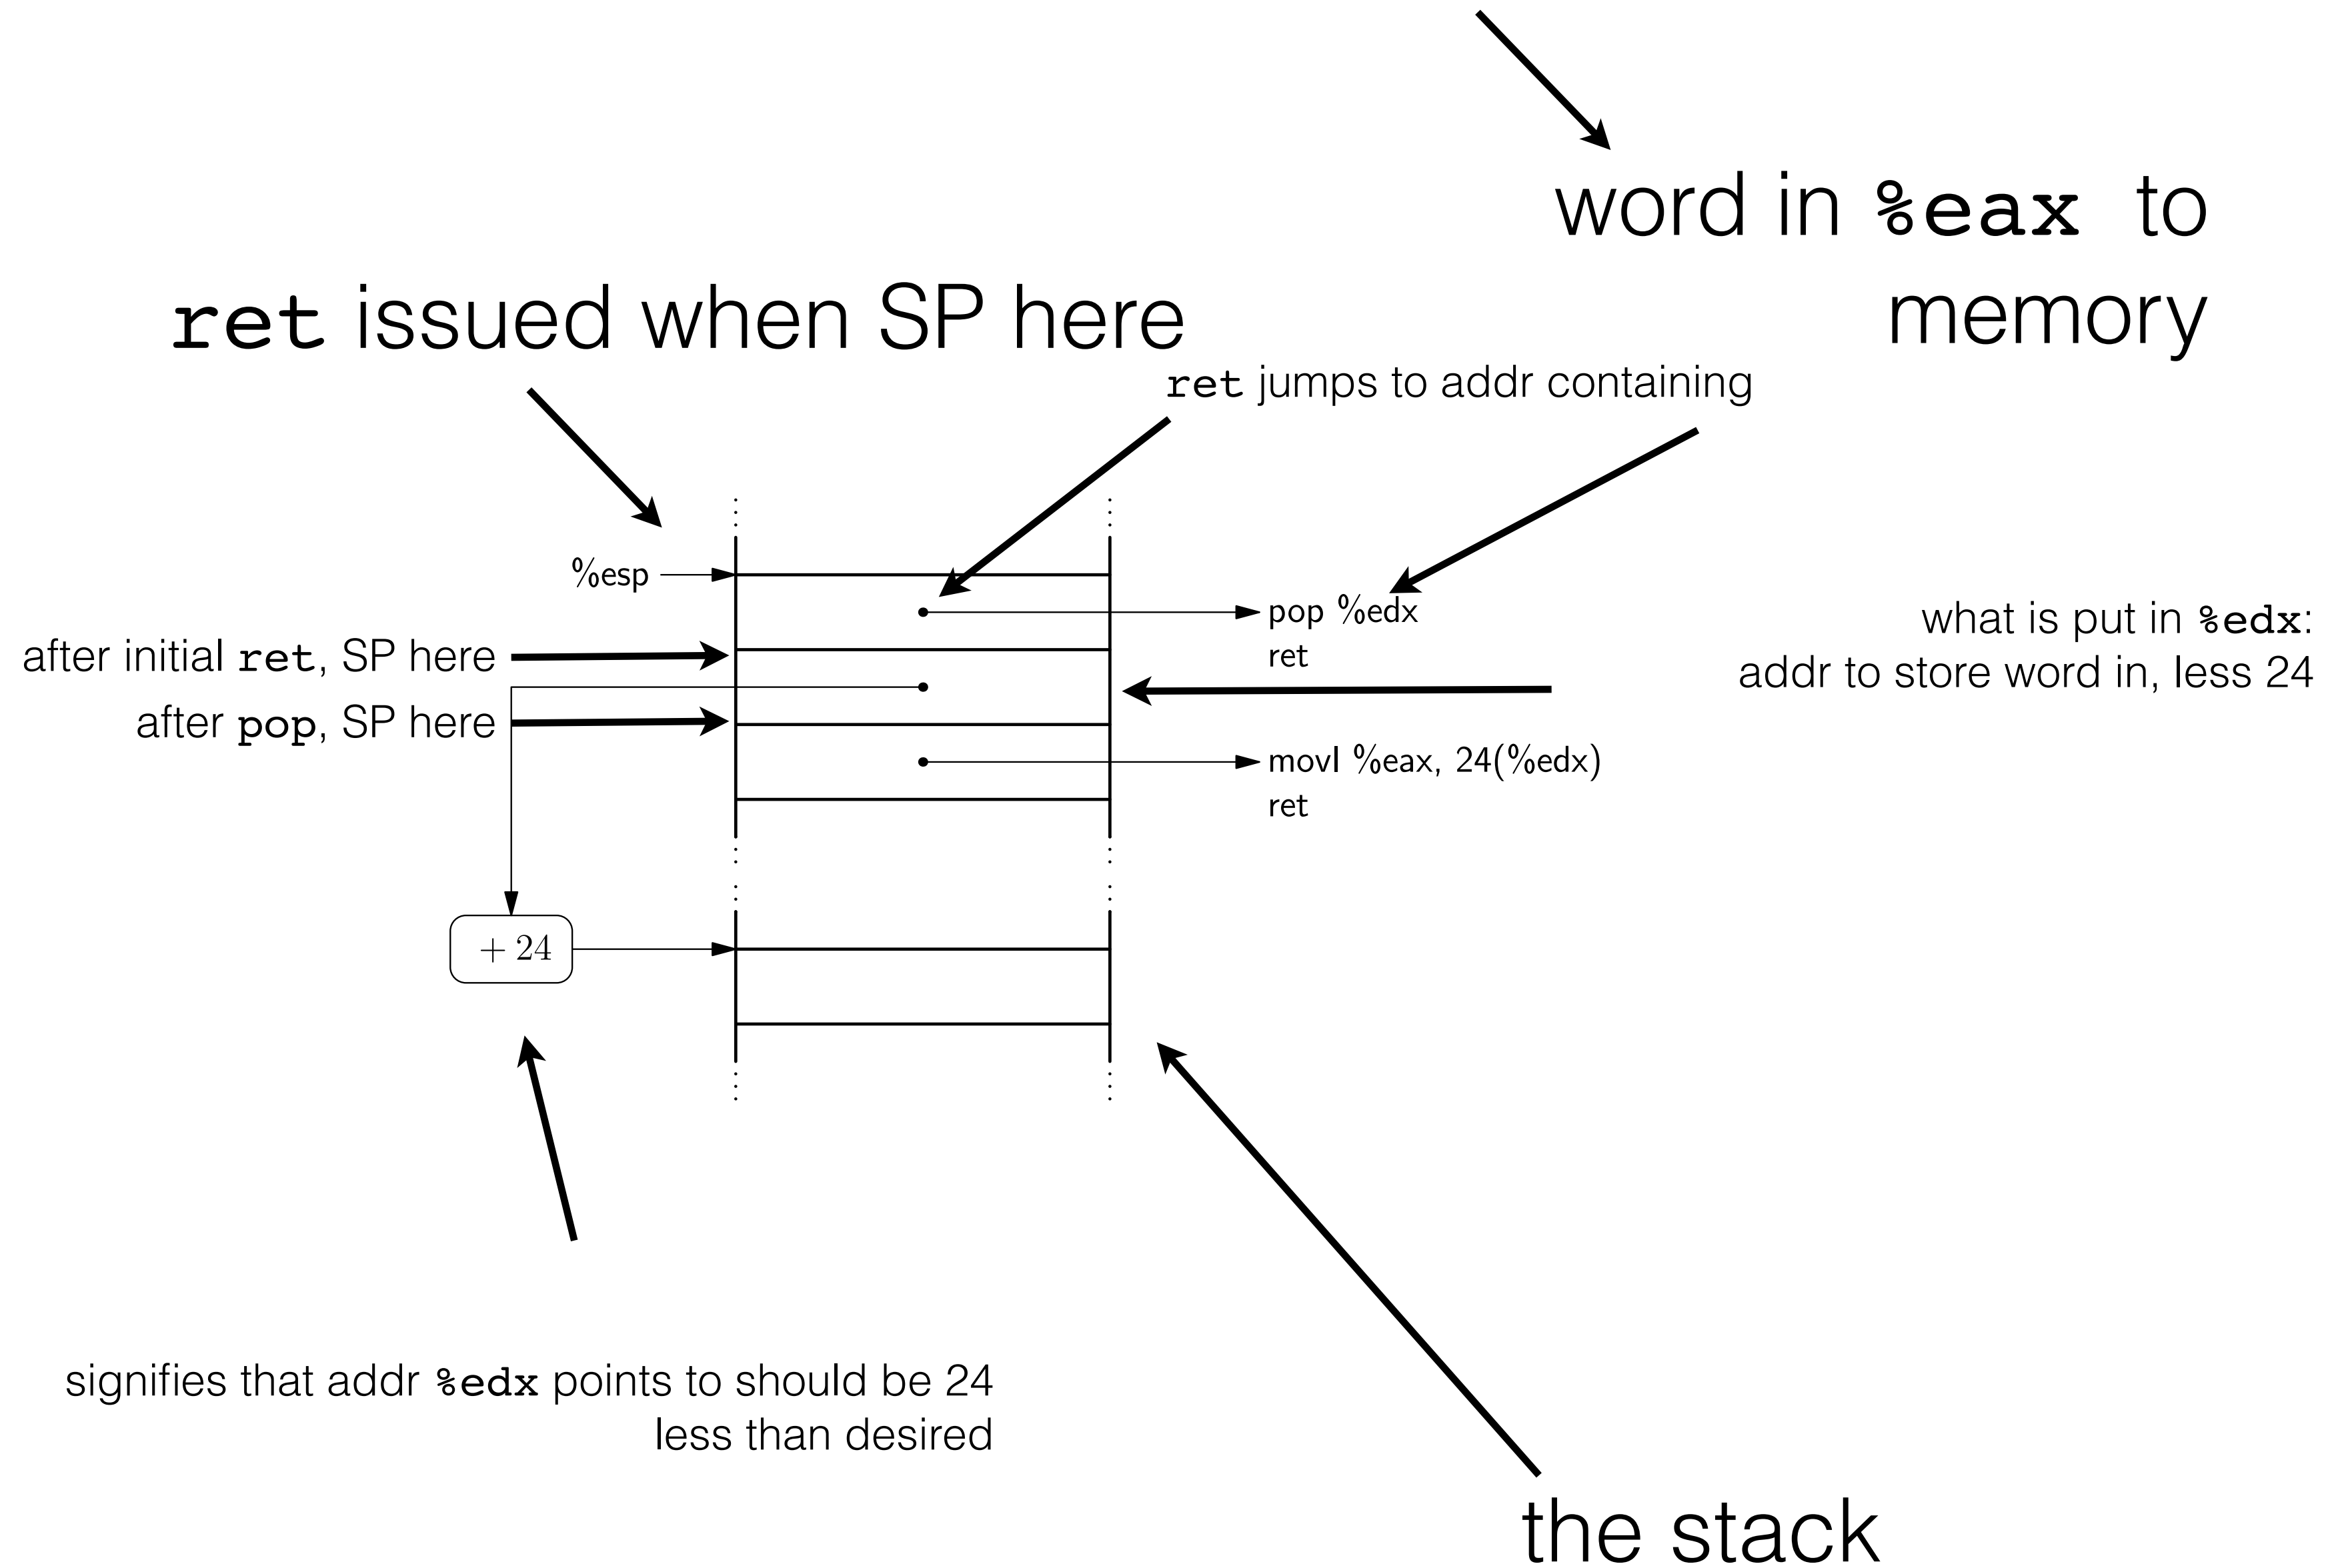
word in `%eax` to
memory

`ret` issued when SP here

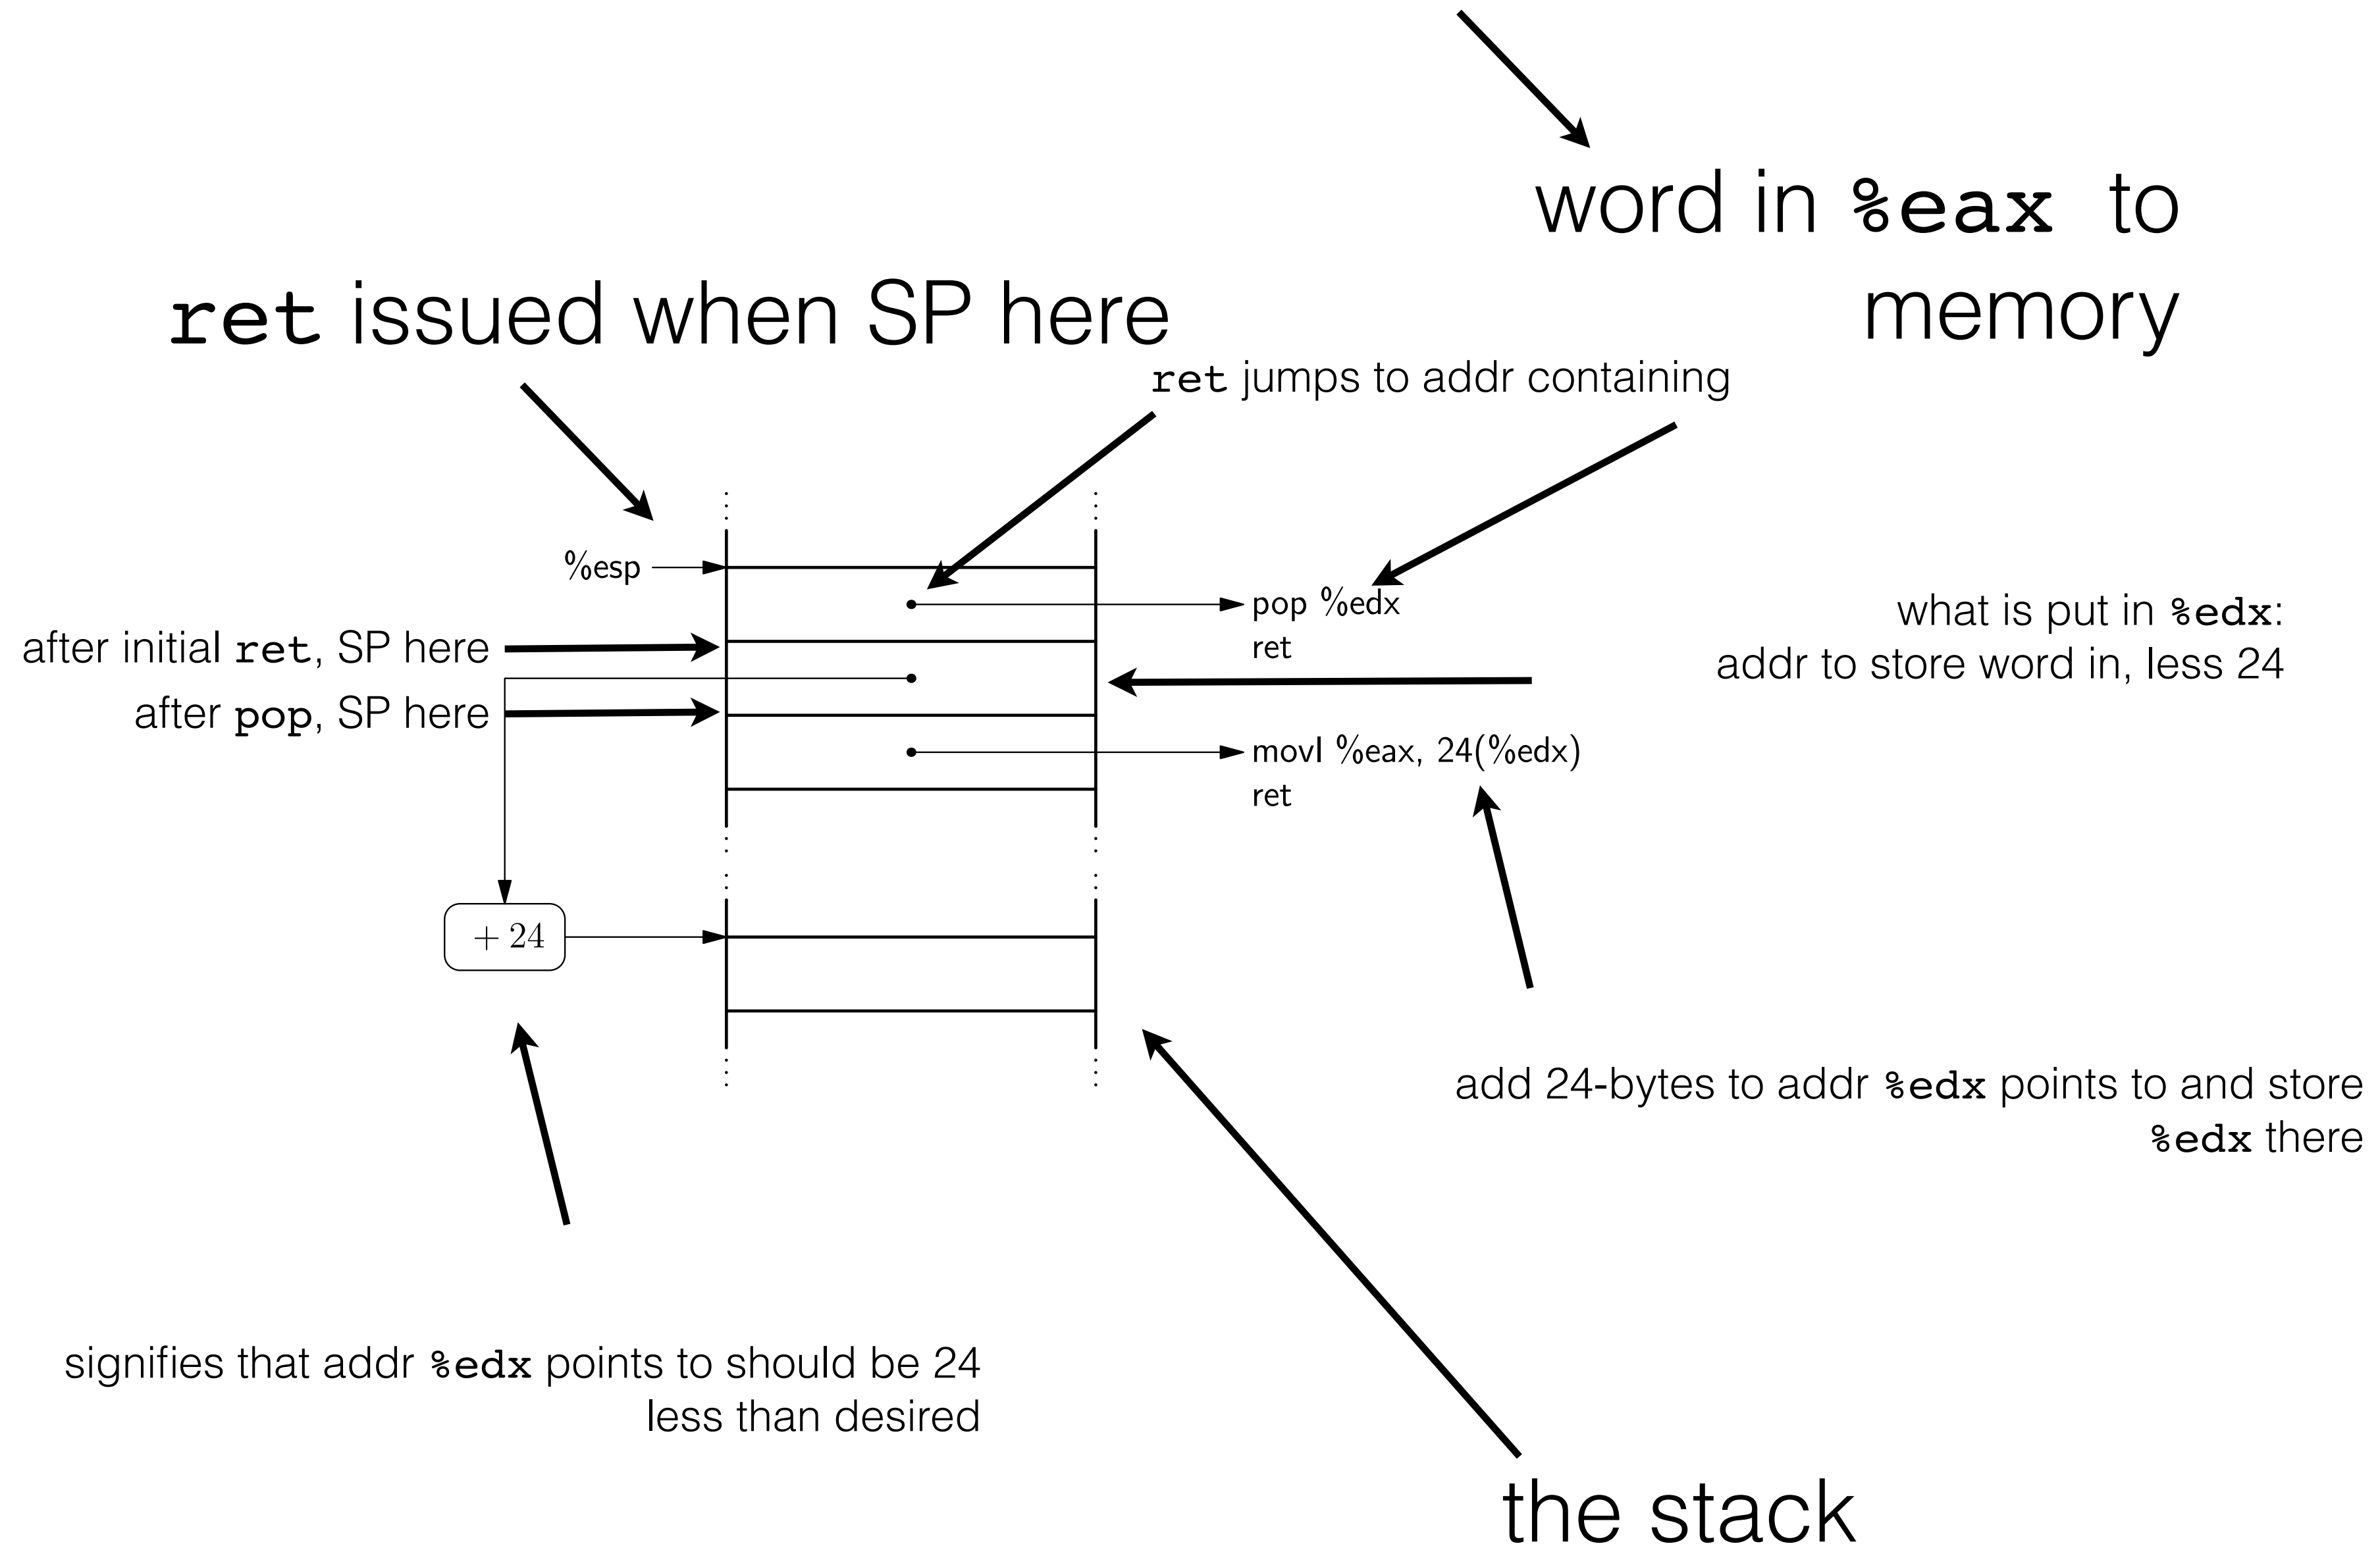
`ret` jumps to addr containing



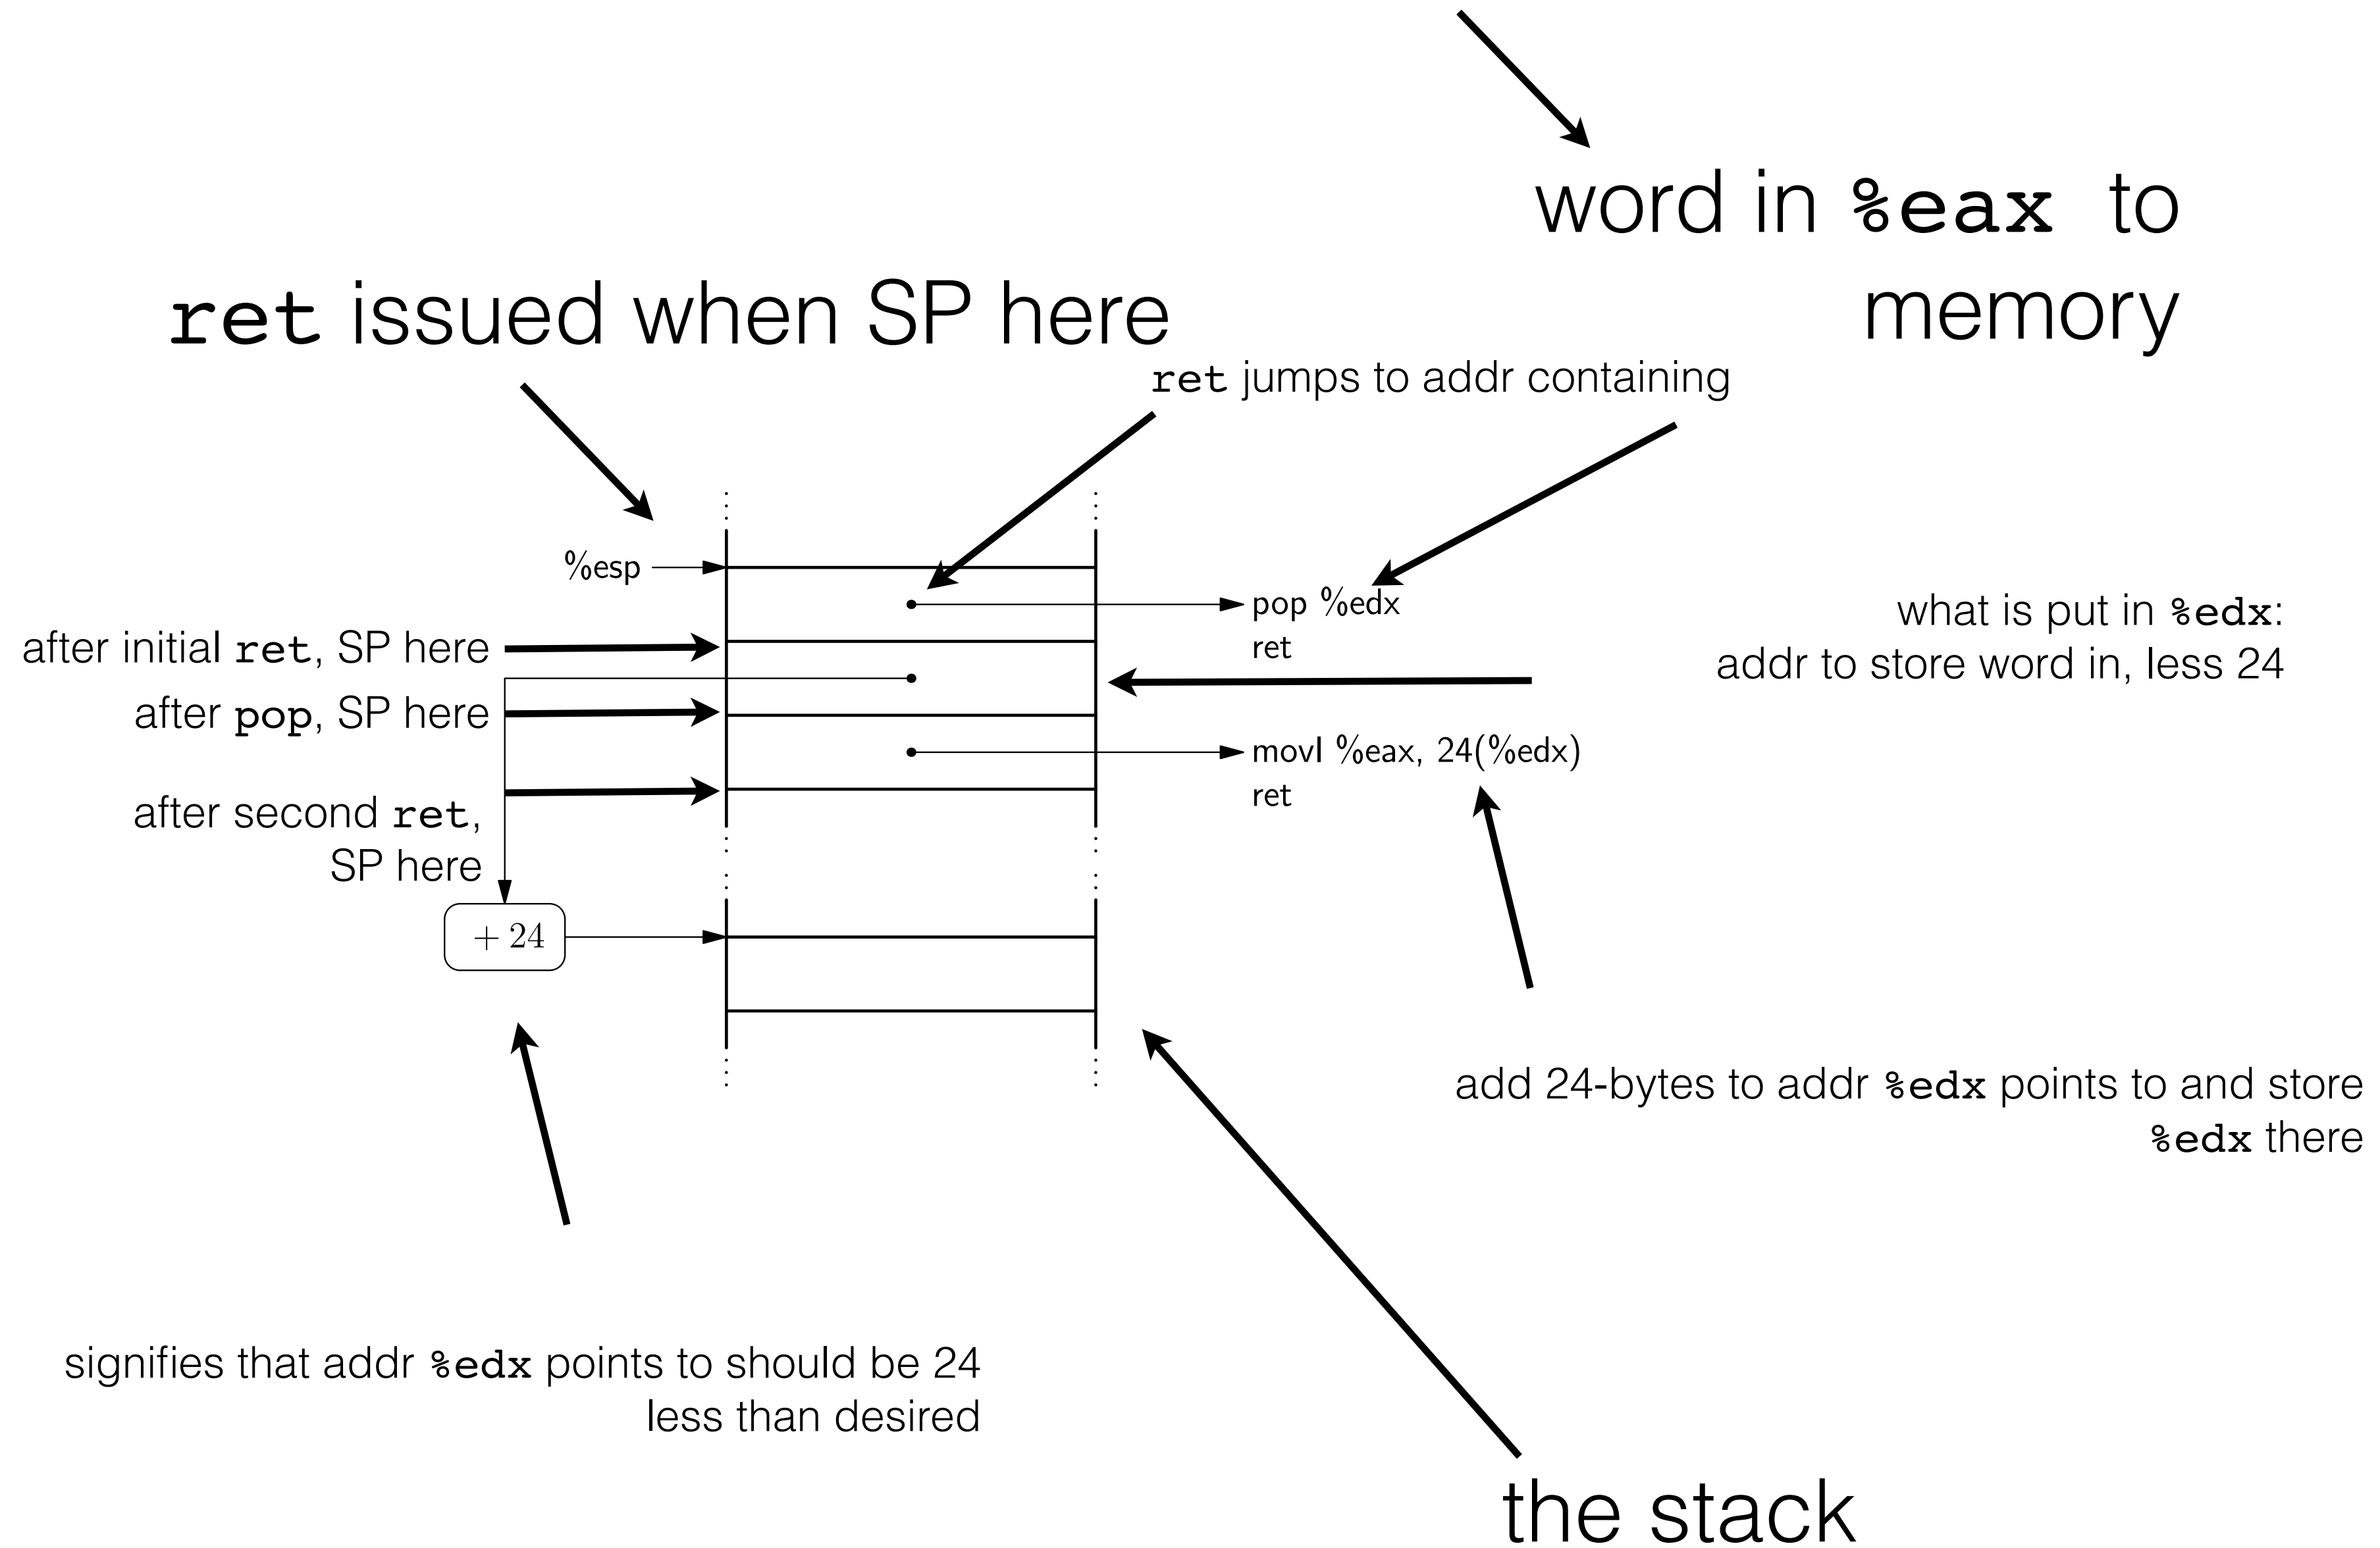
ROP word store



ROP word store



ROP word store



the good news:

can construct sequence of instructions using libc
that are Turing complete

what Turing complete
means to an attacker:



arbitrary computation

what Turing complete
means to an attacker:



arbitrary computation